

Project 2

May 11, 2020

1 Final Results

- Best Model based on Kaggle Public Scoreboard (0.83182) - Cost Sensitive - XGBoost
- Best parameters: {'class_weight': 'balanced', 'learning_rate': 0.01, 'n_estimators': 1140, 'scale_pos_weight': 1150}
- Best Mean cross-validation score: 0.86
- Train score is 0.9926769731489015

- Best Model based on CV scores in Stacking - Combination 3
- Best Mean Cross Validation Score is 0.9978266297321057
- Best Parameters are {'stack_method': 'auto'}
- Train score is 0.9535655058043118

- Note: scores depicted here on heading next to model names are Kaggle Public Scores

2 EDA

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
pd.pandas.set_option('display.max_columns', None)
%matplotlib inline

# for the model
from sklearn.model_selection import train_test_split
```

```
[2]: data = pd.read_csv(r'C:\Users\nabhs\OneDrive\BUAN - Semester 2\BUAN 6341 - Applied Machine Learning\Project 2\mis6341-project2\train.csv')
test= pd.read_csv(r'C:\Users\nabhs\OneDrive\BUAN - Semester 2\BUAN 6341 - Applied Machine Learning\Project 2\mis6341-project2\test.csv')
```

```
[3]: pd.set_option('display.max_columns', 999)
data.head()
```

```
[3]:      Id      V1      V2      V3      V4      V5      V6 \
0  138662 -0.711273  1.272483  1.681631  0.039897 -0.312381 -1.161083
```

1	235999	1.990679	-0.128465	-1.700556	0.529839	0.136889	-1.404270
2	245376	1.812653	-0.476162	-0.338988	1.386750	-0.745965	-0.449870
3	202483	-1.014219	0.522775	-0.337978	-1.957797	3.578395	3.266965
4	9710	-0.967270	-0.053815	2.273463	-0.894434	0.728924	0.153524

	V7	V8	V9	V10	V11	V12	V13 \
0	0.853875	-0.173979	-0.438558	-0.053805	0.321036	0.814690	1.226450
1	0.593353	-0.456700	0.648260	-0.071353	-0.982339	0.229658	-0.592959
2	-0.492226	0.029904	1.129394	0.227346	-1.197477	-0.555129	-1.369130
3	0.602857	0.644645	-0.305879	-0.417906	0.010333	-0.299630	-0.379164
4	-0.038330	0.165252	1.172526	-0.879114	2.505831	-2.249565	0.791677

	V14	V15	V16	V17	V18	V19	V20 \
0	-0.152389	0.687480	-0.030825	-0.284477	-0.696867	-0.227419	0.298008
1	0.574159	-0.263545	-0.780886	-0.066039	-0.575168	0.227544	-0.205092
2	0.320282	1.002745	0.501765	-0.670053	0.478236	-0.820942	-0.202894
3	0.329193	0.957904	-0.417285	-0.705683	-0.063235	-0.381035	0.136061
4	1.546279	-0.133370	0.184501	0.307080	-0.081722	-1.359941	0.007512

	V21	V22	V23	V24	V25	V26	V27 \
0	-0.208181	-0.414631	0.057957	0.930321	-0.165598	0.038799	0.384295
1	0.018179	0.179084	-0.004248	0.039531	0.333765	-0.225201	-0.047285
2	0.252668	0.655719	0.090057	-0.136884	-0.172452	-0.556921	0.036490
3	0.147742	0.590740	-0.680551	0.745346	1.119496	0.013520	0.004530
4	0.106692	0.514430	-0.081546	-0.305042	-0.354883	0.896254	-0.182293

	V28	V29	Target
0	0.192762	21.87	0
1	-0.059682	49.99	0
2	-0.021575	79.00	0
3	-0.132643	6.00	0
4	-0.142810	15.95	0

```
[4]: data.shape
```

```
[4]: (24846, 31)
```


```
[5]: test.shape
```

```
[5]: (24846, 30)
```

2.1 Correlation Matrix

```
[134]: fig, ax = plt.subplots(figsize=(25,30))           # Sample figsize in inches
        sns.heatmap(data.corr(), annot=True, linewidths=.5, ax=ax)
```

```
[134]: <matplotlib.axes._subplots.AxesSubplot at 0x25759d6af08>
```



2.2 Split in Test/Train

```
[6]: X_train = data.drop(['Target', 'Id'], axis=1)
      y_train = data['Target']

      X_train.shape, y_train.shape
```

```
[6]: ((24846, 29), (24846,))
```

```
[7]: X_test=test.drop(['Id'], axis=1)
      X_test.shape
```

```
[7]: (24846, 29)
```

2.3 Check Explanatory Variable Types

```
[8]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24846 entries, 0 to 24845
Data columns (total 31 columns):
Id          24846 non-null int64
V1          22459 non-null float64
V2          24846 non-null float64
V3          24846 non-null float64
V4          24846 non-null float64
V5          24846 non-null float64
V6          24846 non-null float64
V7          24846 non-null float64
V8          24846 non-null float64
V9          24846 non-null float64
V10         24846 non-null float64
V11         24846 non-null float64
V12         24846 non-null float64
V13         24846 non-null float64
V14         24846 non-null float64
V15         24846 non-null float64
V16         24846 non-null float64
V17         24846 non-null float64
V18         24846 non-null float64
V19         24846 non-null float64
V20         22317 non-null float64
V21         24846 non-null float64
V22         24846 non-null float64
V23         24846 non-null float64
V24         24846 non-null float64
V25         24846 non-null float64
```

```
V26      24846 non-null float64
V27      24846 non-null float64
V28      24846 non-null float64
V29      24846 non-null float64
Target    24846 non-null int64
dtypes: float64(29), int64(2)
memory usage: 5.9 MB
```

2.4 Categorize Variables acc. to Types

```
[9]: # make list of variables types

# numerical: discrete vs continuous
discrete = [var for var in data.columns if data[var].dtype!='0' and var!
            ↳='Target' and data[var].nunique()<10]
continuous = [var for var in data.columns if data[var].dtype!='0' and var!
              ↳='Target' and var not in discrete]

# categorical
categorical = [var for var in data.columns if data[var].dtype=='0']

print(f'There are {len(discrete)} discrete variables')
print(f'There are {len(continuous)} continuous variables')
print(f'There are {len(categorical)} categorical variables')
```

```
There are 0 discrete variables
There are 30 continuous variables
There are 0 categorical variables
```


```
[10]: len(data.columns)
```

```
[10]: 31
```

2.5 Outliers in Continuous

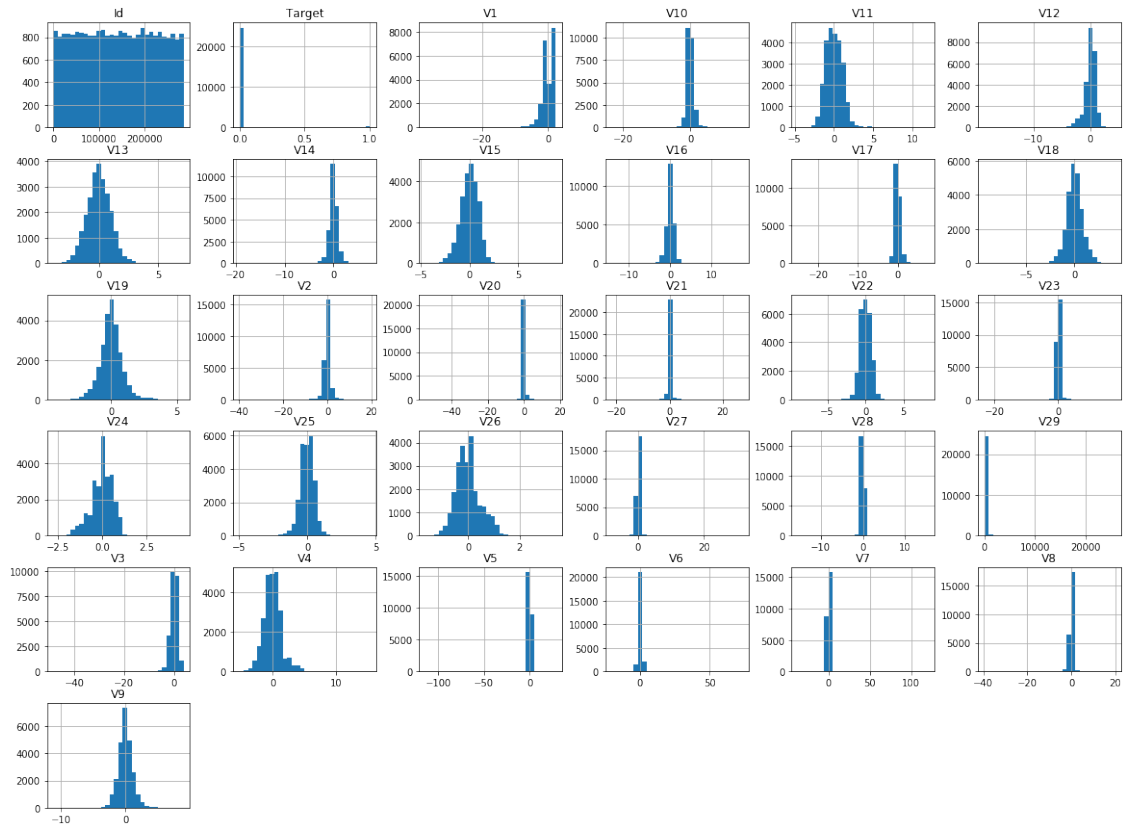
```
[11]: data[continuous].boxplot(figsize=(20,15))
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x2573d6dbbc8>
```



2.6 Check for Normality

```
[12]: # numerical
data.select_dtypes(exclude='O').hist(bins=30, figsize=(20,15))
plt.show()
```



2.7 Check for Null Values

```
[13]: X_train.isnull().mean()
```

```
[13]: V1      0.096072
      V2      0.000000
      V3      0.000000
      V4      0.000000
      V5      0.000000
      V6      0.000000
      V7      0.000000
      V8      0.000000
      V9      0.000000
      V10     0.000000
      V11     0.000000
      V12     0.000000
      V13     0.000000
      V14     0.000000
      V15     0.000000
      V16     0.000000
      V17     0.000000
```

```
V18    0.000000
V19    0.000000
V20    0.101787
V21    0.000000
V22    0.000000
V23    0.000000
V24    0.000000
V25    0.000000
V26    0.000000
V27    0.000000
V28    0.000000
V29    0.000000
dtype: float64
```

2.8 Pre-Processing

```
[14]: # import relevant modules for feature engineering
      from sklearn.pipeline import Pipeline
      from feature_engine import missing_data_imputers as mdi
      from feature_engine import categorical_encoders as ce
      from sklearn.preprocessing import StandardScaler
```

```
[15]: continuous=list(set(list(continuous))-set('Id'))
```

```
[16]: continuous.remove('Id')
```

```
[17]: continuous
```

```
[17]: ['V6',
      'V2',
      'V13',
      'V17',
      'V21',
      'V11',
      'V20',
      'V27',
      'V8',
      'V22',
      'V24',
      'V26',
      'V10',
      'V12',
      'V16',
      'V1',
      'V7',
      'V25',
      'V15',
```



```
'V9',  
'V19',  
'V4',  
'V23',  
'V29',  
'V18',  
'V3',  
'V5',  
'V14',  
'V28']
```

```
[18]: project2_preprocess = Pipeline([  
    ('imputer_num', mdi.MeanMedianImputer(imputation_method='median',  
                                           variables=continuous)),  
    # feature Scaling  
    ('scaler', StandardScaler())  
])
```

```
[19]: project2_preprocess.fit(X_train,y_train)
```

```
[19]: Pipeline(memory=None,  
    steps=[('imputer_num',  
            MeanMedianImputer(imputation_method='median',  
                               variables=['V6', 'V2', 'V13', 'V17', 'V21',  
                                          'V11', 'V20', 'V27', 'V8', 'V22',  
                                          'V24', 'V26', 'V10', 'V12', 'V16',  
                                          'V1', 'V7', 'V25', 'V15', 'V9',  
                                          'V19', 'V4', 'V23', 'V29', 'V18',  
                                          'V3', 'V5', 'V14', 'V28'])),  
            ('scaler',  
             StandardScaler(copy=True, with_mean=True, with_std=True))],  
    verbose=False)
```

```
[20]: # Apply Transformations  
X_train=project2_preprocess.transform(X_train)  
X_test=project2_preprocess.transform(X_test)
```

3 DO NOT CHANGE STEPS BEFORE THIS POINT

3.1 Apply Basic Models

```
[21]: from numpy import mean  
from numpy import std  
from sklearn.datasets import make_classification  
from sklearn.model_selection import cross_val_score  
from sklearn.model_selection import RepeatedStratifiedKfold
```

```

from sklearn.dummy import DummyClassifier
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import auc
from sklearn.metrics import make_scorer
from matplotlib import pyplot
from sklearn.model_selection import GridSearchCV

```

```

[22]: from sklearn.model_selection import cross_val_score
      from imblearn.metrics import geometric_mean_score
      from sklearn.metrics import make_scorer, fbeta_score

```

```

[23]: f2score = make_scorer(fbeta_score, beta=2)

```

3.1.1 Logistic Regression - .80909

```

[24]: from sklearn.linear_model import LogisticRegression

param_logit = {'C': [0.001, 0.01, 0.1, 1, 10, 1000],
               'solver': ['lbfgs'],
               'penalty': ['l2'],
               'max_iter': range(1000, 1500, 10)}
print("Parameter grid:\n{}".format(param_logit))

grid_logit = GridSearchCV(LogisticRegression(random_state=42), param_logit,
    ↪cv=5, return_train_score=True, scoring = f2score)
grid_logit.fit(X_train, y_train)

```

Parameter grid:

```

{'C': [0.001, 0.01, 0.1, 1, 10, 1000], 'solver': ['lbfgs'], 'penalty': ['l2'],
 'max_iter': range(1000, 1500, 10)}

```

```

[24]: GridSearchCV(cv=5, error_score=nan,
                  estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                  fit_intercept=True,
                  intercept_scaling=1, l1_ratio=None,
                  max_iter=100, multi_class='auto',
                  n_jobs=None, penalty='l2',
                  random_state=42, solver='lbfgs',
                  tol=0.0001, verbose=0,
                  warm_start=False),
                  iid='deprecated', n_jobs=None,
                  param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 1000],
                  'max_iter': range(1000, 1500, 10), 'penalty': ['l2'],
                  'solver': ['lbfgs']},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)

```

Results

```
[333]: print("Best parameters: {}".format(grid_logit.best_params_))
print("Best cross-validation score: {:.2f}".format(grid_logit.best_score_))
print('Train score: {:.4f}'.format(grid_logit.score(X_train, y_train)))
```

Best parameters: {'C': 10, 'max_iter': 1000, 'penalty': 'l2', 'solver': 'lbfgs'}
Best cross-validation score: 0.80
Train score: 0.8177

3.1.2 Decision Tree

```
[26]: from sklearn.tree import DecisionTreeClassifier
```

```
[27]: dtree = DecisionTreeClassifier(random_state=42)

#define a list of parameters
param_dtree = {'max_depth': range(1,20)}

#apply grid search
grid_dtree = GridSearchCV(dtree, param_dtree, cv=5, return_train_score = _
    ↪ True, scoring = f2score)
grid_dtree.fit(X_train, y_train)
```

```
[27]: GridSearchCV(cv=5, error_score=nan,
                  estimator=DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None,
                                                    criterion='gini', max_depth=None,
                                                    max_features=None,
                                                    max_leaf_nodes=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    presort='deprecated',
                                                    random_state=42,
                                                    splitter='best'),
                  iid='deprecated', n_jobs=None,
                  param_grid={'max_depth': range(1, 20)}, pre_dispatch='2*n_jobs',
                  refit=True, return_train_score=True,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[335]: #find best parameters
print('Decision Tree parameters: ', grid_dtree.best_params_)
# Mean Cross Validation Score
print("Best Mean Cross-validation score: {:.2f}".format(grid_dtree.best_score_))
```

```
print('Decision Tree Train score: {:.4f}'.format(grid_dtree.score(X_train, y_train)))
```


Decision Tree parameters: {'max_depth': 5}

Best Mean Cross-validation score: 0.81

Decision Tree Train score: 0.8742

```
[137]: %matplotlib inline
def plot_feature_importances_cancer(model):
    n_features = X_train.shape[1]
    plt.barh(range(n_features), model.best_estimator_.feature_importances_, align='center')
    plt.xticks(np.arange(n_features), X_train.dtype.names)
    plt.xlabel("Feature importance")
    plt.ylabel("Feature")
    plt.ylim(-1, n_features)

plot_feature_importances_cancer(grid_dtree)
```



3.1.3 KNN - .82289

```
[360]: from sklearn.neighbors import KNeighborsClassifier
```

```
[361]: # Train a KNN model, report the coefficients, the best parameters, and model
        ↳performance
        # hint: find the optimal k

knn = KNeighborsClassifier()

# define a list of parameters
param_knn = {'n_neighbors': range(1,10)}

#apply grid search
grid_knn = GridSearchCV(knn, param_knn, cv=5,
        ↳return_train_score=True,scoring=f2score)
grid_knn.fit(X_train, y_train)
```

```
[361]: GridSearchCV(cv=5, error_score=nan,
                  estimator=KNeighborsClassifier(algorithm='auto', leaf_size=30,
                                                  metric='minkowski',
                                                  metric_params=None, n_jobs=None,
                                                  n_neighbors=5, p=2,
                                                  weights='uniform'),
                  iid='deprecated', n_jobs=None,
                  param_grid={'n_neighbors': range(1, 10)}, pre_dispatch='2*n_jobs',
                  refit=True, return_train_score=True,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[363]: #find best parameters
print("KNN parameters: {}".format(grid_knn.best_params_))
# Mean Cross Validation Score
print("Best Mean Cross-validation score: {:.2f}".format(grid_knn.best_score_))
print('Train score: {:.4f}'.format(grid_knn.score(X_train, y_train)))
```

```
KNN parameters: {'n_neighbors': 1}
Best Mean Cross-validation score: 0.83
Train score: 1.0000
```

3.1.4 Kernel SVC

```
[32]: from sklearn.svm import SVC
param_svc = {'C':range(1,50,5),
             'kernel':['rbf','sigmoid'],
             'gamma': ['auto']}

[33]: grid_svc = GridSearchCV(SVC(), param_svc, cv=5,
        ↳return_train_score=True,scoring=f2score)
grid_svc.fit(X_train, y_train)
```

```
[33]: GridSearchCV(cv=5, error_score=nan,
                  estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                                class_weight=None, coef0=0.0,
                                decision_function_shape='ovr', degree=3,
                                gamma='scale', kernel='rbf', max_iter=-1,
                                probability=False, random_state=None, shrinking=True,
                                tol=0.001, verbose=False),
                  iid='deprecated', n_jobs=None,
                  param_grid={'C': range(1, 50, 5), 'gamma': ['auto'],
                              'kernel': ['rbf', 'sigmoid']},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[336]: print("Best parameters: {}".format(grid_svc.best_params_))
        print("Best cross-validation score: {:.2f}".format(grid_svc.best_score_))
        print('Train score: {:.4f}'.format(grid_svc.score(X_train, y_train)))
```

Best parameters: {'C': 21, 'gamma': 'auto', 'kernel': 'rbf'}
 Best cross-validation score: 0.73
 Train score: 0.9636

3.1.5 Linear SVC - .81501

```
[35]: from sklearn.svm import LinearSVC
        from sklearn.metrics import accuracy_score
        linear_svm = LinearSVC(max_iter=3000, dual=False).fit(X_train, y_train)
        print("Coefficient shape: ", linear_svm.coef_.shape)
        print("Intercept shape: ", linear_svm.intercept_.shape)
```

Coefficient shape: (1, 29)
 Intercept shape: (1,)

```
[36]: cv_scores_linear = cross_val_score(linear_svm, X_train, y_train, scoring=f2score)
```

Results

```
[337]: print("Cross-validation scores: {}".format(cv_scores_linear))
        print('Train score: {:.4f}'.format(linear_svm.score(X_train, y_train)))
```

Cross-validation scores: [0.78389831 0.84051724 0.75757576 0.80508475
 0.84745763]
 Train score: 0.9976

3.1.6 Random Forest - .81352

```
[57]: from sklearn.ensemble import RandomForestClassifier
```

```
[107]: #random forest
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(random_state=0)
rfc_param = {
    'n_estimators': range(2,15,2),
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth' : range(2,15,3),
    'criterion' :['gini', 'entropy']
}

grid_rf = GridSearchCV(rfc, rfc_param,cv=5, return_train_score=True, scoring_
    ↪=f2score)
grid_rf.fit(X_train,y_train)
```

```
[107]: GridSearchCV(cv=5, error_score=nan,
                  estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
                                                    class_weight=None,
                                                    criterion='gini', max_depth=None,
                                                    max_features='auto',
                                                    max_leaf_nodes=None,
                                                    max_samples=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    n_estimators=100, n_jobs=None,
                                                    oob_score=False, random_state=0,
                                                    verbose=0, warm_start=False),
                  iid='deprecated', n_jobs=None,
                  param_grid={'criterion': ['gini', 'entropy'],
                              'max_depth': range(2, 15, 3),
                              'max_features': ['auto', 'sqrt', 'log2'],
                              'n_estimators': range(2, 15, 2)},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[108]: print(f'Best Mean Cross Validation Score is {grid_rf.best_params_}')
print(f'Best Mean Cross Validation Score is {grid_rf.best_score_}')
print(f'Train score is {grid_rf.score(X_train,y_train)}')
```

Best Mean Cross Validation Score is {'criterion': 'gini', 'max_depth': 5,

```
'max_features': 'auto', 'n_estimators': 12}
Best Mean Cross Validation Score is 0.8500952309334316
Train score is 0.8614864864864864
```

3.1.7 Extra-Trees - .80586

```
[40]: from sklearn.ensemble import ExtraTreesClassifier
```

```
[41]: etc= ExtraTreesClassifier(random_state=42)
etc_param = {
    'n_estimators': range(50,100,10),
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth' : range(2,15,2),
    'criterion' :['gini', 'entropy']
}
etc_grid = GridSearchCV(etc, etc_param,cv=5, return_train_score=True,scoring =_
↪f2score)
etc_grid.fit(X_train,y_train)
```

```
[41]: GridSearchCV(cv=5, error_score=nan,
                  estimator=ExtraTreesClassifier(bootstrap=False, ccp_alpha=0.0,
                                                  class_weight=None, criterion='gini',
                                                  max_depth=None, max_features='auto',
                                                  max_leaf_nodes=None,
                                                  max_samples=None,
                                                  min_impurity_decrease=0.0,
                                                  min_impurity_split=None,
                                                  min_samples_leaf=1,
                                                  min_samples_split=2,
                                                  min_weight_fraction_leaf=0.0,
                                                  n_estimators=100, n_jobs=None,
                                                  oob_score=False, random_state=42,
                                                  verbose=0, warm_start=False),
                  iid='deprecated', n_jobs=None,
                  param_grid={'criterion': ['gini', 'entropy'],
                              'max_depth': range(2, 15, 2),
                              'max_features': ['auto', 'sqrt', 'log2'],
                              'n_estimators': range(50, 100, 10)},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[338]: print(f'Best Mean Cross Validation Score is {etc_grid.best_params_}')
print(f'Best Mean Cross Validation Score is {etc_grid.best_score_}')
print(f'Train score is {etc_grid.score(X_train,y_train)}')
```

```
Best Mean Cross Validation Score is {'criterion': 'gini', 'max_depth': 14,
```



```
'max_features': 'auto', 'n_estimators': 50}
Best Mean Cross Validation Score is 0.8327935241589982
Train score is 0.8608990670059373
```

3.1.8 Gradient Boost - .79684

```
[43]: from sklearn.ensemble import GradientBoostingClassifier
```

```
[44]: gbc= GradientBoostingClassifier(random_state=42)
gbc_param = {
    'max_depth' : range(2,15,5),
    'n_estimators' : range(50,80,10),
    'learning_rate' : [0.01,1.0,2],
}
gbc_grid = GridSearchCV(gbc, gbc_param,cv=5, return_train_score=True,scoring =_
    →f2score)
gbc_grid.fit(X_train,y_train)
```

```
[44]: GridSearchCV(cv=5, error_score=nan,
                estimator=GradientBoostingClassifier(ccp_alpha=0.0,
                criterion='friedman_mse',
                init=None, learning_rate=0.1,
                loss='deviance', max_depth=3,
                max_features=None,
                max_leaf_nodes=None,
                min_impurity_decrease=0.0,
                min_impurity_split=None,
                min_samples_leaf=1,
                min_samples_split=2,
                min_weight_fraction_leaf=0.0,
                n_estimators=100,
                n_iter_no_c...
                presort='deprecated',
                random_state=42,
                subsample=1.0, tol=0.0001,
                validation_fraction=0.1,
                verbose=0, warm_start=False),
                iid='deprecated', n_jobs=None,
                param_grid={'learning_rate': [0.01, 1.0, 2],
                'max_depth': range(2, 15, 5),
                'n_estimators': range(50, 80, 10)},
                pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[340]: print(f'Best Mean Cross Validation Score is {gbc_grid.best_score_}')
print(f'Best Mean Cross Validation Score is {gbc_grid.best_params_}')
```

```
print(f'Train score is {gbc_grid.score(X_train,y_train)}')
```

Best Mean Cross Validation Score is 0.7998976537815454

Best Mean Cross Validation Score is {'learning_rate': 2, 'max_depth': 12,
'n_estimators': 50}

Train score is 1.0

3.1.9 XGBoost - .82116

```
[46]: from xgboost import XGBClassifier
```

```
[47]: from xgboost import XGBClassifier
xgbc= XGBClassifier(random_state=42,early_stopping_rounds=2,objective= 'binary:
↳logistic')
xgbc_param = {
    'max_depth' : range(2,15,2),
    'n_estimators' : range(50,80,10),
    'learning_rate' : [0.1,0.01],
    'min_child_weight' : [1,3,5,7],
    'subsample':[0.6,0.7,0.8,0.9,1]
}
xgbc_grid = GridSearchCV(xgbc, xgbc_param,cv=5, return_train_score=True,scoring_
↳= f2score)
xgbc_grid.fit(X_train,y_train)
```

```
[47]: GridSearchCV(cv=5, error_score=nan,
    estimator=XGBClassifier(base_score=None, booster=None,
        colsample_bylevel=None,
        colsample_bynode=None,
        colsample_bytree=None,
        early_stopping_rounds=2, gamma=None,
        gpu_id=None, importance_type='gain',
        interaction_constraints=None,
        learning_rate=None, max_delta_step=None,
        max_depth=None, min_child_weight=None,
        missing=nan, monotone_...
        subsample=None, tree_method=None,
        validate_parameters=False,
        verbosity=None),
    iid='deprecated', n_jobs=None,
    param_grid={'learning_rate': [0.1, 0.01],
        'max_depth': range(2, 15, 2),
        'min_child_weight': [1, 3, 5, 7],
        'n_estimators': range(50, 80, 10),
        'subsample': [0.6, 0.7, 0.8, 0.9, 1]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
    scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[341]: print(f'Best Mean Cross Validation Score is {xgbc_grid.best_score_}')
print(f'Best Mean Cross Validation Score is {xgbc_grid.best_params_}')
print(f'Train score is {xgbc_grid.score(X_train,y_train)}')
```

Best Mean Cross Validation Score is 0.8536961171789608

Best Mean Cross Validation Score is {'learning_rate': 0.1, 'max_depth': 4, 'min_child_weight': 3, 'n_estimators': 50, 'subsample': 0.7}

Train score is 0.867003367003367

3.2 Cost Sensitive Models

3.2.1 Logistic Regression - .80909

```
[49]: logreg = LogisticRegression(random_state=42)
param_grid = {'class_weight': ['balanced', 'balanced_subsample'],
              'C': [0.001, 0.01, 0.1, 1, 10, 1000],
              'solver': ['lbfgs'],
              'penalty': ['l2'],
              'max_iter': range(1000, 1500, 10)}# [{0:100,1:1}, {0:10,1:1}, {0:1,1:
→1}, {0:1,1:10}, {0:1,1:100}]

#apply grid search
grid_CSlogreg= GridSearchCV(logreg, param_grid, cv=5, n_jobs=-1,
→scoring=f2score)
grid_CSlogreg.fit(X_train,y_train)
```

```
[49]: GridSearchCV(cv=5, error_score=nan,
                  estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                  fit_intercept=True,
                  intercept_scaling=1, l1_ratio=None,
                  max_iter=100, multi_class='auto',
                  n_jobs=None, penalty='l2',
                  random_state=42, solver='lbfgs',
                  tol=0.0001, verbose=0,
                  warm_start=False),
                  iid='deprecated', n_jobs=-1,
                  param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 1000],
                  'class_weight': ['balanced', 'balanced_subsample'],
                  'max_iter': range(1000, 1500, 10), 'penalty': ['l2'],
                  'solver': ['lbfgs']},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[105]: print("Best parameters: {}".format(grid_CSlogreg.best_params_))
```

```
print("Best Mean cross-validation score: {:.2f}".format(grid_CSlogreg.
    ↳best_score_))
print('Train score: {:.4f}'.format(grid_logit.score(X_train, y_train)))
```

Best parameters: {'C': 10, 'class_weight': 'balanced_subsample', 'max_iter': 1000, 'penalty': 'l2', 'solver': 'lbfgs'}
 Best Mean cross-validation score: 0.80
 Train score: 0.8177

3.2.2 Descision Tree - .74823

```
[51]: dtree = DecisionTreeClassifier()
param_grid = {'class_weight': ['balanced', 'balanced_subsample']}#[{0:100,1:1},
    ↳{0:10,1:1}, {0:1,1:1}, {0:1,1:10}, {0:1,1:100}]]
#apply grid search
grid_CSdtree= GridSearchCV(dtree, param_grid, cv=5, n_jobs=2, scoring=f2score)
grid_CSdtree.fit(X_train,y_train)
```

```
[51]: GridSearchCV(cv=5, error_score=nan,
            estimator=DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None,
            criterion='gini', max_depth=None,
            max_features=None,
            max_leaf_nodes=None,
            min_impurity_decrease=0.0,
            min_impurity_split=None,
            min_samples_leaf=1,
            min_samples_split=2,
            min_weight_fraction_leaf=0.0,
            presort='deprecated',
            random_state=None,
            splitter='best'),
            iid='deprecated', n_jobs=2,
            param_grid={'class_weight': ['balanced', 'balanced_subsample']},
            pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
            scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[342]: print("Best parameters: {}".format(grid_CSdtree.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(grid_CSdtree.
    ↳best_score_))
print(f'Train score is {grid_CSdtree.score(X_train,y_train)}')
```

Best parameters: {'class_weight': 'balanced'}
 Best Mean cross-validation score: 0.78
 Train score is 1.0

3.2.3 SVM

```
[53]: # svc = SVC()
# param_grid = {'class_weight': ['balanced', 'balanced_subsample']}#{0:100,1:
↪1}, {0:10,1:1}, {0:1,1:1}, {0:1,1:10}, {0:1,1:100}}

# #apply grid search
# grid_CSsvc= GridSearchCV(svc, param_grid, cv=5, n_jobs=2, scoring=f2score)
# grid_CSsvc.fit(X_train,y_train)
```

Results

```
[54]: # print("Best parameters: {}".format(grid_CSsvc.best_params_))
# print("Best Mean cross-validation score: {:.2f}".format(grid_CSsvc.
↪best_score_))
```

3.2.4 Random Forest - .80586

```
[55]: from numpy import mean
from sklearn.datasets import make_classification
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.ensemble import BaggingClassifier
```

```
[58]: # define model
SCRandom = RandomForestClassifier(random_state=1)
SCRandom_param = {'n_estimators':range(2,10,2),
                  'class_weight':['balanced', 'balanced_subsample']}

# define evaluation procedure
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
#apply grid search
grid_CSRF= GridSearchCV(SCRandom, SCRandom_param, cv=cv, n_jobs=-1,
↪scoring=f2score)
grid_CSRF.fit(X_train,y_train)
```

```
[58]: GridSearchCV(cv=RepeatedStratifiedKFold(n_repeats=3, n_splits=10,
random_state=1),
              error_score=nan,
              estimator=RandomForestClassifier(bootstrap=True, ccp_alpha=0.0,
class_weight=None,
criterion='gini', max_depth=None,
max_features='auto',
max_leaf_nodes=None,
max_samples=None,
min_impurity_decrease=0.0,
min_impurity_split=None,
min_samples_leaf=1,
```

```

min_samples_split=2,
min_weight_fraction_leaf=0.0,
n_estimators=100, n_jobs=None,
oob_score=False, random_state=1,
verbose=0, warm_start=False),

iid='deprecated', n_jobs=-1,
param_grid={'class_weight': ['balanced', 'balanced_subsample'],
            'n_estimators': range(2, 10, 2)},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=make_scorer(fbeta_score, beta=2), verbose=0)

```

Results

```

[343]: print("Best parameters: {}".format(grid_CSRF.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(grid_CSRF.best_score_))
print(f'Train score is {grid_CSRF.score(X_train,y_train)}')

```

Best parameters: {'class_weight': 'balanced', 'n_estimators': 8}
Best Mean cross-validation score: 0.79
Train score is 0.9727947238252268

3.2.5 XGBoost - .83182

```

[400]: # SCxgboost = XGBClassifier()
# param_SCgrid = {'scale_pos_weight': [1, 10, 25, 50, 75, 99, 100, 1000]}

# #apply grid search
# grid_xgboost= GridSearchCV(SCxgboost, param_SCgrid, cv=5, n_jobs=-1,
#                             ↪scoring=f2score)
# grid_xgboost.fit(X_train,y_train)

SCxgboost = XGBClassifier()
param_SCgrid = {'scale_pos_weight': range(1000,1200,50),
                'class_weight':['balanced','balanced_subsample']}

#apply grid search
grid_xgboost= GridSearchCV(SCxgboost, param_SCgrid, cv=5, n_jobs=-1,
                            ↪scoring=f2score)
grid_xgboost.fit(X_train,y_train)

# SCxgboost = XGBClassifier(early_stopping_rounds=5)
# param_SCgrid = {'scale_pos_weight': range(1000,1200,50),
#                 'class_weight':['balanced','balanced_subsample'],
#                 'n_estimators':range(1000,1200,20),
#                 'learning_rate':[0.05,0.01,0.1]}

# #apply grid search

```

```
# grid_xgboost= GridSearchCV(SCxgboost, param_SCgrid, cv=5, n_jobs=-1,
↪scoring=f2score)
# grid_xgboost.fit(X_train,y_train)
```

```
[400]: GridSearchCV(cv=5, error_score=nan,
                  estimator=XGBClassifier(base_score=None, booster=None,
                                          colsample_bylevel=None,
                                          colsample_bynode=None,
                                          colsample_bytreet=None, gamma=None,
                                          gpu_id=None, importance_type='gain',
                                          interaction_constraints=None,
                                          learning_rate=None, max_delta_step=None,
                                          max_depth=None, min_child_weight=None,
                                          missing=nan, monotone_constraints=None,
                                          n_estim...
                                          reg_lambda=None, scale_pos_weight=None,
                                          subsample=None, tree_method=None,
                                          validate_parameters=False,
                                          verbosity=None),
                  iid='deprecated', n_jobs=-1,
                  param_grid={'class_weight': ['balanced', 'balanced_subsample'],
                              'scale_pos_weight': range(1000, 1200, 50)},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[401]: print("Best parameters: {}".format(grid_xgboost.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(grid_xgboost.
↪best_score_))
print(f'Train score is {grid_xgboost.score(X_train,y_train)}')
```

```
Best parameters: {'class_weight': 'balanced', 'scale_pos_weight': 1000}
Best Mean cross-validation score: 0.86
Train score is 1.0
```

3.2.6 Extra Trees - .80586

```
[62]: #extratrees
SCEExtraTree = ExtraTreesClassifier(random_state=42)
SCEExtraTree_param = {'n_estimators':range(2,50,5),
                      'class_weight':['balanced','balanced_subsample']}

cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=42)

grid_SCExTree= GridSearchCV(SCEExtraTree, SCEExtraTree_param, cv=cv, n_jobs=-1,
    ↳scoring=f2score)
grid_SCExTree.fit(X_train,y_train)
```

```
[62]: GridSearchCV(cv=RepeatedStratifiedKFold(n_repeats=3, n_splits=10,
random_state=42),
                error_score=nan,
                estimator=ExtraTreesClassifier(bootstrap=False, ccp_alpha=0.0,
                                                class_weight=None, criterion='gini',
                                                max_depth=None, max_features='auto',
                                                max_leaf_nodes=None,
                                                max_samples=None,
                                                min_impurity_decrease=0.0,
                                                min_impurity_split=None,
                                                min_samples_leaf=1,
                                                min_samples_split=2,
                                                min_weight_fraction_leaf=0.0,
                                                n_estimators=100, n_jobs=None,
                                                oob_score=False, random_state=42,
                                                verbose=0, warm_start=False),
                iid='deprecated', n_jobs=-1,
                param_grid={'class_weight': ['balanced', 'balanced_subsample'],
                            'n_estimators': range(2, 50, 5)},
                pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[346]: print("Best parameters: {}".format(grid_SCExTree.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(grid_SCExTree.
    ↳best_score_))
print(f'Train score is {grid_SCExTree.score(X_train,y_train)}')
```

```
Best parameters: {'class_weight': 'balanced', 'n_estimators': 27}
Best Mean cross-validation score: 0.84
Train score is 1.0
```


3.2.7 Bagging Decision Tree - Under Sampling

```
[66]: from imblearn.sampling import RandomUnderSampler
from sklearn.ensemble import BaggingClassifier
from imblearn.pipeline import Pipeline
from imblearn.ensemble import BalancedBaggingClassifier

[67]: bag_dtrees1 = Pipeline(steps=[('random', RandomUnderSampler(random_state=42)),
↳ ('bagging', BalancedBaggingClassifier(base_estimator=
↳ DecisionTreeClassifier(class_weight='balanced'),
↳ random_state=42))])

bag_dtrees1_param = {'bagging__n_estimators': [100]}

cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=42)

grid_SCBagging = GridSearchCV(bag_dtrees1, bag_dtrees1_param, cv=cv,
↳ scoring=f2score, return_train_score=True)
grid_SCBagging.fit(X_train, y_train)

[67]: GridSearchCV(cv=RepeatedStratifiedKFold(n_repeats=3, n_splits=10,
random_state=42),
error_score=nan,
estimator=Pipeline(memory=None,
steps=[('random',
RandomUnderSampler(random_state=42,
replacement=False,
sampling_strategy='auto')),
('bagging',
BalancedBaggingClassifier(base_estimator=DecisionTreeClassifier(ccp_alpha=0.0,
class_weight='balanced',
criterion='g...
max_features=1.0,
max_samples=1.0,
n_estimators=10,
n_jobs=None,
oob_score=False,
random_state=42,
replacement=False,
sampling_strategy='auto',
warm_start=False))],
verbose=0,
verbose=False),
iid='deprecated', n_jobs=None,
```

```
param_grid={'bagging__n_estimators': [100]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[347]: print("Best parameters: {}".format(grid_ScBagging.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(grid_ScBagging.
↪best_score_))
print(f'Train score is {grid_ScBagging.score(X_train,y_train)}')
```

Best parameters: {'bagging__n_estimators': 100}

Best Mean cross-validation score: 0.47

Train score is 0.3774752475247524

3.3 Data Sampling

3.3.1 Over Sampling

```
[69]: from imblearn.over_sampling import RandomOverSampler
from imblearn.over_sampling import SMOTE
from imblearn.over_sampling import SVMSMOTE
from imblearn.over_sampling import ADASYN
from imblearn.pipeline import Pipeline
```

```
[70]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.svm import LinearSVC
from sklearn.ensemble import RandomForestClassifier

from numpy import mean
from numpy import std
from sklearn.datasets import make_classification
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.dummy import DummyClassifier
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import auc
from sklearn.metrics import make_scorer
from matplotlib import pyplot
from sklearn.model_selection import GridSearchCV
```

Logistic Regression

```
[71]: # GridSearch with oversampling
```

```

pipe_logit_smote =
    Pipeline([('smote', SMOTE()), ('Logistic', LogisticRegression(random_state=42))])

param_grid = {
    # try different feature engineering parameters
    'smote__k_neighbors': [1,2,3],
    'Logistic__C': [0.001, 0.01],
    'Logistic__solver': ['lbfgs'],
    'Logistic__penalty': ['l2'],
    'Logistic__max_iter': [1200] #range(1000,1500,10)
}

#apply grid search
grid_logitsmote= GridSearchCV(pipe_logit_smote, param_grid, cv=5, n_jobs=2,
    scoring=f2score)
grid_logitsmote.fit(X_train,y_train)

```

```

[71]: GridSearchCV(cv=5, error_score=nan,
                estimator=Pipeline(memory=None,
                                   steps=[('smote',
                                           SMOTE(k_neighbors=5, n_jobs=None,
                                                  random_state=None,
                                                  sampling_strategy='auto')),
                                           ('Logistic',
                                            LogisticRegression(C=1.0,
                                                                class_weight=None,
                                                                dual=False,
                                                                fit_intercept=True,
                                                                intercept_scaling=1,
                                                                l1_ratio=None,
                                                                max_iter=100,
                                                                multi_class='auto',
                                                                n_jobs=None,
                                                                penalty='l2',
                                                                random_state=42,
                                                                solver='lbfgs',
                                                                tol=0.0001,
                                                                verbose=0,
                                                                warm_start=False))],
                                   verbose=False),
                iid='deprecated', n_jobs=2,
                param_grid={'Logistic__C': [0.001, 0.01],
                             'Logistic__max_iter': [1200],
                             'Logistic__penalty': ['l2'],
                             'Logistic__solver': ['lbfgs'],
                             'smote__k_neighbors': [1, 2, 3]},
                pre_dispatch='2*n_jobs', refit=True, return_train_score=False,

```

```
scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[348]: print("Best parameters: {}".format(grid_logitstmote.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(grid_logitstmote.
    ↳best_score_))
print(f'Train score is {grid_logitstmote.score(X_train,y_train)}')
```

```
Best parameters: {'Logistic__C': 0.001, 'Logistic__max_iter': 1200,
'Logistic__penalty': 'l2', 'Logistic__solver': 'lbfgs', 'smote__k_neighbors': 2}
Best Mean cross-validation score: 0.70
Train score is 0.7057340894770006
```

Decision Tree

```
[73]: #decisiontreee
pipe_ds_dtree = ␣
    ↳Pipeline([('smote',SMOTE()),('dtree',DecisionTreeClassifier(random_state=42))])
param_ds_dtree = {'smote__k_neighbors': range(2,8,2),
                  'dtree__max_depth': range(1,20,5)}

DOS_dtree= GridSearchCV(pipe_ds_dtree,param_ds_dtree, cv=5, n_jobs=2,␣
    ↳scoring=f2score)
DOS_dtree.fit(X_train, y_train)
```

```
[73]: GridSearchCV(cv=5, error_score=nan,
                  estimator=Pipeline(memory=None,
                                     steps=[('smote',
                                             SMOTE(k_neighbors=5, n_jobs=None,
                                                    random_state=None,
                                                    sampling_strategy='auto')),
                                             ('dtree',
                                              DecisionTreeClassifier(ccp_alpha=0.0,
                                                                      class_weight=None,
                                                                      criterion='gini',
                                                                      max_depth=None,
                                                                      max_features=None,
                                                                      max_leaf_nodes=None,
                                                                      min_impurity_decrease=0.0,
                                                                      min_impurity_split=None,
                                                                      min_samples_leaf=1,
                                                                      min_samples_split=2,
                                                                      min_weight_fraction_leaf=0.0,
                                                                      presort='deprecated',
                                                                      random_state=42,
                                                                      splitter='best'))],
                                     verbose=False),
```

```
iid='deprecated', n_jobs=2,
param_grid={'dtree__max_depth': range(1, 20, 5),
            'smote__k_neighbors': range(2, 8, 2)},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[349]: print("Best parameters: {}".format(DOS_dtree.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(DOS_dtree.best_score_))
print(f'Train score is {DOS_dtree.score(X_train,y_train)}')
```

```
Best parameters: {'dtree__max_depth': 16, 'smote__k_neighbors': 4}
Best Mean cross-validation score: 0.74
Train score is 0.959874114870181
```

KNN - .80279

```
[406]: #knn
DOS_knn = Pipeline([('smote', SMOTE()), ('knn', KNeighborsClassifier())])
param_DOS_knn = {'smote__k_neighbors': range(2,8,2),
                 'knn__n_neighbors': range(1,10,2)}

DOSgrid_knn= GridSearchCV(DOS_knn,param_DOS_knn, cv=5, n_jobs=-1,
↳scoring=f2score)
DOSgrid_knn.fit(X_train, y_train)
```

```
[406]: GridSearchCV(cv=5, error_score=nan,
                    estimator=Pipeline(memory=None,
                                       steps=[('smote',
                                                SMOTE(k_neighbors=5, n_jobs=None,
                                                       random_state=None,
                                                       sampling_strategy='auto')),
                                               ('knn',
                                                KNeighborsClassifier(algorithm='auto',
                                                                       leaf_size=30,
                                                                       metric='minkowski',
                                                                       metric_params=None,
                                                                       n_jobs=None,
                                                                       n_neighbors=5, p=2,
                                                                       weights='uniform'))]),
                    verbose=False),
                    iid='deprecated', n_jobs=-1,
                    param_grid={'knn__n_neighbors': range(1, 10, 2),
                                'smote__k_neighbors': range(2, 8, 2)},
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                    scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[408]: print("Best parameters: {}".format(DOSgrid_knn.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(DOSgrid_knn.
    ↳best_score_))
print(f'Train score is {DOSgrid_knn.score(X_train,y_train)}')
```

Best parameters: {'knn__n_neighbors': 1, 'smote__k_neighbors': 2}

Best Mean cross-validation score: 0.82

Train score is 1.0

```
[411]: test_data_labels = DOSgrid_knn.predict(X_test)

# Create predictions to be submitted!
pd.DataFrame({'Id': test.Id, 'Target': test_data_labels}).
    ↳to_csv('Project2_solution_base.csv', index=False)
print("Done :D")
```

Done :D

SVM

```
[77]: # #sum
# DOS_sum =
    ↳Pipeline([('smote', SMOTE()), ('sum', SVC(random_state=42, probability=True))])
# param_DOS_sum = {'smote__k_neighbors': [1,2,3,4,5],
#                  'sum__C': [1,20,30],
#                  'sum__kernel': ['rbf', 'sigmoid'],
#                  'sum__gamma': ['auto']}

# DOSgrid_sum= GridSearchCV(DOS_sum,param_DOS_sum, cv=5, n_jobs=-1,
    ↳scoring=f2score)
# DOSgrid_sum.fit(X_train, y_train)

[78]: # print("Best parameters: {}".format(DOSgrid_sum.best_params_))
# print("Best Mean cross-validation score: {:.2f}".format(DOSgrid_sum.
    ↳best_score_))
```

Random Forest - .81395

```
[79]: #randomforest
DOS_rf =
    ↳Pipeline([('smote', SMOTE()), ('rf', RandomForestClassifier(random_state=42))])
param_DOS_rf = {'smote__k_neighbors': [1,2,3,4,5],
                'rf__n_estimators': range(20,80,10),
                'rf__max_features': ['auto', 'sqrt', 'log2'],
                'rf__max_depth': range(2,8,4),
                'rf__criterion': ['gini', 'entropy']}
```

```
DOSgrid_rf= GridSearchCV(DOS_rf,param_DOS_rf, cv=5, n_jobs=-1, scoring=f2score)
DOSgrid_rf.fit(X_train, y_train)
```

```
[79]: GridSearchCV(cv=5, error_score=nan,
                estimator=Pipeline(memory=None,
                                   steps=[('smote',
                                           SMOTE(k_neighbors=5, n_jobs=None,
                                                  random_state=None,
                                                  sampling_strategy='auto')),
                                           ('rf',
                                            RandomForestClassifier(bootstrap=True,
                                                                    ccp_alpha=0.0,
                                                                    class_weight=None,
                                                                    criterion='gini',
                                                                    max_depth=None,
                                                                    max_features='auto',
                                                                    max_leaf_nodes=None,
                                                                    max_samples=None,
                                                                    min_impurity_decrease=0.0,
                                                                    min_impurity_split=None,
                                                                    m...
                                                                    warm_start=False))]),
                                   verbose=False),
                iid='deprecated', n_jobs=-1,
                param_grid={'rf__criterion': ['gini', 'entropy'],
                            'rf__max_depth': range(2, 8, 4),
                            'rf__max_features': ['auto', 'sqrt', 'log2'],
                            'rf__n_estimators': range(20, 80, 10),
                            'smote__k_neighbors': [1, 2, 3, 4, 5]},
                pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[350]: print("Best parameters: {}".format(DOSgrid_rf.best_params_))
        print("Best Mean cross-validation score: {:.2f}".format(DOSgrid_rf.best_score_))
        print(f'Train score is {DOSgrid_rf.score(X_train,y_train)}')
```

```
Best parameters: {'rf__criterion': 'gini', 'rf__max_depth': 6,
'rf__max_features': 'auto', 'rf__n_estimators': 60, 'smote__k_neighbors': 1}
Best Mean cross-validation score: 0.86
Train score is 0.9226430298146655
```

Easy Ensembler

```
[81]: #easy ensemble classifier
from imblearn.ensemble import EasyEnsembleClassifier
DOS_easy = Pipeline([('smote',SMOTE()),('easy',EasyEnsembleClassifier())])
param_DOS_easy = {'smote__k_neighbors': range(2,10,2)}
```

```
DOSgrid_easy= GridSearchCV(DOS_easy,param_DOS_easy, cv=5, n_jobs=-1,
    ↳scoring=f2score)
DOSgrid_easy.fit(X_train, y_train)
```

```
[81]: GridSearchCV(cv=5, error_score=nan,
                estimator=Pipeline(memory=None,
                                   steps=[('smote',
                                           SMOTE(k_neighbors=5, n_jobs=None,
                                                  random_state=None,
                                                  sampling_strategy='auto'))],
                                   ('easy',
                                   EasyEnsembleClassifier(base_estimator=None,
                                                         n_estimators=10,
                                                         n_jobs=None,
                                                         random_state=None,
                                                         replacement=False,
                                                         sampling_strategy='auto',
                                                         verbose=0,
                                                         warm_start=False))],
                                   verbose=False),
                iid='deprecated', n_jobs=-1,
                param_grid={'smote__k_neighbors': range(2, 10, 2)},
                pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[351]: print("Best parameters: {}".format(DOSgrid_easy.best_params_))
        print("Best Mean cross-validation score: {:.2f}".format(DOSgrid_easy.
        ↳best_score_))
        print(f'Train score is {DOSgrid_easy.score(X_train,y_train)}')
```

Best parameters: {'smote__k_neighbors': 8}

Best Mean cross-validation score: 0.68

Train score is 0.6595622119815668

XGBoost - .79216

```
[83]: from xgboost import XGBClassifier
        #xgb
        DOS_xgb = Pipeline([('smote', SMOTE()), ('xgboost', XGBClassifier(random_state=42,
        ↳early_stopping_rounds=2,
        #
        ↳n_estimators=100,
        ↳objective='binary:logistic',
```



```
#
↳max_depth = 4,

↳)))]
param_DOS_xgb = {'smote__k_neighbors': [1,2,3,4,5],
                  'xgboost__max_depth' : range(2,8,4),
                  'xgboost__n_estimators' : range(50,70,20),
                  'xgboost__learning_rate' : [0.1],
                  'xgboost__min_child_weight' : [1,3,5,7],
                  'xgboost__subsample': [0.6,0.7,0.8,0.9,1]}

DOSgrid_xgb= GridSearchCV(DOS_xgb,param_DOS_xgb, cv=3, n_jobs=-1,↳
↳scoring=f2score)
DOSgrid_xgb.fit(X_train, y_train)
```

```
[83]: GridSearchCV(cv=3, error_score=nan,
                  estimator=Pipeline(memory=None,
                                     steps=[('smote',
                                             SMOTE(k_neighbors=5, n_jobs=None,
                                                    random_state=None,
                                                    sampling_strategy='auto')),
                                             ('xgboost',
                                              XGBClassifier(base_score=None,
                                                            booster=None,
                                                            colsample_bylevel=None,
                                                            colsample_bynode=None,
                                                            colsample_bytree=None,
                                                            early_stopping_rounds=2,
                                                            gamma=None, gpu_id=None,
                                                            importance_type='gain',
                                                            i...

                  iid='deprecated', n_jobs=-1,
                  param_grid={'smote__k_neighbors': [1, 2, 3, 4, 5],
                              'xgboost__learning_rate': [0.1],
                              'xgboost__max_depth': range(2, 8, 4),
                              'xgboost__min_child_weight': [1, 3, 5, 7],
                              'xgboost__n_estimators': range(50, 70, 20),
                              'xgboost__subsample': [0.6, 0.7, 0.8, 0.9, 1]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[352]: print("Best parameters: {}".format(DOSgrid_xgb.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(DOSgrid_xgb.
↳best_score_))
print(f'Train score is {DOSgrid_xgb.score(X_train,y_train)}')
```

```
Best parameters: {'smote__k_neighbors': 1, 'xgboost__learning_rate': 0.1,
'xgboost__max_depth': 6, 'xgboost__min_child_weight': 5,
'xgboost__n_estimators': 50, 'xgboost__subsample': 0.8}
Best Mean cross-validation score: 0.84
Train score is 0.9345794392523363
```

Neural Network - MLP Classifier

```
[152]: from sklearn.neural_network import MLPClassifier
```

```
[159]: DOS_mlp = Pipeline([('smote', SMOTE()),
                           ('mlp', MLPClassifier(random_state=42))])

parameters = {'mlp__solver': ['lbfgs'],
              'mlp__max_iter': [1200],
              'mlp__alpha': 10.0 ** -np.arange(1, 10),
              'mlp__hidden_layer_sizes': np.arange(10, 15),
              'mlp__random_state': range(2, 12, 2)}
DOSMLPGrid = GridSearchCV(DOS_mlp, parameters, cv=5, n_jobs=-1, scoring=f2score)

DOSMLPGrid.fit(X_train, y_train)
```

```
[159]: GridSearchCV(cv=5, error_score=nan,
                  estimator=Pipeline(memory=None,
                                     steps=[('random',
                                             RandomUnderSampler(random_state=42,
                                                                    replacement=False,
                                                                    sampling_strategy='auto')),
                                             ('mlp',
                                              MLPClassifier(activation='relu',
                                                            alpha=0.0001,
                                                            batch_size='auto',
                                                            beta_1=0.9, beta_2=0.999,
                                                            early_stopping=False,
                                                            epsilon=1e-08,
                                                            hidden_layer_sizes=(100,),
                                                            learning_rate='constant',
                                                            learning_rate_decay=0.0001,
                                                            max_iter=1000,
                                                            n_iter_no_change=10,
                                                            n_layers=2,
                                                            pre_activation_function='tanh',
                                                            random_state=None,
                                                            solver='lbfgs',
                                                            tol=0.0001,
                                                            verbose=0,
                                                            warm_start=False,
                                                            zeta=0.0001)),
                                     iid='deprecated', n_jobs=-1,
                                     param_grid={'mlp__alpha': array([1.e-01, 1.e-02, 1.e-03, 1.e-04,
1.e-05, 1.e-06, 1.e-07, 1.e-08,
1.e-09]),
                                                'mlp__hidden_layer_sizes': array([10, 11, 12, 13, 14]),
                                                'mlp__max_iter': [1200],
                                                'mlp__random_state': range(2, 12, 2),
                                                'mlp__solver': ['lbfgs']}},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

```
[353]: print("Best parameters: {}".format(DOSMLPGrid.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(DOSMLPGrid.best_score_))
print(f'Train score is {DOSMLPGrid.score(X_train,y_train)}')
```

```
Best parameters: {'mlp__alpha': 0.1, 'mlp__hidden_layer_sizes': 10,
'mlp__max_iter': 1200, 'mlp__random_state': 2, 'mlp__solver': 'lbfgs'}
Best Mean cross-validation score: 0.43
Train score is 0.4450930317402408
```

3.3.2 Under Sampling

Logistic Regression

```
[85]: from imblearn.under_sampling import RandomUnderSampler
```

```
[86]: # GridSearch with oversampling
pipe_logit_RS = Pipeline([('random',RandomUnderSampler(random_state=42)),
                           ('Logistic',LogisticRegression(random_state=42))])

param_grid = {
    # try different feature engineering parameters
    'Logistic__C': [0.001, 0.01, 0.1, 1, 10,1000],
    'Logistic__solver':['lbfgs'],
    'Logistic__penalty':['l2'],
    'Logistic__max_iter':[1200]#range(1000,1500,10)
}

#apply grid search
grid_logitRS= GridSearchCV(pipe_logit_smote, param_grid, cv=5, n_jobs=-1,
    ↳scoring=f2score)
grid_logitRS.fit(X_train,y_train)
```

```
[86]: GridSearchCV(cv=5, error_score=nan,
                  estimator=Pipeline(memory=None,
                                     steps=[('smote',
                                              SMOTE(k_neighbors=5, n_jobs=None,
                                                    random_state=None,
                                                    sampling_strategy='auto'))],
                                     ('Logistic',
                                      LogisticRegression(C=1.0,
                                                         class_weight=None,
                                                         dual=False,
                                                         fit_intercept=True,
                                                         intercept_scaling=1,
                                                         l1_ratio=None,
                                                         max_iter=100,
                                                         multi_class='auto',
                                                         n_jobs=None,
```

```

penalty='l2',
random_state=42,
solver='lbfgs',
tol=0.0001,
verbose=0,
warm_start=False))],
verbose=False),
iid='deprecated', n_jobs=-1,
param_grid={'Logistic_C': [0.001, 0.01, 0.1, 1, 10, 1000],
            'Logistic__max_iter': [1200],
            'Logistic__penalty': ['l2'],
            'Logistic__solver': ['lbfgs']},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=make_scorer(fbeta_score, beta=2), verbose=0)

```

Results

```

[354]: print("Best parameters: {}".format(grid_logitRS.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(grid_logitRS.
    ↳best_score_))
print(f'Train score is {grid_logitRS.score(X_train,y_train)}')

```

```

Best parameters: {'Logistic_C': 0.001, 'Logistic__max_iter': 1200,
'Logistic__penalty': 'l2', 'Logistic__solver': 'lbfgs'}
Best Mean cross-validation score: 0.69
Train score is 0.7070707070707071

```

Decision Tree

```

[88]: #decisiontree
pipe_dus_dtree = Pipeline([('random',RandomUnderSampler(random_state=42)),
                           ('dtree',DecisionTreeClassifier(random_state=42))])
param_dus_dtree = {'dtree__max_depth': range(1,20)}

DUS_dtree= GridSearchCV(pipe_dus_dtree,param_dus_dtree, cv=5, n_jobs=-1,
    ↳scoring=f2score)
DUS_dtree.fit(X_train, y_train)

```

```

[88]: GridSearchCV(cv=5, error_score=nan,
                estimator=Pipeline(memory=None,
                                steps=[('random',
                                        RandomUnderSampler(random_state=42,
                                                                replacement=False,
                                                                sampling_strategy='auto')),
                                        ('dtree',
                                         DecisionTreeClassifier(ccp_alpha=0.0,
                                                                class_weight=None,
                                                                criterion='gini',

```

```

max_depth=None,
max_features=None,
max_leaf_nodes=None,
min_impurity_decrease=0.0,
min_impurity_split=None,
min_samples_leaf=1,
min_samples_split=2,
min_weight_fraction_leaf=0.0,
presort='deprecated',
splitter='best'))],
verbose=False),
iid='deprecated', n_jobs=-1,
param_grid={'dtree__max_depth': range(1, 20)},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=make_scorer(fbeta_score, beta=2), verbose=0)

```

Results

```

[355]: print("Best parameters: {}".format(DUS_dtree.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(DUS_dtree.best_score_))
print(f'Train score is {DUS_dtree.score(X_train,y_train)}')

```

```

Best parameters: {'dtree__max_depth': 1}
Best Mean cross-validation score: 0.51
Train score is 0.5549654806160383

```

KNN

```

[404]: #knn
DUS_knn = Pipeline([('random',RandomUnderSampler(random_state=42)),
                    ('knn',KNeighborsClassifier())])
param_DUS_knn = {'knn__n_neighbors': range(1,10)}

DUSgrid_knn= GridSearchCV(DUS_knn,param_DUS_knn, cv=5, n_jobs=-1,
    ↳scoring=f2score)
DUSgrid_knn.fit(X_train, y_train)

```

```

[404]: GridSearchCV(cv=5, error_score=nan,
                  estimator=Pipeline(memory=None,
                                     steps=[('random',
                                             RandomUnderSampler(random_state=42,
                                                                     replacement=False,
                                                                     sampling_strategy='auto')),
                                             ('knn',
                                              KNeighborsClassifier(algorithm='auto',
                                                                      leaf_size=30,
                                                                      metric='minkowski',

```

```

weights='uniform'))],
metric_params=None,
n_jobs=None,
n_neighbors=5, p=2,
verbose=False),
iid='deprecated', n_jobs=-1,
param_grid={'knn__n_neighbors': range(1, 10)},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=make_scorer(fbeta_score, beta=2), verbose=0)

```

Results

```

[405]: print("Best parameters: {}".format(DUSgrid_knn.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(DUSgrid_knn.
↳best_score_))

```

Best parameters: {'knn__n_neighbors': 8}
Best Mean cross-validation score: 0.72

SVM

```

[92]: # #svm
# DUS_sum = Pipeline([('random',RandomUnderSampler(random_state=42)),
# ('sum',SVC(random_state=42))])
# param_DUS_sum = {'sum__C':[1,20,30],
# 'sum__kernel':['rbf','sigmoid'],
# 'sum__gamma':['auto']}

# DUSgrid_sum= GridSearchCV(DUS_sum,param_DUS_sum, cv=5, n_jobs=-1,
↳scoring=f2score)
# DUSgrid_sum.fit(X_train, y_train)

[93]: # print("Best parameters: {}".format(DUSgrid_sum.best_params_))
# print("Best Mean cross-validation score: {:.2f}".format(DUSgrid_sum.
↳best_score_))

```

Random Forest - .76518

```

[143]: #randomforest
DUS_rf = Pipeline([('random',RandomUnderSampler(random_state=42)),
('rf',RandomForestClassifier(random_state=42))])
param_DUS_rf = {'rf__n_estimators': range(20,100,10),
'rf__max_features': ['auto', 'sqrt', 'log2'],
'rf__max_depth': range(2,15,4),
'rf__criterion': ['gini', 'entropy']}

DUSgrid_rf= GridSearchCV(DUS_rf,param_DUS_rf, cv=5, n_jobs=-1, scoring=f2score)
DUSgrid_rf.fit(X_train, y_train)

```

```
[143]: GridSearchCV(cv=5, error_score=nan,
                  estimator=Pipeline(memory=None,
                                     steps=[('random',
                                             RandomUnderSampler(random_state=42,
                                                                    replacement=False,
                                                                    sampling_strategy='auto')),
                                             ('rf',
                                              RandomForestClassifier(bootstrap=True,
                                                                      ccp_alpha=0.0,
                                                                      criterion='gini',
                                                                      max_depth=None,
                                                                      max_features='auto',
                                                                      max_leaf_nodes=None,
                                                                      max_samples=None,
                                                                      min_impurity_decrease=0.0,
                                                                      min_impurity_split=None,
                                                                      n_estimators=100,
                                                                      n_jobs=-1,
                                                                      oob_score=True,
                                                                      random_state=42,
                                                                      verbose=0,
                                                                      warm_start=False))],
                                     verbose=False),
                  iid='deprecated', n_jobs=-1,
                  param_grid={'rf__criterion': ['gini', 'entropy'],
                              'rf__max_depth': range(2, 15, 4),
                              'rf__max_features': ['auto', 'sqrt', 'log2'],
                              'rf__n_estimators': range(20, 100, 10)},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[356]: print("Best parameters: {}".format(DUSgrid_rf.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(DUSgrid_rf.best_score_))
print(f'Train score is {DUSgrid_rf.score(X_train,y_train)}')
```

```
Best parameters: {'rf__criterion': 'gini', 'rf__max_depth': 2,
'rf__max_features': 'auto', 'rf__n_estimators': 90}
Best Mean cross-validation score: 0.82
Train score is 0.8154670750382847
```

Easy Ensembler

```
[96]: #easy ensemble classifier
from imblearn.ensemble import EasyEnsembleClassifier
DUS_easy = Pipeline([('random', RandomUnderSampler()),
                      ('easy', EasyEnsembleClassifier())])
param_DUS_easy = {'easy__n_estimators': range(2, 50, 5)}
```

```
DUSgrid_easy= GridSearchCV(DUS_easy,param_DUS_easy, cv=5, n_jobs=-1,
    ↳scoring=f2score)
DUSgrid_easy.fit(X_train, y_train)
```

```
[96]: GridSearchCV(cv=5, error_score=nan,
            estimator=Pipeline(memory=None,
                                steps=[('random',
                                         RandomUnderSampler(random_state=None,
                                                                replacement=False,
                                                                sampling_strategy='auto')),
                                         ('easy',
                                          EasyEnsembleClassifier(base_estimator=None,
                                                                n_estimators=10,
                                                                n_jobs=None,
                                                                random_state=None,
                                                                replacement=False,
                                                                sampling_strategy='auto',
                                                                verbose=0,
                                                                warm_start=False))],
                                verbose=False),
            iid='deprecated', n_jobs=-1,
            param_grid={'easy__n_estimators': range(2, 50, 5)},
            pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
            scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

Results

```
[357]: print("Best parameters: {}".format(DUSgrid_easy.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(DUSgrid_easy.
    ↳best_score_))
print(f'Train score is {DUSgrid_easy.score(X_train,y_train)}')
```

```
Best parameters: {'easy__n_estimators': 37}
Best Mean cross-validation score: 0.46
Train score is 0.5091819699499165
```

XGBoost

```
[98]: #xgb
DUS_xgb = Pipeline([('random',RandomUnderSampler(random_state=42)),
                    ('xgboost',XGBClassifier(random_state=42,
                                              early_stopping_rounds=2,
                                              objective='binary:logistic',
                                              ))])
param_DUS_xgb = {'xgboost__max_depth' : range(2,8,2),
                 'xgboost__n_estimators' : range(50,80,20),
                 'xgboost__learning_rate' : [0.1],
                 'xgboost__min_child_weight' : [1,3,5,7],
```



```

        'xgboost__subsample': [0.6, 0.7, 0.8, 0.9, 1]}

DUSgrid_xgb= GridSearchCV(DUS_xgb,param_DUS_xgb, cv=5, n_jobs=-1,
    ↳scoring=f2score)
DUSgrid_xgb.fit(X_train, y_train)

```

```

[98]: GridSearchCV(cv=5, error_score=nan,
                  estimator=Pipeline(memory=None,
                                     steps=[('random',
                                             RandomUnderSampler(random_state=42,
                                                                    replacement=False,
                                                                    sampling_strategy='auto')),
                                             ('xgboost',
                                              XGBClassifier(base_score=None,
                                                            booster=None,
                                                            colsample_bylevel=None,
                                                            colsample_bynode=None,
                                                            colsample_bytree=None,
                                                            early_stopping_rounds=2,
                                                            gamma=None, gpu_id=None,
                                                            importance_type='gain',
                                                            verbosity=None))],
                                     verbose=False),
                  iid='deprecated', n_jobs=-1,
                  param_grid={'xgboost__learning_rate': [0.1],
                              'xgboost__max_depth': range(2, 8, 2),
                              'xgboost__min_child_weight': [1, 3, 5, 7],
                              'xgboost__n_estimators': range(50, 80, 20),
                              'xgboost__subsample': [0.6, 0.7, 0.8, 0.9, 1]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)

```

Results

```

[358]: print("Best parameters: {}".format(DUSgrid_xgb.best_params_))
        print("Best Mean cross-validation score: {:.2f}".format(DUSgrid_xgb.
    ↳best_score_))
        print(f'Train score is {DUSgrid_xgb.score(X_train,y_train)}')

```

```

Best parameters: {'xgboost__learning_rate': 0.1, 'xgboost__max_depth': 4,
'xgboost__min_child_weight': 7, 'xgboost__n_estimators': 50,
'xgboost__subsample': 0.6}
Best Mean cross-validation score: 0.59
Train score is 0.5963791267305645

```

3.4 Neural Network - MLP Classifier

```
[146]: from sklearn.neural_network import MLPClassifier
```

```
[157]: DUS_mlp = Pipeline([('random', RandomUnderSampler(random_state=42)),
                           ('mlp', MLPClassifier(random_state=42))])

parameters = {'mlp__solver': ['lbfgs'],
              'mlp__max_iter': [1200],
              'mlp__alpha': 10.0 ** -np.arange(1, 10),
              'mlp__hidden_layer_sizes': np.arange(10, 15),
              'mlp__random_state': range(2, 12, 2)}
MLPGrid = GridSearchCV(DUS_mlp, parameters, cv=5, n_jobs=-1, scoring=f2score)

MLPGrid.fit(X_train, y_train)
```

```
[157]: GridSearchCV(cv=5, error_score=nan,
                   estimator=Pipeline(memory=None,
                                       steps=[('random',
                                               RandomUnderSampler(random_state=42,
                                                                       replacement=False,
                                                                       sampling_strategy='auto')),
                                              ('mlp',
                                               MLPClassifier(activation='relu',
                                                             alpha=0.0001,
                                                             batch_size='auto',
                                                             beta_1=0.9, beta_2=0.999,
                                                             early_stopping=False,
                                                             epsilon=1e-08,
                                                             hidden_layer_sizes=(100,),
                                                             learning_rate='constant',
                                                             learning_rate_decay=0.0001,
                                                             max_iter=1000,
                                                             n_iter_no_change=10,
                                                             random_state=None,
                                                             solver='lbfgs',
                                                             tol=0.0001,
                                                             verbose=0))]),
                   iid='deprecated', n_jobs=-1,
                   param_grid={'mlp__alpha': array([1.e-01, 1.e-02, 1.e-03, 1.e-04,
1.e-05, 1.e-06, 1.e-07, 1.e-08,
1.e-09]),
                              'mlp__hidden_layer_sizes': array([10, 11, 12, 13, 14]),
                              'mlp__max_iter': [1200],
                              'mlp__random_state': range(2, 12, 2),
                              'mlp__solver': ['lbfgs']}},
                   pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                   scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

```
[359]: print("Best parameters: {}".format(MLPGrid.best_params_))
print("Best Mean cross-validation score: {:.2f}".format(MLPGrid.best_score_))
print(f'Train score is {MLPGrid.score(X_train, y_train)}')
```

Best parameters: {'mlp__alpha': 0.1, 'mlp__hidden_layer_sizes': 10,

```
'mlp__max_iter': 1200, 'mlp__random_state': 2, 'mlp__solver': 'lbfgs'}
Best Mean cross-validation score: 0.43
Train score is 0.4450930317402408
```

3.5 Stacking

4 Summary

```
[195]: classifiers={
    'Logistic':grid_logit,
    'DecisionTree':grid_dtree,
    'Kernel SVC':grid_svc,
    'RF':grid_rf,
    'Extra Trees':etc_grid,
    'Gradient Boost':gbc_grid,
    'XGBoost':xgbc_grid,
    'CS Logit':grid_CSlogreg,
    'CS DTree':grid_CSdtree,
    # 'CS RF':SCRandom_param,
    'CS XGBoost':grid_xgboost,
    'CS Extra Tree':grid_SCEXTree,
    'CS Bagging D Tree':grid_SCBagging,
    'Over logit':grid_logitsmote,
    'Over DTree':DOS_dtree,
    'Over RF':DOSgrid_rf,
    'Over Easy':DOSgrid_easy,
    'Over XGboost':DOSgrid_xgb,
    'Over MLP':DOSMLPGrid,
    'Under Logit':grid_logitRS,
    'Under Decision Tree': DUS_dtree,
    'Under RF':DUSgrid_rf,
    'Under Easy':DUSgrid_easy,
    'Under XGboost':DUSgrid_xgb,
    'Under MLP':MLPGrid
}
```

```
[196]: classifiers.keys()
```

```
[196]: dict_keys(['Logistic', 'DecisionTree', 'Kernel SVC', 'RF', 'Extra Trees',
    'Gradient Boost', 'XGBoost', 'CS Logit', 'CS DTree', 'CS XGBoost', 'CS Extra
    Tree', 'CS Bagging D Tree', 'Over logit', 'Over DTree', 'Over RF', 'Over Easy',
    'Over XGboost', 'Over MLP', 'Under Logit', 'Under Decision Tree', 'Under RF',
    'Under Easy', 'Under XGboost', 'Under MLP'])
```

```
[197]: results_mean_std = []
for key, value in classifiers.items():
    mean = value.cv_results_['mean_test_score'][value.best_index_]
```

```

std=value.cv_results_['std_test_score'][value.best_index_]

results_mean_std.append({
    "model": key,
    "mean": mean,
    "std": std
})

```

```

[198]: # Create a Pandas DataFrame with the mean+std results
accuracy_df = pd.DataFrame(results_mean_std, columns=['model', 'mean', 'std'])

```

```

[199]: # Show the accuracy dataframe

accuracy_df.sort_values(by=['mean'], inplace=True, ascending=False)
accuracy_df

```

```

[199]:

```

	model	mean	std
9	CS XGBoost	0.863311	0.046246
14	Over RF	0.861917	0.034663
6	XGBoost	0.853696	0.048157
3	RF	0.850095	0.055296
10	CS Extra Tree	0.842084	0.049814
16	Over XGboost	0.838657	0.050082
4	Extra Trees	0.832794	0.053444
20	Under RF	0.821151	0.029584
1	DecisionTree	0.814990	0.056385
0	Logistic	0.800555	0.027454
7	CS Logit	0.800555	0.027454
5	Gradient Boost	0.799898	0.050013
8	CS DTree	0.781130	0.036795
13	Over DTree	0.735339	0.047498
2	Kernel SVC	0.728533	0.068554
12	Over logit	0.696800	0.038193
18	Under Logit	0.689065	0.037545
15	Over Easy	0.675208	0.025761
22	Under XGboost	0.587185	0.026722
19	Under Decision Tree	0.514595	0.051219
11	CS Bagging D Tree	0.471119	0.053869
21	Under Easy	0.461867	0.032179
17	Over MLP	0.427471	0.045042
23	Under MLP	0.427471	0.045042

```

[200]: # Create a prediction of all models on the test set
predictions_all = {}
for key, value in classifiers.items():
    # Get best estimator
    best_model = value.best_estimator_

```


```
# Predict test labels
predictions = best_model.predict(X_test)

# Save predictions to a list
predictions_all[key] = predictions
```

```
[201]: # Creat a DataFrame for the predictions
pred = pd.DataFrame(predictions_all)
```

```
[206]: # Plot a heatmap of all correlations for easier visualization
fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(20,15))
g = sns.heatmap(pred.corr(), annot=True, cmap='coolwarm', ax=ax)
g.set_title('Correlation of the test set label prediction between models')
```

```
[206]: Text(0.5, 1, 'Correlation of the test set label prediction between models')
```



```
[207]: def get_redundant_pairs(df):
        '''Get diagonal and lower triangular pairs of correlation matrix'''
        pairs_to_drop = set()
        cols = df.columns
        for i in range(0, df.shape[1]):
            for j in range(0, i+1):
                pairs_to_drop.add((cols[i], cols[j]))
        return pairs_to_drop

def get_top_abs_correlations(df, n=5):
    au_corr = df.corr().abs().unstack()
    labels_to_drop = get_redundant_pairs(df)
    au_corr = au_corr.drop(labels=labels_to_drop).sort_values(ascending=True)
    return au_corr[0:n]
```

```
[209]: print("Top Least Correlations")
        print(get_top_abs_correlations(pred, 7))
```

```
Top Least Correlations
CS Bagging D Tree   Over DTree          0.238920
Kernel SVC          CS Bagging D Tree    0.276968
CS Extra Tree       CS Bagging D Tree    0.284768
Extra Trees         CS Bagging D Tree    0.284895
CS DTree            CS Bagging D Tree    0.285106
Logistic            CS Bagging D Tree    0.286314
CS Logit            CS Bagging D Tree    0.286314
dtype: float64
```

4.1 Voting top 5

```
[211]: from sklearn.ensemble import VotingClassifier
```

```
[212]: vclf1 = VotingClassifier(estimators=
                                [ ('CS XGBoost', grid_xgboost.best_estimator_),
                                  ('Over RF', DOSgrid_rf.best_estimator_),
                                  ('XGBoost', xgbc_grid.best_estimator_),
                                  ('RF', grid_rf.best_estimator_),
                                  ('CS Extra Tree', grid_SCEXtree.best_estimator_),
                                  ('Over XGboost', DOSgrid_xgb.best_estimator_),
                                  ('Extra Trees', etc_grid.best_estimator_),
                                ] )
vclf1_param = {'voting' : ['hard', 'soft']}
vclf1_grid = GridSearchCV(vclf1, vclf1_param, cv=5, n_jobs=-1,
    ↳return_train_score=True, scoring=f2score)
vclf1_grid.fit(X_train, y_train)
```

```
[212]: GridSearchCV(cv=5, error_score=nan,
                  estimator=VotingClassifier(estimators=[('CS XGBoost',
XGBClassifier(base_score=0.5,
                                                    booster=None,
                                                    colsample_bylevel=1,
                                                    colsample_bynode=1,
                                                    colsample_bytree=1,
                                                    gamma=0,
                                                    gpu_id=-1,
                                                    importance_type='gain',
                                                    interaction_constraints=None,
                                                    learning_rate=0.300000012,
                                                    max_delta_step=0,
                                                    max_depth=6,
                                                    min_child_weight=1,
                                                    missing=nan,
                                                    monotone...
                                                    min_weight_fraction_leaf=0.0,
                                                    n_estimators=50,
                                                    n_jobs=None,
                                                    oob_score=False,
                                                    random_state=42,
                                                    verbose=0,
                                                    warm_start=False))],
                  flatten_transform=True, n_jobs=None,
                  voting='hard', weights=None),
        iid='deprecated', n_jobs=-1,
        param_grid={'voting': ['hard', 'soft']}, pre_dispatch='2*n_jobs',
        refit=True, return_train_score=True,
        scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

4.1.1 Results

```
[213]: print(f'Best Mean Cross Validation Score is {vc1f1_grid.best_score_}')
print(f'Best Mean Cross Validation Score is {vc1f1_grid.best_params_}')
print(f'Train score is {vc1f1_grid.score(X_train,y_train)}')
```

```
Best Mean Cross Validation Score is 0.8599431114386332
Best Mean Cross Validation Score is {'voting': 'hard'}
Train score is 0.940347970173985
```

4.2 Stacking Top 5

<https://towardsdatascience.com/stacking-made-easy-with-sklearn-e27a0793c92b>

```
[217]: from sklearn.ensemble import StackingClassifier
```

4.2.1 Combination 1 - .81352

```
[222]: sclf1 = StackingClassifier(estimators=[
    ('CS XGBoost', grid_xgboost.best_estimator_),
    ('Over RF', D0Sgrid_rf.best_estimator_),
    ('XGBoost', xgbc_grid.best_estimator_),
    ('RF', grid_rf.best_estimator_),
    ('CS Extra Tree', grid_SCEXtree.best_estimator_),
    ('Over XGboost', D0Sgrid_xgb.best_estimator_),
    ('Extra Trees', etc_grid.best_estimator_),
    ], final_estimator=LogisticRegression(
    multi_class="multinomial", solver="lbfgs", C=30))

sclf1_param = {
    'final_estimator__C': [0.001, 0.01, 0.1, 1, 10, 1000],
    'final_estimator__penalty': ['l2'],
    'stack_method': ['auto', 'predict_proba']
}

sclf1_grid = GridSearchCV(sclf1, sclf1_param, cv=5,
    ↪return_train_score=True, n_jobs=-1, scoring=f2score )
sclf1_grid.fit(X_train, y_train)
```

```
[222]: GridSearchCV(cv=5, error_score=nan,
    estimator=StackingClassifier(cv=None,
    estimators=[('CS XGBoost',
    XGBClassifier(base_score=0.5,
    booster=None,
    colsample_bylevel=1,
    colsample_bynode=1,
    colsample_bytrees=1,
    gamma=0,
    gpu_id=-1,
    importance_type='gain',
    interaction_constraints=None,
    learning_rate=0.300000012,
    max_delta_step=0,
    max_depth=6,
    min_child_weight=1,
    missing=na...,
    warm_start=False),
    n_jobs=None, passthrough=False,
    stack_method='auto', verbose=0),
    iid='deprecated', n_jobs=-1,
    param_grid={'final_estimator__C': [0.001, 0.01, 0.1, 1, 10, 1000],
    'final_estimator__penalty': ['l2'],
    'stack_method': ['auto', 'predict_proba']},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
```



```
scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

```
[223]: print(f'Best Mean Cross Validation Score is {sclf1_grid.best_score_}')
print(f'Best Mean Cross Validation Score is {sclf1_grid.best_params_}')
print(f'Train score is {sclf1_grid.score(X_train,y_train)}')
```

```
Best Mean Cross Validation Score is 0.8478698044874516
Best Mean Cross Validation Score is {'final_estimator__C': 1,
'final_estimator__penalty': 'l2', 'stack_method': 'auto'}
Train score is 0.9835390946502058
```

4.2.2 Combination 2 -.81352

```
[225]: sclf2 = StackingClassifier(estimators=
                                [ ('CS XGBoost', grid_xgboost.best_estimator_),
                                  ('Over RF', DOSgrid_rf.best_estimator_),
                                  ('XGBoost', xgbc_grid.best_estimator_),
                                  ('RF', grid_rf.best_estimator_),
                                  ('CS Extra Tree', grid_SCEXtree.best_estimator_),
                                  ('Over XGboost', DOSgrid_xgb.best_estimator_),
                                  ('Extra Trees', etc_grid.best_estimator_)
                                ] , final_estimator=sclf1_grid.best_estimator_)

sclf2_param = {
#             'final_estimator__C': [0.001, 0.01, 0.1, 1, 10,1000],
#             'final_estimator__penalty':['l2'],
            'stack_method':['auto', 'predict_proba']
}
sclf2_grid = GridSearchCV(sclf2, sclf2_param,cv=5,
    ↪return_train_score=True,scoring=f2score)
sclf2_grid.fit(X_train,y_train)
```

```
[225]: GridSearchCV(cv=5, error_score=nan,
                  estimator=StackingClassifier(cv=None,
                                                estimators=[('CS XGBoost',
XGBClassifier(base_score=0.5,
booster=None,
colsample_bylevel=1,
colsample_bynode=1,
colsample_bytree=1,
importance_type='gain',
interaction_constraints=None,
learning_rate=0.300000012,
max_delta_step=0,
max_depth=6,
gamma=0,
gpu_id=-1,
```

```

min_child_weight=1,
missing=na...

                                tol=0.0001,
                                verbose=0,
                                warm_start=False),

n_jobs=None,
passthrough=False,
stack_method='auto',
verbose=0),

                                n_jobs=None, passthrough=False,
                                stack_method='auto', verbose=0),

iid='deprecated', n_jobs=None,
param_grid={'stack_method': ['auto', 'predict_proba']},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=make_scorer(fbeta_score, beta=2), verbose=0)

```

```

[226]: print(f'Best Mean Cross Validation Score is {sclf2_grid.best_score_}')
        print(f'Best Mean Cross Validation Score is {sclf2_grid.best_params_}')
        print(f'Train score is {sclf2_grid.score(X_train,y_train)}')

```

```

Best Mean Cross Validation Score is 0.8595398428731761
Best Mean Cross Validation Score is {'stack_method': 'auto'}
Train score is 0.9535655058043118

```

4.2.3 Combination 3 - .81204

```

[227]: sclf3 = StackingClassifier(estimators=
                                [ ('CS XGBoost', grid_xgboost.best_estimator_),
                                  ('Over RF', DOSgrid_rf.best_estimator_),
                                  ('XGBoost', xgbc_grid.best_estimator_),
                                  ('RF', grid_rf.best_estimator_),
                                  ('CS Extra Tree', grid_SCEXtree.best_estimator_),
                                  ('Over XGboost', DOSgrid_xgb.best_estimator_),
                                  ('Extra Trees', etc_grid.best_estimator_),
                                ] , final_estimator=sclf2_grid.best_estimator_)

sclf3_param = {
#           'final_estimator__C': [0.001, 0.01, 0.1, 1, 10,1000],
#           'final_estimator__penalty':['l2'],
    'stack_method':['auto', 'predict_proba']
}
sclf3_grid = GridSearchCV(sclf3, sclf3_param,cv=5, return_train_score=True )
sclf3_grid.fit(X_train,y_train)

```

```

[227]: GridSearchCV(cv=5, error_score=nan,
                    estimator=StackingClassifier(cv=None,
                                                  estimators=[('CS XGBoost',

```

```

XGBClassifier(base_score=0.5,
booster=None,
colsample_bylevel=1,
colsample_bynode=1,
colsample_bytree=1,

                                                    gamma=0,
                                                    gpu_id=-1,

importance_type='gain',
interaction_constraints=None,
learning_rate=0.300000012,
max_delta_step=0,
max_depth=6,
min_child_weight=1,
missing=na...
warm_start=False),

                                                    n_jobs=None,
                                                    passthrough=False,
                                                    stack_method='auto',
                                                    verbose=0),

n_jobs=None,
passthrough=False,
stack_method='auto',
verbose=0),

                                                    n_jobs=None, passthrough=False,
                                                    stack_method='auto', verbose=0),

iid='deprecated', n_jobs=None,
param_grid={'stack_method': ['auto', 'predict_proba']},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=None, verbose=0)

```

```

[382]: print(f'Best Mean Cross Validation Score is {sclf3_grid.best_score_}')
       print(f'Best Parameters are {sclf3_grid.best_params_}')
       print(f'Train score is {sclf2_grid.score(X_train,y_train)}')

```

```

Best Mean Cross Validation Score is 0.9978266297321057
Best Parameters are {'stack_method': 'auto'}
Train score is 0.9535655058043118

```

4.3 Other Methods

4.3.1 RandomizedSearchCV with Random Forest

```

[234]: from sklearn.model_selection import RandomizedSearchCV

```

```

[235]: #random forest
       from sklearn.ensemble import RandomForestClassifier
       rfc =RandomForestClassifier(random_state=0)
       rfc_param = {

```

```

    'n_estimators': range(2,15,2),
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth' : range(2,15,3),
    'criterion' :['gini', 'entropy']
}

rsearchcv_rf = RandomizedSearchCV(rfc, rfc_param,cv=5, return_train_score=True,
    ↳scoring =f2score)
rsearchcv_rf.fit(X_train,y_train)

```

```

[235]: RandomizedSearchCV(cv=5, error_score=nan,
                        estimator=RandomForestClassifier(bootstrap=True,
                                                            ccp_alpha=0.0,
                                                            class_weight=None,
                                                            criterion='gini',
                                                            max_depth=None,
                                                            max_features='auto',
                                                            max_leaf_nodes=None,
                                                            max_samples=None,
                                                            min_impurity_decrease=0.0,
                                                            min_impurity_split=None,
                                                            min_samples_leaf=1,
                                                            min_samples_split=2,
                                                            min_weight_fraction_leaf=0.0,
                                                            n_estimators=100,
                                                            n_jobs...
                                                            random_state=0, verbose=0,
                                                            warm_start=False),
                        iid='deprecated', n_iter=10, n_jobs=None,
                        param_distributions={'criterion': ['gini', 'entropy'],
                                             'max_depth': range(2, 15, 3),
                                             'max_features': ['auto', 'sqrt',
                                                             'log2'],
                                             'n_estimators': range(2, 15, 2)},
                        pre_dispatch='2*n_jobs', random_state=None, refit=True,
                        return_train_score=True,
                        scoring=make_scorer(fbeta_score, beta=2), verbose=0)

```

Results

```

[236]: print(f'Best Mean Cross Validation Score is {rsearchcv_rf.best_params_}')
       print(f'Best Mean Cross Validation Score is {rsearchcv_rf.best_score_}')
       print(f'Train score is {rsearchcv_rf.score(X_train,y_train)}')

```

```

Best Mean Cross Validation Score is {'n_estimators': 12, 'max_features': 'auto',
'max_depth': 8, 'criterion': 'gini'}
Best Mean Cross Validation Score is 0.8415861898736801
Train score is 0.871404399323181

```

4.3.2 CatBoost Classifier - .81352

```
[239]: #pip install catboost
```

```
[243]: from catboost import CatBoostClassifier
```

```
[308]: catb = CatBoostClassifier()

catb_param = {'learning_rate': [0.03, 0.1],
              'depth': range(2,10,2),
              'l2_leaf_reg': range(2,10,2)}

catb_grid = GridSearchCV(catb, catb_param,cv=5,n_jobs=-1,
    ↪return_train_score=True, scoring =f2score)
catb_grid.fit(X_train,y_train)
```

0:	learn: 0.6074288	total: 74.3ms	remaining: 1m 14s
1:	learn: 0.5308762	total: 88ms	remaining: 43.9s
2:	learn: 0.4647003	total: 98.8ms	remaining: 32.8s
3:	learn: 0.4103385	total: 110ms	remaining: 27.4s
4:	learn: 0.3589659	total: 122ms	remaining: 24.3s
5:	learn: 0.3156362	total: 133ms	remaining: 22.1s
6:	learn: 0.2759750	total: 146ms	remaining: 20.6s
7:	learn: 0.2440715	total: 159ms	remaining: 19.7s
8:	learn: 0.2144962	total: 171ms	remaining: 18.8s
9:	learn: 0.1885638	total: 182ms	remaining: 18s
10:	learn: 0.1670417	total: 196ms	remaining: 17.6s
11:	learn: 0.1476058	total: 208ms	remaining: 17.1s
12:	learn: 0.1331938	total: 220ms	remaining: 16.7s
13:	learn: 0.1190247	total: 233ms	remaining: 16.4s
14:	learn: 0.1061889	total: 245ms	remaining: 16.1s
15:	learn: 0.0950364	total: 259ms	remaining: 15.9s
16:	learn: 0.0854880	total: 273ms	remaining: 15.8s
17:	learn: 0.0770394	total: 286ms	remaining: 15.6s
18:	learn: 0.0698292	total: 297ms	remaining: 15.3s
19:	learn: 0.0635772	total: 309ms	remaining: 15.1s
20:	learn: 0.0577759	total: 322ms	remaining: 15s
21:	learn: 0.0528676	total: 335ms	remaining: 14.9s
22:	learn: 0.0485986	total: 347ms	remaining: 14.8s
23:	learn: 0.0449189	total: 364ms	remaining: 14.8s
24:	learn: 0.0415947	total: 375ms	remaining: 14.6s
25:	learn: 0.0387336	total: 388ms	remaining: 14.5s
26:	learn: 0.0361942	total: 400ms	remaining: 14.4s
27:	learn: 0.0339457	total: 411ms	remaining: 14.3s
28:	learn: 0.0318975	total: 422ms	remaining: 14.1s
29:	learn: 0.0300825	total: 433ms	remaining: 14s
30:	learn: 0.0282923	total: 446ms	remaining: 13.9s
31:	learn: 0.0268204	total: 461ms	remaining: 13.9s

32:	learn: 0.0255583	total: 474ms	remaining: 13.9s
33:	learn: 0.0243798	total: 490ms	remaining: 13.9s
34:	learn: 0.0233485	total: 504ms	remaining: 13.9s
35:	learn: 0.0224139	total: 516ms	remaining: 13.8s
36:	learn: 0.0215928	total: 529ms	remaining: 13.8s
37:	learn: 0.0208351	total: 541ms	remaining: 13.7s
38:	learn: 0.0201242	total: 553ms	remaining: 13.6s
39:	learn: 0.0195033	total: 565ms	remaining: 13.6s
40:	learn: 0.0189460	total: 577ms	remaining: 13.5s
41:	learn: 0.0182528	total: 590ms	remaining: 13.4s
42:	learn: 0.0177152	total: 603ms	remaining: 13.4s
43:	learn: 0.0172876	total: 616ms	remaining: 13.4s
44:	learn: 0.0168570	total: 629ms	remaining: 13.4s
45:	learn: 0.0164760	total: 644ms	remaining: 13.4s
46:	learn: 0.0161703	total: 656ms	remaining: 13.3s
47:	learn: 0.0158765	total: 668ms	remaining: 13.2s
48:	learn: 0.0155867	total: 680ms	remaining: 13.2s
49:	learn: 0.0152391	total: 692ms	remaining: 13.1s
50:	learn: 0.0150067	total: 704ms	remaining: 13.1s
51:	learn: 0.0148255	total: 716ms	remaining: 13.1s
52:	learn: 0.0145782	total: 727ms	remaining: 13s
53:	learn: 0.0143297	total: 739ms	remaining: 12.9s
54:	learn: 0.0141875	total: 751ms	remaining: 12.9s
55:	learn: 0.0139818	total: 763ms	remaining: 12.9s
56:	learn: 0.0137251	total: 774ms	remaining: 12.8s
57:	learn: 0.0135430	total: 786ms	remaining: 12.8s
58:	learn: 0.0133772	total: 799ms	remaining: 12.7s
59:	learn: 0.0132640	total: 811ms	remaining: 12.7s
60:	learn: 0.0131530	total: 824ms	remaining: 12.7s
61:	learn: 0.0130523	total: 837ms	remaining: 12.7s
62:	learn: 0.0129569	total: 1.18s	remaining: 17.6s
63:	learn: 0.0128589	total: 1.19s	remaining: 17.5s
64:	learn: 0.0126957	total: 1.21s	remaining: 17.4s
65:	learn: 0.0125399	total: 1.22s	remaining: 17.2s
66:	learn: 0.0124000	total: 1.23s	remaining: 17.1s
67:	learn: 0.0122751	total: 1.24s	remaining: 17s
68:	learn: 0.0121453	total: 1.25s	remaining: 16.9s
69:	learn: 0.0120447	total: 1.27s	remaining: 16.8s
70:	learn: 0.0119707	total: 1.28s	remaining: 16.7s
71:	learn: 0.0119007	total: 1.29s	remaining: 16.6s
72:	learn: 0.0118277	total: 1.3s	remaining: 16.5s
73:	learn: 0.0117382	total: 1.31s	remaining: 16.4s
74:	learn: 0.0116948	total: 1.33s	remaining: 16.4s
75:	learn: 0.0116481	total: 1.34s	remaining: 16.3s
76:	learn: 0.0116015	total: 1.35s	remaining: 16.2s
77:	learn: 0.0115148	total: 1.36s	remaining: 16.1s
78:	learn: 0.0114770	total: 1.38s	remaining: 16.1s
79:	learn: 0.0114308	total: 1.39s	remaining: 16s

80:	learn: 0.0113522	total: 1.4s	remaining: 15.9s
81:	learn: 0.0112771	total: 1.42s	remaining: 15.9s
82:	learn: 0.0112396	total: 1.43s	remaining: 15.8s
83:	learn: 0.0111904	total: 1.44s	remaining: 15.7s
84:	learn: 0.0111437	total: 1.45s	remaining: 15.6s
85:	learn: 0.0110683	total: 1.46s	remaining: 15.5s
86:	learn: 0.0110223	total: 1.47s	remaining: 15.5s
87:	learn: 0.0109570	total: 1.49s	remaining: 15.4s
88:	learn: 0.0108920	total: 1.5s	remaining: 15.3s
89:	learn: 0.0108630	total: 1.51s	remaining: 15.3s
90:	learn: 0.0108245	total: 1.52s	remaining: 15.2s
91:	learn: 0.0107678	total: 1.53s	remaining: 15.1s
92:	learn: 0.0106960	total: 1.55s	remaining: 15.1s
93:	learn: 0.0106408	total: 1.56s	remaining: 15s
94:	learn: 0.0105965	total: 1.57s	remaining: 15s
95:	learn: 0.0105478	total: 1.58s	remaining: 14.9s
96:	learn: 0.0105234	total: 1.6s	remaining: 14.9s
97:	learn: 0.0104791	total: 1.61s	remaining: 14.8s
98:	learn: 0.0104317	total: 1.62s	remaining: 14.7s
99:	learn: 0.0104177	total: 1.63s	remaining: 14.7s
100:	learn: 0.0103786	total: 1.64s	remaining: 14.6s
101:	learn: 0.0103596	total: 1.65s	remaining: 14.6s
102:	learn: 0.0103284	total: 1.66s	remaining: 14.5s
103:	learn: 0.0103166	total: 1.68s	remaining: 14.4s
104:	learn: 0.0102865	total: 1.69s	remaining: 14.4s
105:	learn: 0.0102276	total: 1.7s	remaining: 14.3s
106:	learn: 0.0101953	total: 1.71s	remaining: 14.3s
107:	learn: 0.0101559	total: 1.73s	remaining: 14.3s
108:	learn: 0.0101250	total: 1.74s	remaining: 14.2s
109:	learn: 0.0100993	total: 1.75s	remaining: 14.2s
110:	learn: 0.0100575	total: 1.76s	remaining: 14.1s
111:	learn: 0.0100238	total: 1.77s	remaining: 14.1s
112:	learn: 0.0099926	total: 1.79s	remaining: 14s
113:	learn: 0.0099663	total: 1.8s	remaining: 14s
114:	learn: 0.0099471	total: 1.81s	remaining: 13.9s
115:	learn: 0.0099176	total: 1.82s	remaining: 13.9s
116:	learn: 0.0099030	total: 1.84s	remaining: 13.9s
117:	learn: 0.0098836	total: 1.85s	remaining: 13.8s
118:	learn: 0.0098587	total: 1.86s	remaining: 13.8s
119:	learn: 0.0098566	total: 1.87s	remaining: 13.7s
120:	learn: 0.0098288	total: 1.88s	remaining: 13.7s
121:	learn: 0.0098034	total: 1.89s	remaining: 13.6s
122:	learn: 0.0097571	total: 1.91s	remaining: 13.6s
123:	learn: 0.0097284	total: 1.92s	remaining: 13.6s
124:	learn: 0.0097179	total: 1.93s	remaining: 13.5s
125:	learn: 0.0097039	total: 1.94s	remaining: 13.5s
126:	learn: 0.0096849	total: 1.95s	remaining: 13.4s
127:	learn: 0.0096576	total: 1.97s	remaining: 13.4s

128:	learn: 0.0096399	total: 1.98s	remaining: 13.4s
129:	learn: 0.0096171	total: 1.99s	remaining: 13.3s
130:	learn: 0.0096149	total: 2s	remaining: 13.3s
131:	learn: 0.0095878	total: 2.02s	remaining: 13.3s
132:	learn: 0.0095664	total: 2.03s	remaining: 13.2s
133:	learn: 0.0095468	total: 2.04s	remaining: 13.2s
134:	learn: 0.0095134	total: 2.05s	remaining: 13.1s
135:	learn: 0.0094953	total: 2.06s	remaining: 13.1s
136:	learn: 0.0094830	total: 2.07s	remaining: 13.1s
137:	learn: 0.0094601	total: 2.08s	remaining: 13s
138:	learn: 0.0094482	total: 2.1s	remaining: 13s
139:	learn: 0.0094423	total: 2.11s	remaining: 13s
140:	learn: 0.0094401	total: 2.12s	remaining: 12.9s
141:	learn: 0.0094119	total: 2.13s	remaining: 12.9s
142:	learn: 0.0093963	total: 2.15s	remaining: 12.9s
143:	learn: 0.0093849	total: 2.16s	remaining: 12.8s
144:	learn: 0.0093830	total: 2.17s	remaining: 12.8s
145:	learn: 0.0093510	total: 2.18s	remaining: 12.8s
146:	learn: 0.0093376	total: 2.19s	remaining: 12.7s
147:	learn: 0.0093331	total: 2.2s	remaining: 12.7s
148:	learn: 0.0093199	total: 2.22s	remaining: 12.7s
149:	learn: 0.0093172	total: 2.23s	remaining: 12.6s
150:	learn: 0.0092960	total: 2.24s	remaining: 12.6s
151:	learn: 0.0092748	total: 2.25s	remaining: 12.6s
152:	learn: 0.0092636	total: 2.26s	remaining: 12.5s
153:	learn: 0.0092411	total: 2.27s	remaining: 12.5s
154:	learn: 0.0092309	total: 2.29s	remaining: 12.5s
155:	learn: 0.0092167	total: 2.3s	remaining: 12.4s
156:	learn: 0.0091942	total: 2.32s	remaining: 12.4s
157:	learn: 0.0091918	total: 2.33s	remaining: 12.4s
158:	learn: 0.0091694	total: 2.34s	remaining: 12.4s
159:	learn: 0.0091593	total: 2.35s	remaining: 12.3s
160:	learn: 0.0091515	total: 2.36s	remaining: 12.3s
161:	learn: 0.0091335	total: 2.37s	remaining: 12.3s
162:	learn: 0.0090960	total: 2.38s	remaining: 12.2s
163:	learn: 0.0090865	total: 2.4s	remaining: 12.2s
164:	learn: 0.0090712	total: 2.41s	remaining: 12.2s
165:	learn: 0.0090689	total: 2.42s	remaining: 12.2s
166:	learn: 0.0090331	total: 2.43s	remaining: 12.1s
167:	learn: 0.0090159	total: 2.44s	remaining: 12.1s
168:	learn: 0.0090040	total: 2.46s	remaining: 12.1s
169:	learn: 0.0089861	total: 2.47s	remaining: 12.1s
170:	learn: 0.0089605	total: 2.48s	remaining: 12s
171:	learn: 0.0089486	total: 2.5s	remaining: 12s
172:	learn: 0.0089447	total: 2.51s	remaining: 12s
173:	learn: 0.0089322	total: 2.52s	remaining: 12s
174:	learn: 0.0089241	total: 2.54s	remaining: 12s
175:	learn: 0.0089140	total: 2.55s	remaining: 11.9s

176:	learn: 0.0089126	total: 2.56s	remaining: 11.9s
177:	learn: 0.0088983	total: 2.58s	remaining: 11.9s
178:	learn: 0.0088824	total: 2.59s	remaining: 11.9s
179:	learn: 0.0088689	total: 2.6s	remaining: 11.8s
180:	learn: 0.0088522	total: 2.61s	remaining: 11.8s
181:	learn: 0.0088281	total: 2.63s	remaining: 11.8s
182:	learn: 0.0088197	total: 2.64s	remaining: 11.8s
183:	learn: 0.0087891	total: 2.65s	remaining: 11.8s
184:	learn: 0.0087813	total: 2.67s	remaining: 11.7s
185:	learn: 0.0087796	total: 2.68s	remaining: 11.7s
186:	learn: 0.0087594	total: 2.69s	remaining: 11.7s
187:	learn: 0.0087363	total: 2.7s	remaining: 11.7s
188:	learn: 0.0087138	total: 2.72s	remaining: 11.7s
189:	learn: 0.0087033	total: 2.73s	remaining: 11.6s
190:	learn: 0.0086909	total: 2.74s	remaining: 11.6s
191:	learn: 0.0086831	total: 2.75s	remaining: 11.6s
192:	learn: 0.0086796	total: 2.76s	remaining: 11.6s
193:	learn: 0.0086758	total: 2.78s	remaining: 11.5s
194:	learn: 0.0086676	total: 2.79s	remaining: 11.5s
195:	learn: 0.0086558	total: 2.8s	remaining: 11.5s
196:	learn: 0.0086457	total: 2.81s	remaining: 11.5s
197:	learn: 0.0086353	total: 2.83s	remaining: 11.5s
198:	learn: 0.0086154	total: 2.85s	remaining: 11.5s
199:	learn: 0.0085960	total: 2.86s	remaining: 11.5s
200:	learn: 0.0085940	total: 2.88s	remaining: 11.4s
201:	learn: 0.0085924	total: 2.89s	remaining: 11.4s
202:	learn: 0.0085859	total: 2.9s	remaining: 11.4s
203:	learn: 0.0085711	total: 2.91s	remaining: 11.4s
204:	learn: 0.0085569	total: 2.92s	remaining: 11.3s
205:	learn: 0.0085525	total: 2.94s	remaining: 11.3s
206:	learn: 0.0085327	total: 2.95s	remaining: 11.3s
207:	learn: 0.0085221	total: 2.96s	remaining: 11.3s
208:	learn: 0.0085071	total: 2.98s	remaining: 11.3s
209:	learn: 0.0085038	total: 2.99s	remaining: 11.2s
210:	learn: 0.0084978	total: 3s	remaining: 11.2s
211:	learn: 0.0084891	total: 3.01s	remaining: 11.2s
212:	learn: 0.0084724	total: 3.02s	remaining: 11.2s
213:	learn: 0.0084635	total: 3.04s	remaining: 11.2s
214:	learn: 0.0084551	total: 3.05s	remaining: 11.1s
215:	learn: 0.0084476	total: 3.06s	remaining: 11.1s
216:	learn: 0.0084398	total: 3.07s	remaining: 11.1s
217:	learn: 0.0084308	total: 3.08s	remaining: 11.1s
218:	learn: 0.0084178	total: 3.1s	remaining: 11.1s
219:	learn: 0.0084159	total: 3.11s	remaining: 11s
220:	learn: 0.0084010	total: 3.12s	remaining: 11s
221:	learn: 0.0083927	total: 3.13s	remaining: 11s
222:	learn: 0.0083881	total: 3.15s	remaining: 11s
223:	learn: 0.0083672	total: 3.16s	remaining: 11s

224:	learn: 0.0083568	total: 3.18s	remaining: 10.9s
225:	learn: 0.0083512	total: 3.19s	remaining: 10.9s
226:	learn: 0.0083482	total: 3.2s	remaining: 10.9s
227:	learn: 0.0083350	total: 3.22s	remaining: 10.9s
228:	learn: 0.0083301	total: 3.23s	remaining: 10.9s
229:	learn: 0.0083051	total: 3.24s	remaining: 10.8s
230:	learn: 0.0082948	total: 3.25s	remaining: 10.8s
231:	learn: 0.0082875	total: 3.27s	remaining: 10.8s
232:	learn: 0.0082712	total: 3.28s	remaining: 10.8s
233:	learn: 0.0082668	total: 3.29s	remaining: 10.8s
234:	learn: 0.0082515	total: 3.3s	remaining: 10.7s
235:	learn: 0.0082294	total: 3.31s	remaining: 10.7s
236:	learn: 0.0082167	total: 3.33s	remaining: 10.7s
237:	learn: 0.0082069	total: 3.34s	remaining: 10.7s
238:	learn: 0.0081762	total: 3.35s	remaining: 10.7s
239:	learn: 0.0081611	total: 3.36s	remaining: 10.6s
240:	learn: 0.0081485	total: 3.37s	remaining: 10.6s
241:	learn: 0.0081409	total: 3.39s	remaining: 10.6s
242:	learn: 0.0081302	total: 3.4s	remaining: 10.6s
243:	learn: 0.0080969	total: 3.41s	remaining: 10.6s
244:	learn: 0.0080939	total: 3.42s	remaining: 10.5s
245:	learn: 0.0080789	total: 3.43s	remaining: 10.5s
246:	learn: 0.0080545	total: 3.44s	remaining: 10.5s
247:	learn: 0.0080337	total: 3.46s	remaining: 10.5s
248:	learn: 0.0080285	total: 3.47s	remaining: 10.5s
249:	learn: 0.0080108	total: 3.48s	remaining: 10.4s
250:	learn: 0.0079902	total: 3.49s	remaining: 10.4s
251:	learn: 0.0079786	total: 3.5s	remaining: 10.4s
252:	learn: 0.0079688	total: 3.51s	remaining: 10.4s
253:	learn: 0.0079627	total: 3.52s	remaining: 10.4s
254:	learn: 0.0079594	total: 3.54s	remaining: 10.3s
255:	learn: 0.0079555	total: 3.55s	remaining: 10.3s
256:	learn: 0.0079531	total: 3.56s	remaining: 10.3s
257:	learn: 0.0079377	total: 3.57s	remaining: 10.3s
258:	learn: 0.0079306	total: 3.59s	remaining: 10.3s
259:	learn: 0.0078985	total: 3.6s	remaining: 10.2s
260:	learn: 0.0078890	total: 3.61s	remaining: 10.2s
261:	learn: 0.0078833	total: 3.62s	remaining: 10.2s
262:	learn: 0.0078820	total: 3.64s	remaining: 10.2s
263:	learn: 0.0078695	total: 3.65s	remaining: 10.2s
264:	learn: 0.0078604	total: 3.66s	remaining: 10.2s
265:	learn: 0.0078554	total: 3.67s	remaining: 10.1s
266:	learn: 0.0078500	total: 3.68s	remaining: 10.1s
267:	learn: 0.0078415	total: 3.7s	remaining: 10.1s
268:	learn: 0.0078235	total: 3.71s	remaining: 10.1s
269:	learn: 0.0078173	total: 3.72s	remaining: 10.1s
270:	learn: 0.0078102	total: 3.74s	remaining: 10.1s
271:	learn: 0.0077923	total: 3.75s	remaining: 10s

272:	learn: 0.0077823	total: 3.76s	remaining: 10s
273:	learn: 0.0077589	total: 3.77s	remaining: 10s
274:	learn: 0.0077535	total: 3.79s	remaining: 9.98s
275:	learn: 0.0077488	total: 3.8s	remaining: 9.96s
276:	learn: 0.0077324	total: 3.81s	remaining: 9.95s
277:	learn: 0.0077154	total: 3.82s	remaining: 9.93s
278:	learn: 0.0077028	total: 3.85s	remaining: 9.94s
279:	learn: 0.0076991	total: 3.86s	remaining: 9.92s
280:	learn: 0.0076967	total: 3.87s	remaining: 9.9s
281:	learn: 0.0076937	total: 3.88s	remaining: 9.88s
282:	learn: 0.0076839	total: 3.89s	remaining: 9.86s
283:	learn: 0.0076780	total: 3.91s	remaining: 9.86s
284:	learn: 0.0076676	total: 3.92s	remaining: 9.85s
285:	learn: 0.0076652	total: 3.94s	remaining: 9.83s
286:	learn: 0.0076604	total: 3.95s	remaining: 9.81s
287:	learn: 0.0076587	total: 3.96s	remaining: 9.79s
288:	learn: 0.0076537	total: 3.97s	remaining: 9.77s
289:	learn: 0.0076512	total: 3.98s	remaining: 9.76s
290:	learn: 0.0076212	total: 4s	remaining: 9.74s
291:	learn: 0.0076102	total: 4.01s	remaining: 9.72s
292:	learn: 0.0076046	total: 4.02s	remaining: 9.7s
293:	learn: 0.0075898	total: 4.03s	remaining: 9.69s
294:	learn: 0.0075876	total: 4.05s	remaining: 9.67s
295:	learn: 0.0075798	total: 4.06s	remaining: 9.65s
296:	learn: 0.0075728	total: 4.07s	remaining: 9.64s
297:	learn: 0.0075545	total: 4.1s	remaining: 9.65s
298:	learn: 0.0075443	total: 4.11s	remaining: 9.63s
299:	learn: 0.0075409	total: 4.12s	remaining: 9.61s
300:	learn: 0.0075304	total: 4.13s	remaining: 9.59s
301:	learn: 0.0075129	total: 4.14s	remaining: 9.57s
302:	learn: 0.0074973	total: 4.15s	remaining: 9.55s
303:	learn: 0.0074830	total: 4.17s	remaining: 9.54s
304:	learn: 0.0074810	total: 4.18s	remaining: 9.52s
305:	learn: 0.0074606	total: 4.19s	remaining: 9.51s
306:	learn: 0.0074567	total: 4.2s	remaining: 9.48s
307:	learn: 0.0074453	total: 4.21s	remaining: 9.46s
308:	learn: 0.0074216	total: 4.22s	remaining: 9.45s
309:	learn: 0.0074080	total: 4.24s	remaining: 9.44s
310:	learn: 0.0074010	total: 4.25s	remaining: 9.42s
311:	learn: 0.0073730	total: 4.27s	remaining: 9.41s
312:	learn: 0.0073624	total: 4.28s	remaining: 9.39s
313:	learn: 0.0073583	total: 4.29s	remaining: 9.37s
314:	learn: 0.0073540	total: 4.3s	remaining: 9.35s
315:	learn: 0.0073436	total: 4.31s	remaining: 9.33s
316:	learn: 0.0073408	total: 4.33s	remaining: 9.32s
317:	learn: 0.0073362	total: 4.34s	remaining: 9.3s
318:	learn: 0.0073288	total: 4.35s	remaining: 9.28s
319:	learn: 0.0073226	total: 4.36s	remaining: 9.27s

320:	learn: 0.0073187	total: 4.37s	remaining: 9.25s
321:	learn: 0.0073140	total: 4.38s	remaining: 9.23s
322:	learn: 0.0073116	total: 4.4s	remaining: 9.22s
323:	learn: 0.0073033	total: 4.41s	remaining: 9.21s
324:	learn: 0.0073003	total: 4.42s	remaining: 9.19s
325:	learn: 0.0072903	total: 4.44s	remaining: 9.17s
326:	learn: 0.0072885	total: 4.45s	remaining: 9.16s
327:	learn: 0.0072809	total: 4.47s	remaining: 9.15s
328:	learn: 0.0072781	total: 4.48s	remaining: 9.13s
329:	learn: 0.0072611	total: 4.49s	remaining: 9.12s
330:	learn: 0.0072475	total: 4.5s	remaining: 9.1s
331:	learn: 0.0072246	total: 4.52s	remaining: 9.09s
332:	learn: 0.0072221	total: 4.53s	remaining: 9.07s
333:	learn: 0.0072167	total: 4.54s	remaining: 9.05s
334:	learn: 0.0071953	total: 4.55s	remaining: 9.04s
335:	learn: 0.0071836	total: 4.58s	remaining: 9.05s
336:	learn: 0.0071741	total: 4.59s	remaining: 9.03s
337:	learn: 0.0071670	total: 4.6s	remaining: 9.01s
338:	learn: 0.0071651	total: 4.62s	remaining: 9s
339:	learn: 0.0071602	total: 4.63s	remaining: 8.99s
340:	learn: 0.0071519	total: 4.65s	remaining: 8.98s
341:	learn: 0.0071455	total: 4.66s	remaining: 8.97s
342:	learn: 0.0071431	total: 4.67s	remaining: 8.95s
343:	learn: 0.0071150	total: 4.69s	remaining: 8.95s
344:	learn: 0.0071076	total: 4.7s	remaining: 8.93s
345:	learn: 0.0070949	total: 4.71s	remaining: 8.91s
346:	learn: 0.0070879	total: 4.73s	remaining: 8.9s
347:	learn: 0.0070675	total: 4.74s	remaining: 8.88s
348:	learn: 0.0070584	total: 4.75s	remaining: 8.87s
349:	learn: 0.0070539	total: 4.77s	remaining: 8.85s
350:	learn: 0.0070492	total: 4.78s	remaining: 8.84s
351:	learn: 0.0070428	total: 4.79s	remaining: 8.83s
352:	learn: 0.0070359	total: 4.81s	remaining: 8.81s
353:	learn: 0.0070294	total: 4.82s	remaining: 8.79s
354:	learn: 0.0070205	total: 4.84s	remaining: 8.79s
355:	learn: 0.0070182	total: 4.85s	remaining: 8.77s
356:	learn: 0.0069984	total: 4.86s	remaining: 8.76s
357:	learn: 0.0069955	total: 4.87s	remaining: 8.74s
358:	learn: 0.0069871	total: 4.88s	remaining: 8.72s
359:	learn: 0.0069852	total: 4.9s	remaining: 8.71s
360:	learn: 0.0069732	total: 4.91s	remaining: 8.69s
361:	learn: 0.0069688	total: 4.92s	remaining: 8.68s
362:	learn: 0.0069649	total: 4.94s	remaining: 8.66s
363:	learn: 0.0069566	total: 4.95s	remaining: 8.65s
364:	learn: 0.0069529	total: 4.96s	remaining: 8.63s
365:	learn: 0.0069475	total: 4.97s	remaining: 8.62s
366:	learn: 0.0069442	total: 4.99s	remaining: 8.6s
367:	learn: 0.0069227	total: 5s	remaining: 8.59s

368:	learn: 0.0069129	total: 5.01s	remaining: 8.58s
369:	learn: 0.0068947	total: 5.03s	remaining: 8.56s
370:	learn: 0.0068769	total: 5.04s	remaining: 8.54s
371:	learn: 0.0068705	total: 5.05s	remaining: 8.53s
372:	learn: 0.0068624	total: 5.07s	remaining: 8.52s
373:	learn: 0.0068584	total: 5.08s	remaining: 8.5s
374:	learn: 0.0068530	total: 5.09s	remaining: 8.49s
375:	learn: 0.0068508	total: 5.12s	remaining: 8.49s
376:	learn: 0.0068324	total: 5.13s	remaining: 8.47s
377:	learn: 0.0068240	total: 5.14s	remaining: 8.46s
378:	learn: 0.0068066	total: 5.15s	remaining: 8.45s
379:	learn: 0.0068040	total: 5.17s	remaining: 8.43s
380:	learn: 0.0067965	total: 5.18s	remaining: 8.42s
381:	learn: 0.0067800	total: 5.2s	remaining: 8.41s
382:	learn: 0.0067653	total: 5.21s	remaining: 8.39s
383:	learn: 0.0067494	total: 5.22s	remaining: 8.38s
384:	learn: 0.0067429	total: 5.24s	remaining: 8.37s
385:	learn: 0.0067271	total: 5.25s	remaining: 8.36s
386:	learn: 0.0067251	total: 5.26s	remaining: 8.34s
387:	learn: 0.0067224	total: 5.27s	remaining: 8.32s
388:	learn: 0.0067134	total: 5.29s	remaining: 8.3s
389:	learn: 0.0067056	total: 5.3s	remaining: 8.29s
390:	learn: 0.0067017	total: 5.31s	remaining: 8.27s
391:	learn: 0.0066945	total: 5.32s	remaining: 8.25s
392:	learn: 0.0066927	total: 5.33s	remaining: 8.23s
393:	learn: 0.0066781	total: 5.34s	remaining: 8.22s
394:	learn: 0.0066709	total: 5.36s	remaining: 8.21s
395:	learn: 0.0066641	total: 5.37s	remaining: 8.2s
396:	learn: 0.0066568	total: 5.39s	remaining: 8.19s
397:	learn: 0.0066494	total: 5.41s	remaining: 8.18s
398:	learn: 0.0066457	total: 5.42s	remaining: 8.16s
399:	learn: 0.0066387	total: 5.43s	remaining: 8.15s
400:	learn: 0.0066299	total: 5.45s	remaining: 8.13s
401:	learn: 0.0066208	total: 5.46s	remaining: 8.13s
402:	learn: 0.0066129	total: 5.47s	remaining: 8.11s
403:	learn: 0.0066076	total: 5.49s	remaining: 8.1s
404:	learn: 0.0065940	total: 5.5s	remaining: 8.08s
405:	learn: 0.0065810	total: 5.51s	remaining: 8.06s
406:	learn: 0.0065773	total: 5.54s	remaining: 8.07s
407:	learn: 0.0065741	total: 5.55s	remaining: 8.05s
408:	learn: 0.0065662	total: 5.57s	remaining: 8.05s
409:	learn: 0.0065642	total: 5.58s	remaining: 8.03s
410:	learn: 0.0065610	total: 5.6s	remaining: 8.03s
411:	learn: 0.0065538	total: 5.61s	remaining: 8.01s
412:	learn: 0.0065471	total: 5.62s	remaining: 7.99s
413:	learn: 0.0065320	total: 5.65s	remaining: 8s
414:	learn: 0.0065222	total: 5.66s	remaining: 7.98s
415:	learn: 0.0065169	total: 5.68s	remaining: 7.97s

416:	learn: 0.0065105	total: 5.69s	remaining: 7.96s
417:	learn: 0.0065042	total: 5.71s	remaining: 7.94s
418:	learn: 0.0064974	total: 5.72s	remaining: 7.93s
419:	learn: 0.0064897	total: 5.74s	remaining: 7.92s
420:	learn: 0.0064834	total: 5.75s	remaining: 7.91s
421:	learn: 0.0064771	total: 5.76s	remaining: 7.89s
422:	learn: 0.0064643	total: 5.77s	remaining: 7.87s
423:	learn: 0.0064521	total: 5.79s	remaining: 7.87s
424:	learn: 0.0064478	total: 5.8s	remaining: 7.85s
425:	learn: 0.0064343	total: 5.82s	remaining: 7.84s
426:	learn: 0.0064311	total: 5.83s	remaining: 7.83s
427:	learn: 0.0064267	total: 5.85s	remaining: 7.82s
428:	learn: 0.0064164	total: 5.86s	remaining: 7.8s
429:	learn: 0.0064104	total: 5.87s	remaining: 7.79s
430:	learn: 0.0064049	total: 5.89s	remaining: 7.78s
431:	learn: 0.0063990	total: 5.91s	remaining: 7.77s
432:	learn: 0.0063976	total: 5.92s	remaining: 7.75s
433:	learn: 0.0063862	total: 5.93s	remaining: 7.74s
434:	learn: 0.0063820	total: 5.95s	remaining: 7.73s
435:	learn: 0.0063762	total: 5.96s	remaining: 7.72s
436:	learn: 0.0063707	total: 5.98s	remaining: 7.7s
437:	learn: 0.0063644	total: 5.99s	remaining: 7.69s
438:	learn: 0.0063581	total: 6.01s	remaining: 7.68s
439:	learn: 0.0063559	total: 6.02s	remaining: 7.67s
440:	learn: 0.0063492	total: 6.04s	remaining: 7.65s
441:	learn: 0.0063373	total: 6.05s	remaining: 7.64s
442:	learn: 0.0063246	total: 6.07s	remaining: 7.63s
443:	learn: 0.0063218	total: 6.08s	remaining: 7.61s
444:	learn: 0.0062965	total: 6.09s	remaining: 7.6s
445:	learn: 0.0062951	total: 6.1s	remaining: 7.58s
446:	learn: 0.0062912	total: 6.12s	remaining: 7.57s
447:	learn: 0.0062856	total: 6.13s	remaining: 7.55s
448:	learn: 0.0062797	total: 6.14s	remaining: 7.54s
449:	learn: 0.0062670	total: 6.16s	remaining: 7.52s
450:	learn: 0.0062642	total: 6.17s	remaining: 7.51s
451:	learn: 0.0062584	total: 6.18s	remaining: 7.49s
452:	learn: 0.0062510	total: 6.19s	remaining: 7.48s
453:	learn: 0.0062442	total: 6.21s	remaining: 7.46s
454:	learn: 0.0062421	total: 6.23s	remaining: 7.46s
455:	learn: 0.0062396	total: 6.24s	remaining: 7.45s
456:	learn: 0.0062371	total: 6.25s	remaining: 7.43s
457:	learn: 0.0062350	total: 6.27s	remaining: 7.42s
458:	learn: 0.0062330	total: 6.28s	remaining: 7.41s
459:	learn: 0.0062304	total: 6.3s	remaining: 7.39s
460:	learn: 0.0062280	total: 6.31s	remaining: 7.38s
461:	learn: 0.0062198	total: 6.32s	remaining: 7.36s
462:	learn: 0.0062179	total: 6.33s	remaining: 7.34s
463:	learn: 0.0062126	total: 6.34s	remaining: 7.33s

464:	learn: 0.0062107	total: 6.35s	remaining: 7.31s
465:	learn: 0.0062084	total: 6.37s	remaining: 7.3s
466:	learn: 0.0062065	total: 6.38s	remaining: 7.29s
467:	learn: 0.0062047	total: 6.4s	remaining: 7.27s
468:	learn: 0.0061970	total: 6.41s	remaining: 7.26s
469:	learn: 0.0061946	total: 6.42s	remaining: 7.24s
470:	learn: 0.0061928	total: 6.44s	remaining: 7.23s
471:	learn: 0.0061911	total: 6.45s	remaining: 7.21s
472:	learn: 0.0061791	total: 6.46s	remaining: 7.2s
473:	learn: 0.0061658	total: 6.47s	remaining: 7.18s
474:	learn: 0.0061587	total: 6.49s	remaining: 7.17s
475:	learn: 0.0061487	total: 6.5s	remaining: 7.15s
476:	learn: 0.0061412	total: 6.51s	remaining: 7.14s
477:	learn: 0.0061388	total: 6.53s	remaining: 7.13s
478:	learn: 0.0061321	total: 6.54s	remaining: 7.11s
479:	learn: 0.0061264	total: 6.55s	remaining: 7.1s
480:	learn: 0.0061116	total: 6.57s	remaining: 7.09s
481:	learn: 0.0061080	total: 6.58s	remaining: 7.07s
482:	learn: 0.0060966	total: 6.59s	remaining: 7.06s
483:	learn: 0.0060871	total: 6.6s	remaining: 7.04s
484:	learn: 0.0060849	total: 6.62s	remaining: 7.03s
485:	learn: 0.0060825	total: 6.63s	remaining: 7.01s
486:	learn: 0.0060802	total: 6.64s	remaining: 7s
487:	learn: 0.0060701	total: 6.66s	remaining: 6.98s
488:	learn: 0.0060663	total: 6.68s	remaining: 6.98s
489:	learn: 0.0060617	total: 6.69s	remaining: 6.97s
490:	learn: 0.0060565	total: 6.7s	remaining: 6.95s
491:	learn: 0.0060529	total: 6.72s	remaining: 6.94s
492:	learn: 0.0060441	total: 6.73s	remaining: 6.92s
493:	learn: 0.0060359	total: 6.75s	remaining: 6.91s
494:	learn: 0.0060310	total: 6.76s	remaining: 6.9s
495:	learn: 0.0060287	total: 6.77s	remaining: 6.88s
496:	learn: 0.0060272	total: 6.79s	remaining: 6.87s
497:	learn: 0.0060223	total: 6.8s	remaining: 6.86s
498:	learn: 0.0060208	total: 6.81s	remaining: 6.84s
499:	learn: 0.0060145	total: 6.83s	remaining: 6.83s
500:	learn: 0.0060102	total: 6.84s	remaining: 6.82s
501:	learn: 0.0059963	total: 6.85s	remaining: 6.8s
502:	learn: 0.0059940	total: 6.86s	remaining: 6.78s
503:	learn: 0.0059917	total: 6.88s	remaining: 6.78s
504:	learn: 0.0059871	total: 6.89s	remaining: 6.76s
505:	learn: 0.0059849	total: 6.91s	remaining: 6.75s
506:	learn: 0.0059834	total: 6.92s	remaining: 6.73s
507:	learn: 0.0059728	total: 6.93s	remaining: 6.71s
508:	learn: 0.0059687	total: 6.95s	remaining: 6.7s
509:	learn: 0.0059578	total: 6.96s	remaining: 6.68s
510:	learn: 0.0059532	total: 6.97s	remaining: 6.67s
511:	learn: 0.0059510	total: 6.98s	remaining: 6.65s

512:	learn: 0.0059496	total: 7s	remaining: 6.64s
513:	learn: 0.0059443	total: 7.01s	remaining: 6.63s
514:	learn: 0.0059402	total: 7.02s	remaining: 6.61s
515:	learn: 0.0059368	total: 7.03s	remaining: 6.6s
516:	learn: 0.0059346	total: 7.05s	remaining: 6.59s
517:	learn: 0.0059325	total: 7.07s	remaining: 6.57s
518:	learn: 0.0059312	total: 7.08s	remaining: 6.56s
519:	learn: 0.0059215	total: 7.09s	remaining: 6.55s
520:	learn: 0.0059123	total: 7.1s	remaining: 6.53s
521:	learn: 0.0059095	total: 7.12s	remaining: 6.52s
522:	learn: 0.0059025	total: 7.13s	remaining: 6.5s
523:	learn: 0.0058997	total: 7.14s	remaining: 6.48s
524:	learn: 0.0058951	total: 7.15s	remaining: 6.47s
525:	learn: 0.0058846	total: 7.16s	remaining: 6.46s
526:	learn: 0.0058740	total: 7.17s	remaining: 6.44s
527:	learn: 0.0058689	total: 7.18s	remaining: 6.42s
528:	learn: 0.0058546	total: 7.2s	remaining: 6.41s
529:	learn: 0.0058510	total: 7.21s	remaining: 6.4s
530:	learn: 0.0058466	total: 7.22s	remaining: 6.38s
531:	learn: 0.0058416	total: 7.24s	remaining: 6.37s
532:	learn: 0.0058343	total: 7.25s	remaining: 6.35s
533:	learn: 0.0058322	total: 7.26s	remaining: 6.33s
534:	learn: 0.0058302	total: 7.27s	remaining: 6.32s
535:	learn: 0.0058291	total: 7.29s	remaining: 6.31s
536:	learn: 0.0058231	total: 7.3s	remaining: 6.29s
537:	learn: 0.0058186	total: 7.31s	remaining: 6.28s
538:	learn: 0.0058133	total: 7.32s	remaining: 6.26s
539:	learn: 0.0058113	total: 7.33s	remaining: 6.25s
540:	learn: 0.0058081	total: 7.35s	remaining: 6.23s
541:	learn: 0.0058056	total: 7.36s	remaining: 6.22s
542:	learn: 0.0058036	total: 7.37s	remaining: 6.2s
543:	learn: 0.0057967	total: 7.38s	remaining: 6.19s
544:	learn: 0.0057930	total: 7.39s	remaining: 6.17s
545:	learn: 0.0057881	total: 7.41s	remaining: 6.16s
546:	learn: 0.0057810	total: 7.42s	remaining: 6.15s
547:	learn: 0.0057752	total: 7.43s	remaining: 6.13s
548:	learn: 0.0057692	total: 7.44s	remaining: 6.11s
549:	learn: 0.0057645	total: 7.45s	remaining: 6.1s
550:	learn: 0.0057625	total: 7.46s	remaining: 6.08s
551:	learn: 0.0057594	total: 7.48s	remaining: 6.07s
552:	learn: 0.0057574	total: 7.49s	remaining: 6.05s
553:	learn: 0.0057554	total: 7.5s	remaining: 6.04s
554:	learn: 0.0057470	total: 7.51s	remaining: 6.02s
555:	learn: 0.0057443	total: 7.52s	remaining: 6s
556:	learn: 0.0057385	total: 7.53s	remaining: 5.99s
557:	learn: 0.0057338	total: 7.55s	remaining: 5.98s
558:	learn: 0.0057163	total: 7.56s	remaining: 5.96s
559:	learn: 0.0057144	total: 7.57s	remaining: 5.95s

560:	learn: 0.0057053	total: 7.58s	remaining: 5.93s
561:	learn: 0.0057027	total: 7.59s	remaining: 5.92s
562:	learn: 0.0057008	total: 7.61s	remaining: 5.9s
563:	learn: 0.0056984	total: 7.62s	remaining: 5.89s
564:	learn: 0.0056975	total: 7.63s	remaining: 5.87s
565:	learn: 0.0056956	total: 7.64s	remaining: 5.86s
566:	learn: 0.0056920	total: 7.65s	remaining: 5.84s
567:	learn: 0.0056782	total: 7.66s	remaining: 5.83s
568:	learn: 0.0056748	total: 7.67s	remaining: 5.81s
569:	learn: 0.0056651	total: 7.68s	remaining: 5.8s
570:	learn: 0.0056605	total: 7.7s	remaining: 5.78s
571:	learn: 0.0056548	total: 7.71s	remaining: 5.77s
572:	learn: 0.0056494	total: 7.72s	remaining: 5.75s
573:	learn: 0.0056428	total: 7.73s	remaining: 5.74s
574:	learn: 0.0056265	total: 7.74s	remaining: 5.72s
575:	learn: 0.0056223	total: 7.75s	remaining: 5.71s
576:	learn: 0.0056205	total: 7.77s	remaining: 5.69s
577:	learn: 0.0056187	total: 7.78s	remaining: 5.68s
578:	learn: 0.0056169	total: 7.79s	remaining: 5.66s
579:	learn: 0.0056146	total: 7.8s	remaining: 5.65s
580:	learn: 0.0056085	total: 7.81s	remaining: 5.63s
581:	learn: 0.0056031	total: 7.83s	remaining: 5.62s
582:	learn: 0.0056010	total: 7.84s	remaining: 5.6s
583:	learn: 0.0055970	total: 7.85s	remaining: 5.59s
584:	learn: 0.0055918	total: 7.86s	remaining: 5.58s
585:	learn: 0.0055880	total: 7.87s	remaining: 5.56s
586:	learn: 0.0055836	total: 7.88s	remaining: 5.55s
587:	learn: 0.0055798	total: 7.9s	remaining: 5.53s
588:	learn: 0.0055662	total: 7.91s	remaining: 5.52s
589:	learn: 0.0055644	total: 7.93s	remaining: 5.51s
590:	learn: 0.0055622	total: 7.94s	remaining: 5.49s
591:	learn: 0.0055604	total: 7.95s	remaining: 5.48s
592:	learn: 0.0055570	total: 7.96s	remaining: 5.46s
593:	learn: 0.0055553	total: 7.97s	remaining: 5.45s
594:	learn: 0.0055535	total: 7.99s	remaining: 5.44s
595:	learn: 0.0055513	total: 8s	remaining: 5.42s
596:	learn: 0.0055471	total: 8.01s	remaining: 5.41s
597:	learn: 0.0055375	total: 8.02s	remaining: 5.39s
598:	learn: 0.0055358	total: 8.03s	remaining: 5.38s
599:	learn: 0.0055341	total: 8.04s	remaining: 5.36s
600:	learn: 0.0055248	total: 8.05s	remaining: 5.35s
601:	learn: 0.0055170	total: 8.07s	remaining: 5.33s
602:	learn: 0.0055116	total: 8.08s	remaining: 5.32s
603:	learn: 0.0055026	total: 8.09s	remaining: 5.3s
604:	learn: 0.0054975	total: 8.1s	remaining: 5.29s
605:	learn: 0.0054936	total: 8.12s	remaining: 5.28s
606:	learn: 0.0054869	total: 8.13s	remaining: 5.26s
607:	learn: 0.0054836	total: 8.14s	remaining: 5.25s

608:	learn: 0.0054819	total: 8.15s	remaining: 5.23s
609:	learn: 0.0054783	total: 8.16s	remaining: 5.22s
610:	learn: 0.0054767	total: 8.18s	remaining: 5.21s
611:	learn: 0.0054736	total: 8.19s	remaining: 5.19s
612:	learn: 0.0054661	total: 8.2s	remaining: 5.17s
613:	learn: 0.0054645	total: 8.21s	remaining: 5.16s
614:	learn: 0.0054597	total: 8.22s	remaining: 5.14s
615:	learn: 0.0054510	total: 8.23s	remaining: 5.13s
616:	learn: 0.0054464	total: 8.24s	remaining: 5.12s
617:	learn: 0.0054448	total: 8.25s	remaining: 5.1s
618:	learn: 0.0054403	total: 8.27s	remaining: 5.09s
619:	learn: 0.0054387	total: 8.28s	remaining: 5.07s
620:	learn: 0.0054372	total: 8.29s	remaining: 5.06s
621:	learn: 0.0054337	total: 8.3s	remaining: 5.04s
622:	learn: 0.0054290	total: 8.31s	remaining: 5.03s
623:	learn: 0.0054274	total: 8.32s	remaining: 5.01s
624:	learn: 0.0054191	total: 8.33s	remaining: 5s
625:	learn: 0.0054172	total: 8.34s	remaining: 4.98s
626:	learn: 0.0054128	total: 8.36s	remaining: 4.97s
627:	learn: 0.0054113	total: 8.37s	remaining: 4.96s
628:	learn: 0.0054008	total: 8.38s	remaining: 4.94s
629:	learn: 0.0053965	total: 8.39s	remaining: 4.93s
630:	learn: 0.0053923	total: 8.4s	remaining: 4.91s
631:	learn: 0.0053889	total: 8.41s	remaining: 4.9s
632:	learn: 0.0053819	total: 8.42s	remaining: 4.88s
633:	learn: 0.0053778	total: 8.44s	remaining: 4.87s
634:	learn: 0.0053758	total: 8.45s	remaining: 4.86s
635:	learn: 0.0053727	total: 8.46s	remaining: 4.84s
636:	learn: 0.0053643	total: 8.47s	remaining: 4.83s
637:	learn: 0.0053540	total: 8.48s	remaining: 4.81s
638:	learn: 0.0053525	total: 8.5s	remaining: 4.8s
639:	learn: 0.0053510	total: 8.51s	remaining: 4.79s
640:	learn: 0.0053367	total: 8.52s	remaining: 4.77s
641:	learn: 0.0053352	total: 8.53s	remaining: 4.76s
642:	learn: 0.0053273	total: 8.54s	remaining: 4.74s
643:	learn: 0.0053230	total: 8.55s	remaining: 4.73s
644:	learn: 0.0053174	total: 8.56s	remaining: 4.71s
645:	learn: 0.0053046	total: 8.57s	remaining: 4.7s
646:	learn: 0.0053031	total: 8.59s	remaining: 4.68s
647:	learn: 0.0052984	total: 8.6s	remaining: 4.67s
648:	learn: 0.0052970	total: 8.61s	remaining: 4.66s
649:	learn: 0.0052955	total: 8.62s	remaining: 4.64s
650:	learn: 0.0052923	total: 8.63s	remaining: 4.63s
651:	learn: 0.0052884	total: 8.64s	remaining: 4.61s
652:	learn: 0.0052752	total: 8.66s	remaining: 4.6s
653:	learn: 0.0052689	total: 8.67s	remaining: 4.58s
654:	learn: 0.0052659	total: 8.68s	remaining: 4.57s
655:	learn: 0.0052614	total: 8.69s	remaining: 4.56s

656:	learn: 0.0052596	total: 8.7s	remaining: 4.54s
657:	learn: 0.0052551	total: 8.71s	remaining: 4.53s
658:	learn: 0.0052510	total: 8.72s	remaining: 4.51s
659:	learn: 0.0052382	total: 8.73s	remaining: 4.5s
660:	learn: 0.0052325	total: 8.75s	remaining: 4.49s
661:	learn: 0.0052311	total: 8.76s	remaining: 4.47s
662:	learn: 0.0052192	total: 8.77s	remaining: 4.46s
663:	learn: 0.0052178	total: 8.78s	remaining: 4.44s
664:	learn: 0.0052045	total: 8.79s	remaining: 4.43s
665:	learn: 0.0051997	total: 8.8s	remaining: 4.42s
666:	learn: 0.0051959	total: 8.82s	remaining: 4.4s
667:	learn: 0.0051922	total: 8.83s	remaining: 4.39s
668:	learn: 0.0051909	total: 8.84s	remaining: 4.37s
669:	learn: 0.0051892	total: 8.85s	remaining: 4.36s
670:	learn: 0.0051876	total: 8.86s	remaining: 4.34s
671:	learn: 0.0051822	total: 8.87s	remaining: 4.33s
672:	learn: 0.0051695	total: 8.88s	remaining: 4.32s
673:	learn: 0.0051659	total: 8.9s	remaining: 4.3s
674:	learn: 0.0051645	total: 8.91s	remaining: 4.29s
675:	learn: 0.0051609	total: 8.92s	remaining: 4.27s
676:	learn: 0.0051547	total: 8.93s	remaining: 4.26s
677:	learn: 0.0051534	total: 8.94s	remaining: 4.25s
678:	learn: 0.0051486	total: 8.95s	remaining: 4.23s
679:	learn: 0.0051473	total: 8.97s	remaining: 4.22s
680:	learn: 0.0051431	total: 8.98s	remaining: 4.21s
681:	learn: 0.0051410	total: 8.99s	remaining: 4.19s
682:	learn: 0.0051372	total: 9s	remaining: 4.18s
683:	learn: 0.0051313	total: 9.01s	remaining: 4.16s
684:	learn: 0.0051278	total: 9.02s	remaining: 4.15s
685:	learn: 0.0051241	total: 9.03s	remaining: 4.13s
686:	learn: 0.0051208	total: 9.04s	remaining: 4.12s
687:	learn: 0.0051196	total: 9.06s	remaining: 4.11s
688:	learn: 0.0051155	total: 9.07s	remaining: 4.09s
689:	learn: 0.0051035	total: 9.08s	remaining: 4.08s
690:	learn: 0.0050971	total: 9.09s	remaining: 4.06s
691:	learn: 0.0050959	total: 9.1s	remaining: 4.05s
692:	learn: 0.0050945	total: 9.11s	remaining: 4.04s
693:	learn: 0.0050830	total: 9.12s	remaining: 4.02s
694:	learn: 0.0050811	total: 9.13s	remaining: 4.01s
695:	learn: 0.0050796	total: 9.14s	remaining: 3.99s
696:	learn: 0.0050778	total: 9.16s	remaining: 3.98s
697:	learn: 0.0050764	total: 9.17s	remaining: 3.97s
698:	learn: 0.0050751	total: 9.18s	remaining: 3.95s
699:	learn: 0.0050738	total: 9.19s	remaining: 3.94s
700:	learn: 0.0050704	total: 9.2s	remaining: 3.92s
701:	learn: 0.0050691	total: 9.21s	remaining: 3.91s
702:	learn: 0.0050679	total: 9.22s	remaining: 3.9s
703:	learn: 0.0050667	total: 9.23s	remaining: 3.88s

704:	learn: 0.0050657	total: 9.25s	remaining: 3.87s
705:	learn: 0.0050641	total: 9.26s	remaining: 3.85s
706:	learn: 0.0050580	total: 9.27s	remaining: 3.84s
707:	learn: 0.0050570	total: 9.28s	remaining: 3.83s
708:	learn: 0.0050556	total: 9.29s	remaining: 3.81s
709:	learn: 0.0050513	total: 9.3s	remaining: 3.8s
710:	learn: 0.0050418	total: 9.31s	remaining: 3.78s
711:	learn: 0.0050366	total: 9.32s	remaining: 3.77s
712:	learn: 0.0050324	total: 9.33s	remaining: 3.76s
713:	learn: 0.0050307	total: 9.35s	remaining: 3.74s
714:	learn: 0.0050275	total: 9.36s	remaining: 3.73s
715:	learn: 0.0050222	total: 9.37s	remaining: 3.72s
716:	learn: 0.0050181	total: 9.38s	remaining: 3.7s
717:	learn: 0.0050074	total: 9.39s	remaining: 3.69s
718:	learn: 0.0049968	total: 9.4s	remaining: 3.67s
719:	learn: 0.0049956	total: 9.41s	remaining: 3.66s
720:	learn: 0.0049855	total: 9.43s	remaining: 3.65s
721:	learn: 0.0049837	total: 9.44s	remaining: 3.63s
722:	learn: 0.0049782	total: 9.45s	remaining: 3.62s
723:	learn: 0.0049749	total: 9.46s	remaining: 3.61s
724:	learn: 0.0049732	total: 9.47s	remaining: 3.59s
725:	learn: 0.0049630	total: 9.48s	remaining: 3.58s
726:	learn: 0.0049618	total: 9.5s	remaining: 3.57s
727:	learn: 0.0049609	total: 9.51s	remaining: 3.55s
728:	learn: 0.0049597	total: 9.53s	remaining: 3.54s
729:	learn: 0.0049565	total: 9.54s	remaining: 3.53s
730:	learn: 0.0049512	total: 9.55s	remaining: 3.51s
731:	learn: 0.0049495	total: 9.56s	remaining: 3.5s
732:	learn: 0.0049486	total: 9.57s	remaining: 3.49s
733:	learn: 0.0049474	total: 9.58s	remaining: 3.47s
734:	learn: 0.0049466	total: 9.6s	remaining: 3.46s
735:	learn: 0.0049455	total: 9.61s	remaining: 3.45s
736:	learn: 0.0049446	total: 9.62s	remaining: 3.43s
737:	learn: 0.0049438	total: 9.63s	remaining: 3.42s
738:	learn: 0.0049398	total: 9.64s	remaining: 3.4s
739:	learn: 0.0049366	total: 9.65s	remaining: 3.39s
740:	learn: 0.0049326	total: 9.66s	remaining: 3.38s
741:	learn: 0.0049275	total: 9.68s	remaining: 3.36s
742:	learn: 0.0049177	total: 9.69s	remaining: 3.35s
743:	learn: 0.0049166	total: 9.71s	remaining: 3.34s
744:	learn: 0.0049139	total: 9.72s	remaining: 3.33s
745:	learn: 0.0049112	total: 9.73s	remaining: 3.31s
746:	learn: 0.0049082	total: 9.74s	remaining: 3.3s
747:	learn: 0.0049030	total: 9.75s	remaining: 3.29s
748:	learn: 0.0048951	total: 9.76s	remaining: 3.27s
749:	learn: 0.0048859	total: 9.78s	remaining: 3.26s
750:	learn: 0.0048816	total: 9.79s	remaining: 3.24s
751:	learn: 0.0048765	total: 9.8s	remaining: 3.23s

752:	learn: 0.0048675	total: 9.81s	remaining: 3.22s
753:	learn: 0.0048659	total: 9.82s	remaining: 3.2s
754:	learn: 0.0048648	total: 9.83s	remaining: 3.19s
755:	learn: 0.0048609	total: 9.84s	remaining: 3.18s
756:	learn: 0.0048562	total: 9.85s	remaining: 3.16s
757:	learn: 0.0048551	total: 9.87s	remaining: 3.15s
758:	learn: 0.0048448	total: 9.88s	remaining: 3.14s
759:	learn: 0.0048409	total: 9.89s	remaining: 3.12s
760:	learn: 0.0048399	total: 9.91s	remaining: 3.11s
761:	learn: 0.0048341	total: 9.92s	remaining: 3.1s
762:	learn: 0.0048325	total: 9.93s	remaining: 3.08s
763:	learn: 0.0048239	total: 9.94s	remaining: 3.07s
764:	learn: 0.0048223	total: 9.96s	remaining: 3.06s
765:	learn: 0.0048167	total: 9.97s	remaining: 3.04s
766:	learn: 0.0048129	total: 9.98s	remaining: 3.03s
767:	learn: 0.0048033	total: 10s	remaining: 3.02s
768:	learn: 0.0048017	total: 10s	remaining: 3.01s
769:	learn: 0.0047980	total: 10s	remaining: 2.99s
770:	learn: 0.0047950	total: 10s	remaining: 2.98s
771:	learn: 0.0047930	total: 10s	remaining: 2.97s
772:	learn: 0.0047899	total: 10.1s	remaining: 2.95s
773:	learn: 0.0047843	total: 10.1s	remaining: 2.94s
774:	learn: 0.0047816	total: 10.1s	remaining: 2.93s
775:	learn: 0.0047762	total: 10.1s	remaining: 2.91s
776:	learn: 0.0047698	total: 10.1s	remaining: 2.9s
777:	learn: 0.0047669	total: 10.1s	remaining: 2.89s
778:	learn: 0.0047633	total: 10.1s	remaining: 2.88s
779:	learn: 0.0047597	total: 10.2s	remaining: 2.86s
780:	learn: 0.0047568	total: 10.2s	remaining: 2.85s
781:	learn: 0.0047541	total: 10.2s	remaining: 2.84s
782:	learn: 0.0047473	total: 10.2s	remaining: 2.83s
783:	learn: 0.0047454	total: 10.2s	remaining: 2.81s
784:	learn: 0.0047416	total: 10.2s	remaining: 2.8s
785:	learn: 0.0047390	total: 10.2s	remaining: 2.79s
786:	learn: 0.0047381	total: 10.2s	remaining: 2.77s
787:	learn: 0.0047264	total: 10.3s	remaining: 2.76s
788:	learn: 0.0047256	total: 10.3s	remaining: 2.75s
789:	learn: 0.0047246	total: 10.3s	remaining: 2.73s
790:	learn: 0.0047238	total: 10.3s	remaining: 2.72s
791:	learn: 0.0047173	total: 10.3s	remaining: 2.71s
792:	learn: 0.0047119	total: 10.3s	remaining: 2.69s
793:	learn: 0.0047098	total: 10.3s	remaining: 2.68s
794:	learn: 0.0047088	total: 10.3s	remaining: 2.67s
795:	learn: 0.0047080	total: 10.4s	remaining: 2.65s
796:	learn: 0.0047072	total: 10.4s	remaining: 2.64s
797:	learn: 0.0047062	total: 10.4s	remaining: 2.63s
798:	learn: 0.0047055	total: 10.4s	remaining: 2.61s
799:	learn: 0.0047008	total: 10.4s	remaining: 2.6s

800:	learn: 0.0046922	total: 10.4s	remaining: 2.59s
801:	learn: 0.0046820	total: 10.4s	remaining: 2.58s
802:	learn: 0.0046740	total: 10.4s	remaining: 2.56s
803:	learn: 0.0046730	total: 10.5s	remaining: 2.55s
804:	learn: 0.0046709	total: 10.5s	remaining: 2.54s
805:	learn: 0.0046689	total: 10.5s	remaining: 2.52s
806:	learn: 0.0046653	total: 10.5s	remaining: 2.51s
807:	learn: 0.0046479	total: 10.5s	remaining: 2.5s
808:	learn: 0.0046408	total: 10.5s	remaining: 2.48s
809:	learn: 0.0046389	total: 10.5s	remaining: 2.47s
810:	learn: 0.0046338	total: 10.5s	remaining: 2.46s
811:	learn: 0.0046320	total: 10.6s	remaining: 2.44s
812:	learn: 0.0046250	total: 10.6s	remaining: 2.43s
813:	learn: 0.0046236	total: 10.6s	remaining: 2.42s
814:	learn: 0.0046183	total: 10.6s	remaining: 2.4s
815:	learn: 0.0046158	total: 10.6s	remaining: 2.39s
816:	learn: 0.0046132	total: 10.6s	remaining: 2.38s
817:	learn: 0.0046112	total: 10.6s	remaining: 2.37s
818:	learn: 0.0046102	total: 10.6s	remaining: 2.35s
819:	learn: 0.0046082	total: 10.7s	remaining: 2.34s
820:	learn: 0.0046061	total: 10.7s	remaining: 2.33s
821:	learn: 0.0046041	total: 10.7s	remaining: 2.31s
822:	learn: 0.0046021	total: 10.7s	remaining: 2.3s
823:	learn: 0.0045957	total: 10.7s	remaining: 2.29s
824:	learn: 0.0045932	total: 10.7s	remaining: 2.27s
825:	learn: 0.0045844	total: 10.7s	remaining: 2.26s
826:	learn: 0.0045747	total: 10.7s	remaining: 2.25s
827:	learn: 0.0045712	total: 10.7s	remaining: 2.23s
828:	learn: 0.0045693	total: 10.8s	remaining: 2.22s
829:	learn: 0.0045620	total: 10.8s	remaining: 2.21s
830:	learn: 0.0045598	total: 10.8s	remaining: 2.19s
831:	learn: 0.0045547	total: 10.8s	remaining: 2.18s
832:	learn: 0.0045528	total: 10.8s	remaining: 2.17s
833:	learn: 0.0045466	total: 10.8s	remaining: 2.15s
834:	learn: 0.0045446	total: 10.8s	remaining: 2.14s
835:	learn: 0.0045427	total: 10.8s	remaining: 2.13s
836:	learn: 0.0045403	total: 10.9s	remaining: 2.11s
837:	learn: 0.0045289	total: 10.9s	remaining: 2.1s
838:	learn: 0.0045243	total: 10.9s	remaining: 2.09s
839:	learn: 0.0045109	total: 10.9s	remaining: 2.07s
840:	learn: 0.0045073	total: 10.9s	remaining: 2.06s
841:	learn: 0.0045030	total: 10.9s	remaining: 2.05s
842:	learn: 0.0045021	total: 10.9s	remaining: 2.04s
843:	learn: 0.0045012	total: 10.9s	remaining: 2.02s
844:	learn: 0.0044950	total: 11s	remaining: 2.01s
845:	learn: 0.0044932	total: 11s	remaining: 2s
846:	learn: 0.0044913	total: 11s	remaining: 1.99s
847:	learn: 0.0044895	total: 11s	remaining: 1.97s

848:	learn: 0.0044877	total: 11s	remaining: 1.96s
849:	learn: 0.0044869	total: 11s	remaining: 1.95s
850:	learn: 0.0044850	total: 11s	remaining: 1.93s
851:	learn: 0.0044841	total: 11.1s	remaining: 1.92s
852:	learn: 0.0044834	total: 11.1s	remaining: 1.91s
853:	learn: 0.0044709	total: 11.1s	remaining: 1.89s
854:	learn: 0.0044614	total: 11.1s	remaining: 1.88s
855:	learn: 0.0044592	total: 11.1s	remaining: 1.87s
856:	learn: 0.0044531	total: 11.1s	remaining: 1.85s
857:	learn: 0.0044513	total: 11.1s	remaining: 1.84s
858:	learn: 0.0044495	total: 11.1s	remaining: 1.83s
859:	learn: 0.0044478	total: 11.2s	remaining: 1.82s
860:	learn: 0.0044456	total: 11.2s	remaining: 1.8s
861:	learn: 0.0044434	total: 11.2s	remaining: 1.79s
862:	learn: 0.0044417	total: 11.2s	remaining: 1.78s
863:	learn: 0.0044396	total: 11.2s	remaining: 1.76s
864:	learn: 0.0044368	total: 11.2s	remaining: 1.75s
865:	learn: 0.0044350	total: 11.2s	remaining: 1.74s
866:	learn: 0.0044333	total: 11.3s	remaining: 1.73s
867:	learn: 0.0044231	total: 11.3s	remaining: 1.71s
868:	learn: 0.0044209	total: 11.3s	remaining: 1.7s
869:	learn: 0.0044152	total: 11.3s	remaining: 1.69s
870:	learn: 0.0044132	total: 11.3s	remaining: 1.67s
871:	learn: 0.0044055	total: 11.3s	remaining: 1.66s
872:	learn: 0.0044035	total: 11.3s	remaining: 1.65s
873:	learn: 0.0044028	total: 11.3s	remaining: 1.63s
874:	learn: 0.0044015	total: 11.4s	remaining: 1.62s
875:	learn: 0.0043952	total: 11.4s	remaining: 1.61s
876:	learn: 0.0043925	total: 11.4s	remaining: 1.59s
877:	learn: 0.0043879	total: 11.4s	remaining: 1.58s
878:	learn: 0.0043804	total: 11.4s	remaining: 1.57s
879:	learn: 0.0043788	total: 11.4s	remaining: 1.56s
880:	learn: 0.0043773	total: 11.4s	remaining: 1.54s
881:	learn: 0.0043725	total: 11.4s	remaining: 1.53s
882:	learn: 0.0043695	total: 11.5s	remaining: 1.52s
883:	learn: 0.0043606	total: 11.5s	remaining: 1.5s
884:	learn: 0.0043535	total: 11.5s	remaining: 1.49s
885:	learn: 0.0043467	total: 11.5s	remaining: 1.48s
886:	learn: 0.0043379	total: 11.5s	remaining: 1.47s
887:	learn: 0.0043362	total: 11.5s	remaining: 1.45s
888:	learn: 0.0043305	total: 11.5s	remaining: 1.44s
889:	learn: 0.0043241	total: 11.6s	remaining: 1.43s
890:	learn: 0.0043223	total: 11.6s	remaining: 1.42s
891:	learn: 0.0043157	total: 11.6s	remaining: 1.4s
892:	learn: 0.0043130	total: 11.6s	remaining: 1.39s
893:	learn: 0.0043090	total: 11.6s	remaining: 1.38s
894:	learn: 0.0043070	total: 11.6s	remaining: 1.36s
895:	learn: 0.0043053	total: 11.6s	remaining: 1.35s

896:	learn: 0.0043010	total: 11.6s	remaining: 1.34s
897:	learn: 0.0042984	total: 11.7s	remaining: 1.32s
898:	learn: 0.0042964	total: 11.7s	remaining: 1.31s
899:	learn: 0.0042937	total: 11.7s	remaining: 1.3s
900:	learn: 0.0042922	total: 11.7s	remaining: 1.28s
901:	learn: 0.0042897	total: 11.7s	remaining: 1.27s
902:	learn: 0.0042871	total: 11.7s	remaining: 1.26s
903:	learn: 0.0042845	total: 11.7s	remaining: 1.25s
904:	learn: 0.0042822	total: 11.7s	remaining: 1.23s
905:	learn: 0.0042751	total: 11.8s	remaining: 1.22s
906:	learn: 0.0042629	total: 11.8s	remaining: 1.21s
907:	learn: 0.0042587	total: 11.8s	remaining: 1.19s
908:	learn: 0.0042560	total: 11.8s	remaining: 1.18s
909:	learn: 0.0042536	total: 11.8s	remaining: 1.17s
910:	learn: 0.0042513	total: 11.8s	remaining: 1.16s
911:	learn: 0.0042490	total: 11.8s	remaining: 1.14s
912:	learn: 0.0042468	total: 11.8s	remaining: 1.13s
913:	learn: 0.0042446	total: 11.9s	remaining: 1.12s
914:	learn: 0.0042424	total: 11.9s	remaining: 1.1s
915:	learn: 0.0042403	total: 11.9s	remaining: 1.09s
916:	learn: 0.0042381	total: 11.9s	remaining: 1.08s
917:	learn: 0.0042360	total: 11.9s	remaining: 1.06s
918:	learn: 0.0042320	total: 11.9s	remaining: 1.05s
919:	learn: 0.0042299	total: 12s	remaining: 1.04s
920:	learn: 0.0042266	total: 12s	remaining: 1.03s
921:	learn: 0.0042247	total: 12s	remaining: 1.01s
922:	learn: 0.0042227	total: 12s	remaining: 1s
923:	learn: 0.0042206	total: 12s	remaining: 988ms
924:	learn: 0.0042179	total: 12s	remaining: 974ms
925:	learn: 0.0042159	total: 12s	remaining: 962ms
926:	learn: 0.0042139	total: 12.1s	remaining: 949ms
927:	learn: 0.0042120	total: 12.1s	remaining: 936ms
928:	learn: 0.0042102	total: 12.1s	remaining: 923ms
929:	learn: 0.0042083	total: 12.1s	remaining: 910ms
930:	learn: 0.0042064	total: 12.1s	remaining: 897ms
931:	learn: 0.0042042	total: 12.1s	remaining: 884ms
932:	learn: 0.0042023	total: 12.1s	remaining: 871ms
933:	learn: 0.0042005	total: 12.1s	remaining: 858ms
934:	learn: 0.0041978	total: 12.2s	remaining: 845ms
935:	learn: 0.0041939	total: 12.2s	remaining: 832ms
936:	learn: 0.0041921	total: 12.2s	remaining: 819ms
937:	learn: 0.0041904	total: 12.2s	remaining: 806ms
938:	learn: 0.0041887	total: 12.2s	remaining: 793ms
939:	learn: 0.0041866	total: 12.2s	remaining: 780ms
940:	learn: 0.0041849	total: 12.2s	remaining: 767ms
941:	learn: 0.0041822	total: 12.3s	remaining: 754ms
942:	learn: 0.0041817	total: 12.3s	remaining: 741ms
943:	learn: 0.0041800	total: 12.3s	remaining: 728ms

944:	learn: 0.0041784	total: 12.3s	remaining: 715ms
945:	learn: 0.0041768	total: 12.3s	remaining: 702ms
946:	learn: 0.0041727	total: 12.3s	remaining: 689ms
947:	learn: 0.0041708	total: 12.3s	remaining: 676ms
948:	learn: 0.0041688	total: 12.3s	remaining: 663ms
949:	learn: 0.0041672	total: 12.4s	remaining: 651ms
950:	learn: 0.0041654	total: 12.4s	remaining: 638ms
951:	learn: 0.0041639	total: 12.4s	remaining: 625ms
952:	learn: 0.0041612	total: 12.4s	remaining: 612ms
953:	learn: 0.0041601	total: 12.4s	remaining: 599ms
954:	learn: 0.0041505	total: 12.4s	remaining: 586ms
955:	learn: 0.0041386	total: 12.4s	remaining: 573ms
956:	learn: 0.0041377	total: 12.5s	remaining: 560ms
957:	learn: 0.0041362	total: 12.5s	remaining: 547ms
958:	learn: 0.0041335	total: 12.5s	remaining: 534ms
959:	learn: 0.0041300	total: 12.5s	remaining: 520ms
960:	learn: 0.0041278	total: 12.5s	remaining: 507ms
961:	learn: 0.0041264	total: 12.5s	remaining: 494ms
962:	learn: 0.0041249	total: 12.5s	remaining: 481ms
963:	learn: 0.0041235	total: 12.5s	remaining: 468ms
964:	learn: 0.0041218	total: 12.6s	remaining: 455ms
965:	learn: 0.0041214	total: 12.6s	remaining: 442ms
966:	learn: 0.0041176	total: 12.6s	remaining: 429ms
967:	learn: 0.0041139	total: 12.6s	remaining: 416ms
968:	learn: 0.0041126	total: 12.6s	remaining: 404ms
969:	learn: 0.0041112	total: 12.6s	remaining: 391ms
970:	learn: 0.0041108	total: 12.6s	remaining: 377ms
971:	learn: 0.0041095	total: 12.7s	remaining: 365ms
972:	learn: 0.0041081	total: 12.7s	remaining: 352ms
973:	learn: 0.0041062	total: 12.7s	remaining: 338ms
974:	learn: 0.0041059	total: 12.7s	remaining: 326ms
975:	learn: 0.0041046	total: 12.7s	remaining: 313ms
976:	learn: 0.0041033	total: 12.7s	remaining: 300ms
977:	learn: 0.0041029	total: 12.7s	remaining: 286ms
978:	learn: 0.0041007	total: 12.7s	remaining: 273ms
979:	learn: 0.0040994	total: 12.8s	remaining: 260ms
980:	learn: 0.0040973	total: 12.8s	remaining: 247ms
981:	learn: 0.0040947	total: 12.8s	remaining: 234ms
982:	learn: 0.0040883	total: 12.8s	remaining: 221ms
983:	learn: 0.0040848	total: 12.8s	remaining: 208ms
984:	learn: 0.0040806	total: 12.8s	remaining: 195ms
985:	learn: 0.0040784	total: 12.8s	remaining: 182ms
986:	learn: 0.0040758	total: 12.8s	remaining: 169ms
987:	learn: 0.0040734	total: 12.9s	remaining: 156ms
988:	learn: 0.0040680	total: 12.9s	remaining: 143ms
989:	learn: 0.0040673	total: 12.9s	remaining: 130ms
990:	learn: 0.0040585	total: 12.9s	remaining: 117ms
991:	learn: 0.0040562	total: 12.9s	remaining: 104ms

```

992:   learn: 0.0040549      total: 12.9s   remaining: 91.1ms
993:   learn: 0.0040536      total: 12.9s   remaining: 78.1ms
994:   learn: 0.0040523      total: 13s     remaining: 65.2ms
995:   learn: 0.0040511      total: 13s     remaining: 52.2ms
996:   learn: 0.0040485      total: 13s     remaining: 39.1ms
997:   learn: 0.0040462      total: 13s     remaining: 26.1ms
998:   learn: 0.0040319      total: 13s     remaining: 13ms
999:   learn: 0.0040296      total: 13.1s   remaining: 0us

```

```

[308]: GridSearchCV(cv=5, error_score=nan,
                  estimator=<catboost.core.CatBoostClassifier object at
                  0x00000025775ED1CC8>,
                  iid='deprecated', n_jobs=-1,
                  param_grid={'depth': range(2, 10, 2),
                              'l2_leaf_reg': range(2, 10, 2),
                              'learning_rate': [0.03, 0.1]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)

```

```

[313]: print(f'Best Mean Cross Validation Score is {catb_grid.best_params_}')
      print(f'Best Mean Cross Validation Score is {catb_grid.best_score_}')
      print(f'Train score is {catb_grid.score(X_train,y_train)}')

```

```

Best Mean Cross Validation Score is {'depth': 4, 'l2_leaf_reg': 6,
'learning_rate': 0.03}
Best Mean Cross Validation Score is 0.8564556204653645
Train score is 0.9181969949916529

```

4.3.3 LightGBM Classifier - .79816

```

[310]: #pip install lightgbm

```

```

[248]: from lightgbm import LGBMClassifier

```

```

[311]: lgbm = LGBMClassifier()
      param_lgbm = {
          'n_estimators': range(50,70,10),
          'colsample_bytree': [0.7, 0.8],
          'max_depth': range(2,10,5),
          'num_leaves': range(2,10,2),
          'reg_alpha': [1.1, 1.2, 1.3],
          'reg_lambda': [1.1, 1.2, 1.3],
          'min_split_gain': [0.3, 0.4],
          'subsample': [0.7, 0.8, 0.9],
          'subsample_freq': [20]
      }

```

```
grid_lgbm = GridSearchCV(lgbm, param_lgbm, cv=5, n_jobs=-1,
    ↪return_train_score=True, scoring = f2score)
grid_lgbm.fit(X_train, y_train)
```

```
[311]: GridSearchCV(cv=5, error_score=nan,
    estimator=LGBMClassifier(boosting_type='gbdt', class_weight=None,
    colsample_bytree=1.0,
    importance_type='split',
    learning_rate=0.1, max_depth=-1,
    min_child_samples=20,
    min_child_weight=0.001,
    min_split_gain=0.0, n_estimators=100,
    n_jobs=-1, num_leaves=31, objective=None,
    random_state=None, reg_alpha=0.0,
    reg_lambda=0.0, silent=True,
    param_grid={'colsample_bytree': [0.7, 0.8],
    'max_depth': range(2, 10, 5),
    'min_split_gain': [0.3, 0.4],
    'n_estimators': range(50, 70, 10),
    'num_leaves': range(2, 10, 2),
    'reg_alpha': [1.1, 1.2, 1.3],
    'reg_lambda': [1.1, 1.2, 1.3],
    'subsample': [0.7, 0.8, 0.9], 'subsample_freq': [20]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
    scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

```
[312]: print(f'Best Mean Cross Validation Score is {grid_lgbm.best_params_}')
print(f'Best Mean Cross Validation Score is {grid_lgbm.best_score_}')
print(f'Train score is {grid_lgbm.score(X_train, y_train)}')
```

```
Best Mean Cross Validation Score is {'colsample_bytree': 0.7, 'max_depth': 7,
'min_split_gain': 0.4, 'n_estimators': 60, 'num_leaves': 8, 'reg_alpha': 1.3,
'reg_lambda': 1.3, 'subsample': 0.9, 'subsample_freq': 20}
Best Mean Cross Validation Score is 0.8398943998871164
Train score is 0.8912133891213391
```

4.3.4 One Class Classification for imbalanced data (Isolation forest)

```
[253]: from sklearn.ensemble import IsolationForest
```

```
[273]: f2score2 = make_scorer(fbeta_score, beta=2, average='macro') #[None, 'micro',
    ↪'macro', 'weighted']
```

```
[274]: isoforest = IsolationForest(random_state=42)
iso_param = {
    'n_estimators': range(10, 60, 10),
    'max_samples': ['auto'],
```

```

        'contamination':['auto']
    }
    grid_iso = GridSearchCV(isoforest, iso_param,cv=5,n_jobs=-1,
        ↪return_train_score=True, scoring =f2score2)
    grid_iso.fit(X_train,y_train)

```

```

[274]: GridSearchCV(cv=5, error_score=nan,
                estimator=IsolationForest(behaviour='deprecated', bootstrap=False,
                                           contamination='auto', max_features=1.0,
                                           max_samples='auto', n_estimators=100,
                                           n_jobs=None, random_state=42, verbose=0,
                                           warm_start=False),
                iid='deprecated', n_jobs=-1,
                param_grid={'contamination': ['auto'], 'max_samples': ['auto'],
                            'n_estimators': range(10, 60, 10)},
                pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                scoring=make_scorer(fbeta_score, beta=2, average=macro),
                verbose=0)

```

```

[275]: print(f'Best Mean Cross Validation Score is {grid_iso.best_params_}')
        print(f'Best Mean Cross Validation Score is {grid_iso.best_score_}')
        print(f'Train score is {grid_iso.score(X_train,y_train)}')

```

Best Mean Cross Validation Score is {'contamination': 'auto', 'max_samples': 'auto', 'n_estimators': 30}
 Best Mean Cross Validation Score is 0.003558269018006486
 Train score is 0.0031987205117952823

4.3.5 Naive Bayes

```

[284]: from sklearn.naive_bayes import GaussianNB
        from sklearn.naive_bayes import MultinomialNB
        from sklearn.naive_bayes import BernoulliNB

```

GaussianNB - .79670

```

[304]: naiveGaussian = GaussianNB()
        naive_param = {}
        grid_naiveGB = GridSearchCV(naivebayes, naive_param,cv=5,n_jobs=-1,
        ↪return_train_score=True, scoring =f2score2)
        grid_naiveGB.fit(X_train,y_train)

```

```

[304]: GridSearchCV(cv=5, error_score=nan,
                estimator=BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None,
                                       fit_prior=True),
                iid='deprecated', n_jobs=-1, param_grid={},
                pre_dispatch='2*n_jobs', refit=True, return_train_score=True,

```

```
scoring=make_scorer(fbeta_score, beta=2, average=macro),
verbose=0)
```

```
[305]: print(f'Best Mean Cross Validation Score is {grid_naiveGB.best_params_}')
print(f'Best Mean Cross Validation Score is {grid_naiveGB.best_score_}')
print(f'Train score is {grid_naiveGB.score(X_train,y_train)}')
```

```
Best Mean Cross Validation Score is {}
Best Mean Cross Validation Score is 0.8851045876061363
Train score is 0.8873985829077566
```

BernoulliNB - .79670

```
[318]: naivebernoulli = BernoulliNB()
naive_param = {}
grid_naiveber = GridSearchCV(naivebernoulli, naive_param, cv=5, n_jobs=-1,
↪return_train_score=True, scoring = f2score2)
grid_naiveber.fit(X_train, y_train)
```

```
[318]: GridSearchCV(cv=5, error_score=nan,
estimator=BernoulliNB(alpha=1.0, binarize=0.0, class_prior=None,
fit_prior=True),
iid='deprecated', n_jobs=-1, param_grid={},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=make_scorer(fbeta_score, beta=2, average=macro),
verbose=0)
```

```
[320]: print(f'Best Mean Cross Validation Score is {grid_naiveber.best_params_}')
print(f'Best Mean Cross Validation Score is {grid_naiveber.best_score_}')
print(f'Train score is {grid_naiveber.score(X_train,y_train)}')
```

```
Best Mean Cross Validation Score is {}
Best Mean Cross Validation Score is 0.8851045876061363
Train score is 0.8873985829077566
```

5 Saving and Restoring Model

```
[377]: import pickle
```

```
[378]: # save the model to disk
filename = 'XGBoost_Cost_Sensitive.sav'
pickle.dump(grid_xgboost, open(filename, 'wb'))
```

```
[379]: # load the model from disk
loaded_model = pickle.load(open(filename, 'rb'))
```

```
[380]: loaded_model
```

```
[380]: GridSearchCV(cv=5, error_score=nan,
                  estimator=XGBClassifier(base_score=None, booster=None,
                                          colsample_bylevel=None,
                                          colsample_bynode=None,
                                          colsample_bytree=None,
                                          early_stopping_rounds=5, gamma=None,
                                          gpu_id=None, importance_type='gain',
                                          interaction_constraints=None,
                                          learning_rate=None, max_delta_step=None,
                                          max_depth=None, min_child_weight=None,
                                          missing=nan, monotone_...
                                          subsample=None, tree_method=None,
                                          validate_parameters=False,
                                          verbosity=None),
                  iid='deprecated', n_jobs=-1,
                  param_grid={'class_weight': ['balanced', 'balanced_subsample'],
                              'learning_rate': [0.05, 0.01, 0.1],
                              'n_estimators': range(1000, 1200, 20),
                              'scale_pos_weight': range(1000, 1200, 50)},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=make_scorer(fbeta_score, beta=2), verbose=0)
```

6 The End