

HW3_RegressionV2

April 20, 2020

1 Final Results

- Best Model based on CV scores - Ensemble: XGBoost
- Best Mean Cross Validation Score is 0.8984774624760332
- Best parameters {'learning_rate': 0.1, 'max_depth': 4, 'min_child_weight': 1, 'n_estimators': 192, 'subsample': 0.8}
- Train score is 0.9841860823241139
- Test score is 0.8767229954557099

- Best Model based on Stacking: Voting Top 5
- Best Mean Cross Validation Score is 0.8999858356957512
- Train score is 0.9491284183151113
- Test score is 0.8802582634262037

1.1 Data PreProcessing

```
[1]: from math import sqrt
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from sklearn.metrics import mean_squared_error, r2_score
from math import sqrt
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error, r2_score
from math import sqrt

pd.pandas.set_option('display.max_columns', None)
%matplotlib inline
```

1.1.1 Load Datasets

```
[2]: # load dataset
# your code here
data = pd.read_csv(r'C:\Users\nabhs\OneDrive\BUAN - Semester 2\BUAN 6341 - Applied Machine Learning\Datasets\houseprice.csv')
```

```
[3]: data.head()
```

```
[3]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

		LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	
0		Lvl	AllPub	Inside	Gtl	CollgCr	Norm	
1		Lvl	AllPub	FR2	Gtl	Veenker	Feedr	
2		Lvl	AllPub	Inside	Gtl	CollgCr	Norm	
3		Lvl	AllPub	Corner	Gtl	Crawfor	Norm	
4		Lvl	AllPub	FR2	Gtl	NoRidge	Norm	

		Condition2	BldgType	HouseStyle	OverallQual	OverallCond	YearBuilt	
0		Norm	1Fam	2Story	7	5	2003	
1		Norm	1Fam	1Story	6	8	1976	
2		Norm	1Fam	2Story	7	5	2001	
3		Norm	1Fam	2Story	7	5	1915	
4		Norm	1Fam	2Story	8	5	2000	

		YearRemodAdd	RoofStyle	RoofMatl	Exterior1st	Exterior2nd	MasVnrType	
0		2003	Gable	CompShg	VinylSd	VinylSd	BrkFace	
1		1976	Gable	CompShg	MetalSd	MetalSd	None	
2		2002	Gable	CompShg	VinylSd	VinylSd	BrkFace	
3		1970	Gable	CompShg	Wd Sdng	Wd Shng	None	
4		2000	Gable	CompShg	VinylSd	VinylSd	BrkFace	

		MasVnrArea	ExterQual	ExterCond	Foundation	BsmtQual	BsmtCond	BsmtExposure	
0		196.0	Gd	TA	PConc	Gd	TA	No	
1		0.0	TA	TA	CBlock	Gd	TA	Gd	
2		162.0	Gd	TA	PConc	Gd	TA	Mn	
3		0.0	TA	TA	BrkTil	TA	Gd	No	
4		350.0	Gd	TA	PConc	Gd	TA	Av	

		BsmtFinType1	BsmtFinSF1	BsmtFinType2	BsmtFinSF2	BsmtUnfSF	TotalBsmtSF	
0		GLQ	706	Unf	0	150	856	
1		ALQ	978	Unf	0	284	1262	

2	GLQ	486	Unf	0	434	920
3	ALQ	216	Unf	0	540	756
4	GLQ	655	Unf	0	490	1145

	Heating	HeatingQC	CentralAir	Electrical	1stFlrSF	2ndFlrSF	LowQualFinSF	\
0	GasA	Ex	Y	SBrkr	856	854	0	
1	GasA	Ex	Y	SBrkr	1262	0	0	
2	GasA	Ex	Y	SBrkr	920	866	0	
3	GasA	Gd	Y	SBrkr	961	756	0	
4	GasA	Ex	Y	SBrkr	1145	1053	0	

	GrLivArea	BsmtFullBath	BsmtHalfBath	FullBath	HalfBath	BedroomAbvGr	\
0	1710	1	0	2	1	3	
1	1262	0	1	2	0	3	
2	1786	1	0	2	1	3	
3	1717	1	0	1	0	3	
4	2198	1	0	2	1	4	

	KitchenAbvGr	KitchenQual	TotRmsAbvGrd	Functional	Fireplaces	FireplaceQu	\
0	1	Gd	8	Typ	0	NaN	
1	1	TA	6	Typ	1	TA	
2	1	Gd	6	Typ	1	TA	
3	1	Gd	7	Typ	1	Gd	
4	1	Gd	9	Typ	1	TA	

	GarageType	GarageYrBltd	GarageFinish	GarageCars	GarageArea	GarageQual	\
0	Attchd	2003.0	RFn	2	548	TA	
1	Attchd	1976.0	RFn	2	460	TA	
2	Attchd	2001.0	RFn	2	608	TA	
3	Detchd	1998.0	Unf	3	642	TA	
4	Attchd	2000.0	RFn	3	836	TA	

	GarageCond	PavedDrive	WoodDeckSF	OpenPorchSF	EnclosedPorch	3SsnPorch	\
0	TA	Y	0	61	0	0	
1	TA	Y	298	0	0	0	
2	TA	Y	0	42	0	0	
3	TA	Y	0	35	272	0	
4	TA	Y	192	84	0	0	

	ScreenPorch	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	YrSold	\
0	0	0	NaN	NaN	NaN	0	2	2008	
1	0	0	NaN	NaN	NaN	0	5	2007	
2	0	0	NaN	NaN	NaN	0	9	2008	
3	0	0	NaN	NaN	NaN	0	2	2006	
4	0	0	NaN	NaN	NaN	0	12	2008	

SaleType	SaleCondition	SalePrice
----------	---------------	-----------

0	WD	Normal	208500
1	WD	Normal	181500
2	WD	Normal	223500
3	WD	Abnorml	140000
4	WD	Normal	250000

1.1.2 Types of variables

[4]: *# we have an Id variable, that we should not use for predictions:*

```
print('Number of House Id labels: ', len(data.Id.unique()))
print('Number of Houses in the Dataset: ', len(data))
```

Number of House Id labels: 1460

Number of Houses in the Dataset: 1460

Find categorical variables

[5]: *# find categorical variables- hint data type = 'O'*

```
categorical = [var for var in data.columns if data[var].dtype=='O']

print(f'There are {len(categorical)} categorical variables')
categorical
```

There are 43 categorical variables

```
[5]: ['MSZoning',
      'Street',
      'Alley',
      'LotShape',
      'LandContour',
      'Utilities',
      'LotConfig',
      'LandSlope',
      'Neighborhood',
      'Condition1',
      'Condition2',
      'BldgType',
      'HouseStyle',
      'RoofStyle',
      'RoofMatl',
      'Exterior1st',
      'Exterior2nd',
      'MasVnrType',
      'ExterQual',
      'ExterCond',
      'Foundation',
```

```

'BsmtQual',
'BsmtCond',
'BsmtExposure',
'BsmtFinType1',
'BsmtFinType2',
'Heating',
'HeatingQC',
'CentralAir',
'Electrical',
'KitchenQual',
'Functional',
'FireplaceQu',
'GarageType',
'GarageFinish',
'GarageQual',
'GarageCond',
'PavedDrive',
'PoolQC',
'Fence',
'MiscFeature',
'SaleType',
'SaleCondition']

```

Find temporal variables

```

[6]: # make a list of the numerical variables first= Hint data type != 0
numerical = [var for var in data.columns if data[var].dtype!='O']

# list of variables that contain year information= Hint variable name has Yr
→ or
year_vars = [var for var in numerical if 'Yr' in var or 'Year' in var]

year_vars

```

```

[6]: ['YearBuilt', 'YearRemodAdd', 'GarageYrBlt', 'YrSold']

```

```

[7]: numerical

```

```

[7]: ['Id',
'MSSubClass',
'LotFrontage',
'LotArea',
'OverallQual',
'OverallCond',
'YearBuilt',
'YearRemodAdd',
'MasVnrArea',

```

```

'BsmtFinSF1',
'BsmtFinSF2',
'BsmtUnfSF',
'TotalBsmtSF',
'1stFlrSF',
'2ndFlrSF',
'LowQualFinSF',
'GrLivArea',
'BsmtFullBath',
'BsmtHalfBath',
'FullBath',
'HalfBath',
'BedroomAbvGr',
'KitchenAbvGr',
'TotRmsAbvGrd',
'Fireplaces',
'GarageYrBlt',
'GarageCars',
'GarageArea',
'WoodDeckSF',
'OpenPorchSF',
'EnclosedPorch',
'3SsnPorch',
'ScreenPorch',
'PoolArea',
'MiscVal',
'MoSold',
'YrSold',
'SalePrice']

```

Find discrete variables To identify discrete variables- numerical variables with less than 20 unique values

```

[8]: # let's visualise the values of the discrete variables
discrete = [var for var in numerical if len(data[var].unique()) < 20 and var_
↪not in year_vars]

print(f'There are {len(discrete)} discrete variables')
discrete

```

There are 14 discrete variables

```

[8]: ['MSSubClass',
'OverallQual',
'OverallCond',
'BsmtFullBath',
'BsmtHalfBath',

```

```
'FullBath',  
'HalfBath',  
'BedroomAbvGr',  
'KitchenAbvGr',  
'TotRmsAbvGrd',  
'Fireplaces',  
'GarageCars',  
'PoolArea',  
'MoSold']
```

Continuous variables

```
[9]: # find continuous variables- hint numerical variables not in discrete and   
      ↳ year_years  
      # Also remove the Id variable and the target variable SalePrice  
      # which are both also numerical  
  
      continuous = [var for var in numerical if var not in discrete and var not in [  
                      'Id', 'SalePrice'] and var not in year_vars]  
  
      print('There are {} numerical and continuous variables'.format(len(numerical)))  
      continuous
```

There are 38 numerical and continuous variables

```
[9]: ['LotFrontage',  
      'LotArea',  
      'MasVnrArea',  
      'BsmtFinSF1',  
      'BsmtFinSF2',  
      'BsmtUnfSF',  
      'TotalBsmtSF',  
      '1stFlrSF',  
      '2ndFlrSF',  
      'LowQualFinSF',  
      'GrLivArea',  
      'GarageArea',  
      'WoodDeckSF',  
      'OpenPorchSF',  
      'EnclosedPorch',  
      '3SsnPorch',  
      'ScreenPorch',  
      'MiscVal']
```

1.1.3 Separate train and test set

```
[10]: # Let's separate into train and test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(data.drop(['Id', 'SalePrice'], axis=1),
                                                    data['SalePrice'],
                                                    test_size=0.1,
                                                    random_state=0)

X_train.shape, X_test.shape
```

```
[10]: ((1314, 79), (146, 79))
```

Now we will move on and engineer the features of this dataset. The most important part for this course.

1.1.4 Craete New Variables

Replace 'YearBuilt', 'YearRemodAdd', 'GarageYrBlt with time elapsed since YrSold So $\text{YearBuilt} = \text{YrSold} - \text{YearBuilt}$.

Similarly transform 'YearRemodAdd', 'GarageYrBlt. After making transformation drop YrSold

```
[11]: X_train.head()
```

```
[11]:
```

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
930	20	RL	73.0	8925	Pave	NaN	IR1	
656	20	RL	72.0	10007	Pave	NaN	IR1	
45	120	RL	61.0	7658	Pave	NaN	Reg	
1348	20	RL	NaN	16196	Pave	NaN	IR3	
55	20	RL	100.0	10175	Pave	NaN	IR1	

	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	\
930	HLS	AllPub	Inside	Gtl	Timber	Norm	
656	Lvl	AllPub	Inside	Gtl	NAMES	Norm	
45	Lvl	AllPub	Inside	Gtl	NridgHt	Norm	
1348	Low	AllPub	Inside	Gtl	SawyerW	Norm	
55	Lvl	AllPub	Inside	Gtl	NAMES	Norm	

	Condition2	BldgType	HouseStyle	OverallQual	OverallCond	YearBuilt	\
930	Norm	1Fam	1Story	8	5	2007	
656	Norm	1Fam	1Story	5	7	1959	
45	Norm	TwnhsE	1Story	9	5	2005	
1348	Norm	1Fam	1Story	7	5	1998	
55	Norm	1Fam	1Story	6	5	1964	

	YearRemodAdd	RoofStyle	RoofMatl	Exterior1st	Exterior2nd	MasVnrType	\
930	2007	Gable	CompShg	VinylSd	VinylSd	None	

656	2006	Gable	CompShg	HdBoard	HdBoard	BrkFace
45	2005	Hip	CompShg	MetalSd	MetalSd	BrkFace
1348	1998	Gable	CompShg	VinylSd	VinylSd	None
55	1964	Gable	CompShg	HdBoard	Plywood	BrkFace

	MasVnrArea	ExterQual	ExterCond	Foundation	BsmtQual	BsmtCond	\
930	0.0	Gd	TA	PConc	Gd	TA	
656	54.0	Gd	TA	CBlock	TA	TA	
45	412.0	Ex	TA	PConc	Ex	TA	
1348	0.0	Gd	TA	PConc	Gd	TA	
55	272.0	TA	TA	CBlock	TA	TA	

	BsmtExposure	BsmtFinType1	BsmtFinSF1	BsmtFinType2	BsmtFinSF2	\
930	Av	GLQ	16	Unf	0	
656	No	ALQ	806	Unf	0	
45	No	GLQ	456	Unf	0	
1348	Gd	GLQ	1443	Unf	0	
55	No	BLQ	490	Unf	0	

	BsmtUnfSF	TotalBsmtSF	Heating	HeatingQC	CentralAir	Electrical	\
930	1450	1466	GasA	Ex	Y	SBrkr	
656	247	1053	GasA	Ex	Y	SBrkr	
45	1296	1752	GasA	Ex	Y	SBrkr	
1348	39	1482	GasA	Ex	Y	SBrkr	
55	935	1425	GasA	Gd	Y	SBrkr	

	1stFlrSF	2ndFlrSF	LowQualFinSF	GrLivArea	BsmtFullBath	BsmtHalfBath	\
930	1466	0	0	1466	0	0	
656	1053	0	0	1053	1	0	
45	1752	0	0	1752	1	0	
1348	1494	0	0	1494	1	0	
55	1425	0	0	1425	0	0	

	FullBath	HalfBath	BedroomAbvGr	KitchenAbvGr	KitchenQual	\
930	2	0	3	1	Gd	
656	1	1	3	1	Gd	
45	2	0	2	1	Ex	
1348	2	0	3	1	Gd	
55	2	0	3	1	TA	

	TotRmsAbvGrd	Functional	Fireplaces	FireplaceQu	GarageType	GarageYrBlt	\
930	7	Typ	0	NaN	Attchd	2007.0	
656	5	Typ	0	NaN	Attchd	1959.0	
45	6	Typ	1	Gd	Attchd	2005.0	
1348	5	Typ	1	Fa	Attchd	1998.0	
55	7	Typ	1	Gd	Attchd	1964.0	

	GarageFinish	GarageCars	GarageArea	GarageQual	GarageCond	PavedDrive	\
930	Fin	3	610	TA	TA	Y	
656	RFn	1	312	TA	TA	Y	
45	RFn	2	576	TA	TA	Y	
1348	RFn	2	514	TA	TA	Y	
55	RFn	2	576	TA	TA	Y	

	WoodDeckSF	OpenPorchSF	EnclosedPorch	3SsnPorch	ScreenPorch	\
930	100	18	0	0	0	
656	0	0	0	0	0	
45	196	82	0	0	0	
1348	402	25	0	0	0	
55	0	0	0	407	0	

	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	YrSold	SaleType	\
930	0	NaN	NaN	NaN	0	7	2009	WD	
656	0	NaN	MnPrv	NaN	0	8	2008	WD	
45	0	NaN	NaN	NaN	0	2	2010	WD	
1348	0	NaN	NaN	NaN	0	8	2007	WD	
55	0	NaN	NaN	NaN	0	7	2008	WD	

	SaleCondition
930	Normal
656	Normal
45	Normal
1348	Normal
55	Normal

```
[12]: # function to calculate elapsed time

def elapsed_years(df, var):
    # capture difference between year variable and
    # year the house was sold

    df[var] = df['YrSold'] - df[var]
    return df
```

```
[13]: for var in ['YearBuilt', 'YearRemodAdd', 'GarageYrBlt']:
    X_train = elapsed_years(X_train, var)
    X_test = elapsed_years(X_test, var)
```

```
[14]: # drop YrSold
X_train.drop('YrSold', axis=1, inplace=True)
X_test.drop('YrSold', axis=1, inplace=True)
```

```
[15]: year_vars.remove('YrSold')
```

```
[16]: # capture the column names for use later in the notebook
final_columns = X_train.columns
final_columns
```

```
[16]: Index(['MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street', 'Alley',
        'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope',
        'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
        'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle',
        'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea',
        'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
        'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2',
        'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC',
        'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
        'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
        'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
        'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType', 'GarageYrBlt',
        'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond',
        'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
        'ScreenPorch', 'PoolArea', 'PoolQC', 'Fence', 'MiscFeature', 'MiscVal',
        'MoSold', 'SaleType', 'SaleCondition'],
        dtype='object')
```

1.1.5 Feature Engineering Pipeline

```
[17]: # I will treat discrete variables as if they were categorical
# to treat discrete as categorical using Feature-engine
# we need to re-cast them as object

X_train[discrete] = X_train[discrete].astype('O')
X_test[discrete] = X_test[discrete].astype('O')
```

```
[18]: # import relevant modules for feature engineering
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from feature_engine import missing_data_imputers as mdi
from feature_engine import categorical_encoders as ce
from feature_engine.variable_transformers import YeoJohnsonTransformer
from sklearn.preprocessing import StandardScaler
from feature_engine.discretisers import DecisionTreeDiscretiser
```

```
[19]: house_preprocess = Pipeline([

    # missing data imputation
    ('missing_ind', mdi.AddNaNBinaryImputer(
        variables=['LotFrontage', 'MasVnrArea', 'GarageYrBlt'])),
    ('imputer_num', mdi.MeanMedianImputer(imputation_method='mean',
```

```

variables=['LotFrontage',
↳ 'MasVnrArea', 'GarageYrBlt'])),
    ('imputer_cat', mdi.CategoricalVariableImputer(variables=categorical)),

    # categorical encoding
    ('rare_label_enc', ce.RareLabelCategoricalEncoder(
        tol=0.01,n_categories=6, variables=categorical+discrete)),
    ('categorical_enc', ce.MeanCategoricalEncoder(variables = categorical +
↳ discrete))),

    # Transforming Numerical Variables
    ('yjt', YeoJohnsonTransformer(variables = ['LotFrontage','MasVnrArea',
↳ 'GarageYrBlt'])),

    # discretisation and encoding
    ('treeDisc', DecisionTreeDiscretiser(cv=2,
↳ scoring='neg_mean_squared_error',
        regression=True,
        param_grid={'max_depth': [1,2,3,4,5,6]})),

    # feature Scaling
    ('scaler', StandardScaler()),

])

```

```
[20]: house_preprocess.fit(X_train,y_train)
```

```

[20]: Pipeline(memory=None,
    steps=[('missing_ind',
        AddNaNBinaryImputer(variables=['LotFrontage', 'MasVnrArea',
        'GarageYrBlt'])),
    ('imputer_num',
        MeanMedianImputer(imputation_method='mean',
            variables=['LotFrontage', 'MasVnrArea',
            'GarageYrBlt'])),
    ('imputer_cat',
        CategoricalVariableImputer(variables=['MSZoning', 'Street',
        'Alley', 'LotShape',
        'LandContour',
        'Utilities', '...
        'Utilities', 'LotConfig',
        'LandSlope', 'Neighborhood',
        'Condition1', 'Condition2',
        'BldgType', 'HouseStyle',

```

```

        'OverallQual',
        'OverallCond', 'YearBuilt',
        'YearRemodAdd', 'RoofStyle',
        'RoofMatl', 'Exterior1st',
        'Exterior2nd', 'MasVnrType',
        'MasVnrArea', 'ExterQual',
        'ExterCond', 'Foundation',
        'BsmtQual', ...]],
        ('scaler',
         StandardScaler(copy=True, with_mean=True, with_std=True)]],
        verbose=False)

```

```

[21]: # Apply Transformations
X_train=house_preprocess.transform(X_train)
X_test=house_preprocess.transform(X_test)

```

1.2 DO NOT CHANGE STEPS BEFORE THIS POINT

1.3 Simple Models

1.3.1 Linear Regression

```

[22]: # Train a linear regression model, report the coefficients and model
      ↪ performance

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score

lr = LinearRegression().fit(X_train, y_train)
cv_scores = cross_val_score(lr, X_train, y_train)

```

Results

```

[23]: x_lrtrain= lr.predict(X_train)
      x_lrtest = lr.predict(X_test)

      # check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_lrtrain)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_lrtrain))))
print('train r2: {}'.format(r2_score(y_train, x_lrtrain)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_lrtest)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_lrtest))))
print('test r2: {}'.format(r2_score(y_test, x_lrtest)))

#print("Best parameters: {}".format(lr.best_params_))
print("Mean Cross-validation scores: {}".format(cv_scores))

```

```
print('Train score: {:.4f}'.format(lr.score(X_train, y_train)))
print('Test score: {:.4f}'.format(lr.score(X_test, y_test)))
```

```
train mse: 552937051.405446
train rmse: 23514.61357125492
train r2: 0.9114426745730106
```

```
test mse: 861097775.8563162
test rmse: 29344.467551078793
test r2: 0.8746968378820517
```

```
Mean Cross-validation scores: [ 8.68312903e-01 -6.26561684e+21  8.76203279e-01
 8.98626806e-01
 8.92290590e-01]
```

```
Train score: 0.9114
```

```
Test score: 0.8747
```

1.3.2 Ridge Regression

```
[24]: # Train a Ridge regression model, report the coefficients, the best parameters,
      ↪ and model performance
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
ridge = Ridge()

#define a list of parameters
#param_ridge = {'alpha':[0.001, 0.01, 0.1, 1, 10, 100] }
param_ridge = {'alpha':[1, 1e3, 1e4] }
grid_ridge = GridSearchCV(ridge, param_ridge, cv=10, return_train_score =
      ↪ True,scoring='r2')
grid_ridge.fit(X_train, y_train)
```

```
[24]: GridSearchCV(cv=10, error_score=nan,
                  estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
                                  max_iter=None, normalize=False, random_state=None,
                                  solver='auto', tol=0.001),
                  iid='deprecated', n_jobs=None,
                  param_grid={'alpha': [1, 1000.0, 10000.0]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring='r2', verbose=0)
```

Results

```
[25]: x_gridtrain= grid_ridge.predict(X_train)
      x_gridtest = grid_ridge.predict(X_test)

      # check model performance:
```

```

print('train mse: {}'.format(mean_squared_error(y_train, x_gridtrain)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_gridtrain))))
print('train r2: {}'.format(r2_score(y_train, x_gridtrain)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_gridtest)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_gridtest))))
print('test r2: {}'.format(r2_score(y_test, x_gridtest)))

print("Best parameters: {}".format(grid_ridge.best_params_))
print("Cross-validation scores: {}".format(grid_ridge.best_score_))
print('Train score: {:.4f}'.format(grid_ridge.score(X_train, y_train)))
print('Test score: {:.4f}'.format(grid_ridge.score(X_test, y_test)))

```

```

train mse: 552540520.7585117
train rmse: 23506.1804800038
train r2: 0.911506182152134

```

```

test mse: 858035459.791463
test rmse: 29292.242314159957
test r2: 0.8751424526508844
Best parameters: {'alpha': 1}
Cross-validation scores: 0.8843424460850325
Train score: 0.9115
Test score: 0.8751

```

1.3.3 Lasso

```

[26]: # Train a Lasso regression model, report the coefficients, the best parameters,
      ↪and model performance

      # YOUR CODE HERE

      from sklearn.linear_model import Lasso
      lasso = Lasso(random_state=0)

      #define a list of parameters
      #param_lasso = {'alpha':[0.001, 0.01, 0.1, 1, 10, 100] }
      param_lasso = {'alpha':[1, 1e4, 1e5] }

      grid_lasso = GridSearchCV(lasso, param_lasso, cv=10, return_train_score =
      ↪True,scoring='r2')
      grid_lasso.fit(X_train, y_train)

```

```

[26]: GridSearchCV(cv=10, error_score=nan,
                  estimator=Lasso(alpha=1.0, copy_X=True, fit_intercept=True,
                                  max_iter=1000, normalize=False, positive=False,
                                  precompute=False, random_state=0,

```

```

        selection='cyclic', tol=0.0001, warm_start=False),
iid='deprecated', n_jobs=None,
param_grid={'alpha': [1, 10000.0, 100000.0]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring='r2', verbose=0)

```

Results

```

[27]: x_lassotrain= grid_lasso.predict(X_train)
      x_lassotest = grid_lasso.predict(X_test)

      # check model performance:

      print('train mse: {}'.format(mean_squared_error(y_train, x_lassotrain)))
      print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_lassotrain))))
      print('train r2: {}'.format(r2_score(y_train, x_lassotrain)))
      print()
      print('test mse: {}'.format(mean_squared_error(y_test, x_lassotest)))
      print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_lassotest))))
      print('test r2: {}'.format(r2_score(y_test, x_lassotest)))

      print("Best parameters: {}".format(grid_lasso.best_params_))
      print("Cross-validation scores: {}".format(grid_lasso.best_score_))
      print('Train score: {:.4f}'.format(grid_lasso.score(X_train, y_train)))
      print('Test score: {:.4f}'.format(grid_lasso.score(X_test, y_test)))

```

```

train mse: 552539222.3512008
train rmse: 23506.152861563733
train r2: 0.9115063901025299

test mse: 857849072.1153543
test rmse: 29289.060621934503
test r2: 0.87516957496597
Best parameters: {'alpha': 1}
Cross-validation scores: 0.8843049751400216
Train score: 0.9115
Test score: 0.8752

```

1.4 Linear Regression with SGD

```

[28]: from sklearn.linear_model import SGDRegressor

```

```

[29]: # create pipeline
      reg_sgd_pipe = Pipeline([

          # feature Scaling
          ('scaler', MinMaxScaler()),

```



```

# regression
('sgd_reg', SGDRegressor(max_iter=1000, tol = 1e-6))
])

param_sgd = {'sgd_reg__eta0':[0.01, 0.05, 0.1 ,0.5],
             'sgd_reg__penalty' :['l1','l2'],#lasso,ridge
             'sgd_reg__alpha' :[0.1,0.01,0.001] }

grid_linearsgd = GridSearchCV(reg_sgd_pipe,
                              param_sgd,cv=5,
                              n_jobs=-1,
                              return_train_score = True,
                              scoring='r2')

# let's fit the pipeline
grid_linearsgd.fit(X_train, y_train)

# let's get the predictions
X_train_preds = grid_linearsgd.predict(X_train)
X_test_preds = grid_linearsgd.predict(X_test)

```

1.4.1 Results

```

[30]: # check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, X_train_preds)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, X_train_preds))))
print('train r2: {}'.format(r2_score(y_train, X_train_preds)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, X_test_preds)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, X_test_preds))))
print('test r2: {}'.format(r2_score(y_test, X_test_preds)))

print("Best parameters: {}".format(grid_linearsgd.best_params_))
print('Train score: {:.4f}'.format(grid_linearsgd.score(X_train, y_train)))
print('Test score: {:.4f}'.format(grid_linearsgd.score(X_test, y_test)))

```

train mse: 594505167.346368

train rmse: 24382.476645049166

train r2: 0.9047852057681796

test mse: 957543094.9385668

test rmse: 30944.19323457257

test r2: 0.8606625391167776

Best parameters: {'sgd_reg__alpha': 0.01, 'sgd_reg__eta0': 0.05,
'sgd_reg__penalty': 'l1'}

Train score: 0.9048

Test score: 0.8607

1.5 Polynomial Regression

```
[31]: #apply polynomial regression in pipeline
#pipe_poly = make_pipeline(PolynomialFeatures(),MinMaxScaler(),
#                           ↪LinearRegression())
pipe_poly=Pipeline([
    ('polynomialfeatures', PolynomialFeatures()),
    ('scaler',MinMaxScaler()),
    ('norm_reg', LinearRegression())

])
#define a list of parameters
param_poly = {'polynomialfeatures__degree':range(1,3)}

grid_poly = GridSearchCV(pipe_poly, param_poly,cv=5, n_jobs=-1,
#                           ↪return_train_score = True,scoring='r2')

grid_poly.fit(X_train, y_train)

# let's get the predictions
X_train_preds = grid_poly.predict(X_train)
X_test_preds = grid_poly.predict(X_test)
```

1.5.1 Results

```
[32]: x_polytrain= grid_lasso.predict(X_train)
x_polytest = grid_lasso.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_polytrain)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_polytrain))))
print('train r2: {}'.format(r2_score(y_train, x_polytrain)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_polytest)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_polytest))))
print('test r2: {}'.format(r2_score(y_test, x_polytest)))

print("Best parameters: {}".format(grid_poly.best_params_))
print("Cross-validation scores: {}".format(grid_poly.best_score_))
print('Train score: {:.4f}'.format(grid_poly.score(X_train, y_train)))
print('Test score: {:.4f}'.format(grid_poly.score(X_test, y_test)))
```

train mse: 552539222.3512008

```

train rmse: 23506.152861563733
train r2: 0.9115063901025299

test mse: 857849072.1153543
test rmse: 29289.060621934503
test r2: 0.87516957496597
Best parameters: {'polynomialfeatures__degree': 2}
Cross-validation scores: -9.790048692779738e+22
Train score: 1.0000
Test score: -312603571663290826752.0000

```

```

[33]: print('train mse: {}'.format(mean_squared_error(y_train, X_train_preds)))
      print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, X_train_preds))))
      print('train r2: {}'.format(r2_score(y_train, X_train_preds)))
      print()
      print('test mse: {}'.format(mean_squared_error(y_test, X_test_preds)))
      print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, X_test_preds))))
      print('test r2: {}'.format(r2_score(y_test, X_test_preds)))

```

```

train mse: 27597.318873668188
train rmse: 166.12440782036873
train r2: 0.9999955800669494

test mse: 2.1482477834886397e+30
test rmse: 1465690207202272.2
test r2: -3.126035716632908e+20

```

1.6 ElasticNet

```

[34]: from sklearn.linear_model import ElasticNet

      elasticnet = ElasticNet().fit(X_train, y_train)

      #define a list of parameters
      param_elasticnet = {'alpha':[1, 1e4, 1e5],
                          'l1_ratio':[0.2,0.4,0.6,0.8]}

      grid_elasticnet = GridSearchCV(elasticnet , param_elasticnet, cv=5,
      ↪return_train_score = True,scoring='r2')
      grid_elasticnet.fit(X_train, y_train)

      x_elastictrain_predict = grid_elasticnet.predict(X_train)
      x_elastictest_predict = grid_elasticnet.predict(X_test)

```

1.6.1 Results

```
[35]: print('train mse: {}'.format(mean_squared_error(y_train,
    ↪x_elastictrain_predict)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train,
    ↪x_elastictrain_predict))))
print('train r2: {}'.format(r2_score(y_train, x_elastictrain_predict)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_elastictest_predict)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test,
    ↪x_elastictest_predict))))
print('test r2: {}'.format(r2_score(y_test, x_elastictest_predict)))
print()
print('Best parameters: ', grid_elasticnet.best_params_)
print('Best cross-validation score:', grid_elasticnet.score(X_test, y_test))
print("Training set score: {:.2f}".format(grid_elasticnet.score(X_train,
    ↪y_train)))
print("Test set score: {:.2f}".format(grid_elasticnet.score(X_test, y_test)))
```

```
train mse: 580812656.1496782
train rmse: 24100.055106776796
train r2: 0.9069781717972707
```

```
test mse: 908630644.7418916
test rmse: 30143.500870699998
test r2: 0.8677800638026185
```

```
Best parameters: {'alpha': 1, 'l1_ratio': 0.8}
Best cross-validation score: 0.8677800638026185
Training set score: 0.91
Test set score: 0.87
```

1.7 Tune Multiple Models with one GridSearch

1.7.1 Model - Linear Regression with Model Parameter - Ridge *GridSearch

```
[36]: model_linear = Pipeline([("regressor", LinearRegression())] )
```

```
[37]: model_parm_gd1 = [{ 'regressor': [Ridge()]}]
```

```
[38]: from sklearn import svm
grid_search_linear = GridSearchCV(svm.SVC(gamma='auto'),{
    'C':[30],
    'kernel':['linear']
},cv=5,return_train_score=False,scoring='r2')
```

```
[39]: grid_search_linear = GridSearchCV(model_linear,model_parm_gd1,scoring='r2')
```

```
[40]: grid_search_linear.fit(X_train,y_train)
```

```
[40]: GridSearchCV(cv=None, error_score=nan,
                  estimator=Pipeline(memory=None,
                                     steps=[('regressor',
                                             LinearRegression(copy_X=True,
                                                                fit_intercept=True,
                                                                n_jobs=None,
                                                                normalize=False))],
                                     verbose=False),
                  iid='deprecated', n_jobs=None,
                  param_grid=[{'regressor': [Ridge(alpha=1.0, copy_X=True,
                                                    fit_intercept=True, max_iter=None,
                                                    normalize=False,
                                                    random_state=None, solver='auto',
                                                    tol=0.001)]]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring='r2', verbose=0)
```

Results

```
[41]: x_lrtrain= grid_search_linear.predict(X_train)
      x_lrtest = grid_search_linear.predict(X_test)

      # check model performance:

      print('train mse: {}'.format(mean_squared_error(y_train, x_lrtrain)))
      print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_lrtrain))))
      print('train r2: {}'.format(r2_score(y_train, x_lrtrain)))
      print()
      print('test mse: {}'.format(mean_squared_error(y_test, x_lrtest)))
      print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_lrtest))))
      print('test r2: {}'.format(r2_score(y_test, x_lrtest)))
      print()

      print("Best parameters: {}".format(grid_search_linear.best_params_))
      print("Cross-validation scores: {}".format(grid_search_linear.best_score_))
      print('Train score: {:.4f}'.format(grid_search_linear.score(X_train, y_train)))
      print('Test score: {:.4f}'.format(grid_search_linear.score(X_test, y_test)))
```

```
train mse: 552540520.7585117
train rmse: 23506.1804800038
train r2: 0.911506182152134
```

```
test mse: 858035459.791463
test rmse: 29292.242314159957
test r2: 0.8751424526508844
```

```
Best parameters: {'regressor': Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
max_iter=None,
    normalize=False, random_state=None, solver='auto', tol=0.001)}
Cross-validation scores: 0.8818969790895617
Train score: 0.9115
Test score: 0.8751
```

1.8 Linear Regression with Model Parameter - Lasso *GridSearchCV

```
[42]: model_linear2 = Pipeline([("regressor", LinearRegression())] )

[43]: model_parm_gd2 = [{ 'regressor': [Lasso(random_state=0)]]}

[44]: grid_search_linear2 = GridSearchCV(svm.SVC(gamma='auto'),{
    'C':[1,10,20,30],
    'kernel':['polynomial']

    },cv=5,return_train_score=False,scoring='r2')

[45]: grid_search_linear2 = GridSearchCV(model_linear2,model_parm_gd2,scoring='r2')

[46]: grid_search_linear2.fit(X_train,y_train)

[46]: GridSearchCV(cv=None, error_score=nan,
    estimator=Pipeline(memory=None,
        steps=[('regressor',
            LinearRegression(copy_X=True,
                fit_intercept=True,
                n_jobs=None,
                normalize=False))],
            verbose=False),
    iid='deprecated', n_jobs=None,
    param_grid=[{'regressor': [Lasso(alpha=1.0, copy_X=True,
        fit_intercept=True, max_iter=1000,
        normalize=False, positive=False,
        precompute=False, random_state=0,
        selection='cyclic', tol=0.0001,
        warm_start=False)]]},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
    scoring='r2', verbose=0)
```

Results

```
[47]: x_lrtrain= grid_search_linear2.predict(X_train)
x_lrtest = grid_search_linear2.predict(X_test)

# check model performance:
```

```

print('train mse: {}'.format(mean_squared_error(y_train, x_lrtrain)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_lrtrain))))
print('train r2: {}'.format(r2_score(y_train, x_lrtrain)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_lrtest)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_lrtest))))
print('test r2: {}'.format(r2_score(y_test, x_lrtest)))
print()
print("Best parameters: {}".format(grid_search_linear2.best_params_))
print("Cross-validation scores: {}".format(grid_search_linear2.best_score_))
print('Train score: {:.4f}'.format(grid_search_linear2.score(X_train, y_train)))
print('Test score: {:.4f}'.format(grid_search_linear2.score(X_test, y_test)))

```

```

train mse: 552539222.3512008
train rmse: 23506.152861563733
train r2: 0.9115063901025299

```

```

test mse: 857849072.1153543
test rmse: 29289.060621934503
test r2: 0.87516957496597

```

```

Best parameters: {'regressor': Lasso(alpha=1.0, copy_X=True, fit_intercept=True,
max_iter=1000,
    normalize=False, positive=False, precompute=False, random_state=0,
    selection='cyclic', tol=0.0001, warm_start=False)}
Cross-validation scores: 0.881849542095595
Train score: 0.9115
Test score: 0.8752

```

1.9 Pipeline - RandomForest Regressor with Model Parameter - Linear Regression *GridSearchCV

```
[48]: from sklearn.ensemble import RandomForestRegressor
```

```
[49]: model_rf = Pipeline([('regressor', RandomForestRegressor(random_state = 42))])
```

```
[50]: model_parm_rf = [{ 'regressor': [LinearRegression()]}]
```

```
[51]: from sklearn import svm
```

```
[52]: grid_rf = GridSearchCV(svm.SVC(gamma='auto'),{
    'C':[30],
    'kernel':['linear']

},cv=5,return_train_score=False,scoring='r2')
```

```
[53]: grid_rf = GridSearchCV(model_rf, model_parm_rf)
```

```
[54]: grid_rf.fit(X_train,y_train)
```

```
[54]: GridSearchCV(cv=None, error_score=nan,
                  estimator=Pipeline(memory=None,
                                     steps=[('regressor',
                                             RandomForestRegressor(bootstrap=True,
                                                                      ccp_alpha=0.0,
                                                                      criterion='mse',
                                                                      max_depth=None,
                                                                      max_features='auto',
                                                                      max_leaf_nodes=None,
                                                                      max_samples=None,
                                                                      min_impurity_decrease=0.0,
                                                                      min_impurity_split=None,
                                                                      min_samples_leaf=1,
                                                                      min_samples_split=2,
                                                                      min_weight_fraction_leaf=0.0,
                                                                      n_estimators=100,
                                                                      n_jobs=None,
                                                                      oob_score=False,
                                                                      random_state=42,
                                                                      verbose=0,
                                                                      warm_start=False))),
                                     verbose=False),
                  iid='deprecated', n_jobs=None,
                  param_grid=[{'regressor': [LinearRegression(copy_X=True,
                                                                fit_intercept=True,
                                                                n_jobs=None,
                                                                normalize=False)]}],
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring=None, verbose=0)
```

1.9.1 Results

```
[55]: x_rftrain= grid_rf.predict(X_train)
x_rftest = grid_rf.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_rftrain)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_rftrain))))
print('train r2: {}'.format(r2_score(y_train, x_rftrain)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_rftest)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_rftest))))
print('test r2: {}'.format(r2_score(y_test, x_rftest)))
```



```
print("Best parameters: {}".format(grid_rf.best_params_))
print("Cross-validation scores: {}".format(grid_rf.best_score_))
print('Train score: {:.4f}'.format(grid_rf.score(X_train, y_train)))
print('Test score: {:.4f}'.format(grid_rf.score(X_test, y_test)))
```

```
train mse: 552937051.405446
train rmse: 23514.61357125492
train r2: 0.9114426745730106
```

```
test mse: 861097775.8563162
test rmse: 29344.467551078793
test r2: 0.8746968378820517
Best parameters: {'regressor': LinearRegression(copy_X=True, fit_intercept=True,
n_jobs=None, normalize=False)}
Cross-validation scores: -1.2531233677405604e+21
Train score: 0.9114
Test score: 0.8747
```

1.10 KNN Regressor

```
[56]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import MinMaxScaler
knnreg = KNeighborsRegressor().fit(X_train, y_train)
print("Training set score: {:.2f}".format(knnreg.score(X_train, y_train)))
print("Test set score: {:.2f}".format(knnreg.score(X_test, y_test)))
```

```
Training set score: 0.88
Test set score: 0.59
```

```
[57]: pipe_knn = make_pipeline(MinMaxScaler(), KNeighborsRegressor())
knnreg = pipe_knn.fit(X_train, y_train)
print("Training set score: {:.2f}".format(knnreg.score(X_train, y_train)))
print("Test set score: {:.2f}".format(knnreg.score(X_test, y_test)))
```

```
Training set score: 0.86
Test set score: 0.71
```

```
[58]: from sklearn.model_selection import GridSearchCV

pipe_knn=Pipeline([
    ('scaler',MinMaxScaler()),
    ('knnreg', KNeighborsRegressor())

])
# define a list of parameters
#param_knn = {'n_neighbors': [5, 10, 15, 20, 25, 30]}
```

```

param_knn = {'knnreg__n_neighbors': range(1,25)}

#apply grid search
grid_knn = GridSearchCV(pipe_knn, param_knn, cv=5,
→return_train_score=True,scoring='r2')
grid_knn.fit(X_train, y_train)

```

```

[58]: GridSearchCV(cv=5, error_score=nan,
                estimator=Pipeline(memory=None,
                                steps=[('scaler',
                                        MinMaxScaler(copy=True,
                                                    feature_range=(0, 1))),
                                        ('knnreg',
                                        KNeighborsRegressor(algorithm='auto',
                                                            leaf_size=30,
                                                            metric='minkowski',
                                                            metric_params=None,
                                                            n_jobs=None,
                                                            n_neighbors=5, p=2,
                                                            weights='uniform')))],
                                verbose=False),
                iid='deprecated', n_jobs=None,
                param_grid={'knnreg__n_neighbors': range(1, 25)},
                pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                scoring='r2', verbose=0)

```

1.10.1 Results

```

[59]: x_knntrain= grid_knn.predict(X_train)
x_knnctest = grid_knn.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_knntrain)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_knntrain))))
print('train r2: {}'.format(r2_score(y_train, x_knntrain)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_knnctest)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_knnctest))))
print('test r2: {}'.format(r2_score(y_test, x_knnctest)))
print()
print('train score: ', grid_knn.score(X_train, y_train))
print('test score: ', grid_knn.score(X_test, y_test))

#find best parameters
print('Best parameters: ', grid_knn.best_params_)
print('Best cross-validation score:', grid_knn.best_score_)

```

```

train mse: 926540858.0310757
train rmse: 30439.13366098115
train r2: 0.8516070137143069

test mse: 1657648286.6700912
test rmse: 40714.227079364915
test r2: 0.7587863099604379

train score: 0.8516070137143069
test score: 0.7587863099604379
Best parameters: {'knnreg__n_neighbors': 6}
Best cross-validation score: 0.78587417331835

```

1.11 ElasticNet with GridSearch

```

[60]: elasticnet = ElasticNet()

#define a list of parameters
param_elasticnet = {'alpha':[0.00001, 0.0001, 0.001, 0.01, 0.1, 1, 10],
                    ↳'l1_ratio': [0.2,0.4,0.6,0.8]}

grid_elasticnet = GridSearchCV(elasticnet , param_elasticnet, cv=5,
↳return_train_score = True,scoring='r2')
grid_elasticnet.fit(X_train, y_train)

grid_elasticnet_train_score = grid_elasticnet.score(X_train, y_train)
grid_elasticnet_test_score = grid_elasticnet.score(X_test, y_test)

```

```

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 275261918430.7595, tolerance: 657118734.8147435
positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 270472843507.742, tolerance: 655974723.0327224
positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269415475236.1935, tolerance: 635955900.9288123
positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 293367949607.249, tolerance: 642749145.1567798
positive)

```

```

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 299702168310.0032, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 275132322643.92615, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 270265546612.78378, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269359988115.02875, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 293114639456.2878, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 299461621874.3815, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 274876218400.22867, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269853488069.79288, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269250900786.1243, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.

```

```

Duality gap: 292610782725.36975, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 298983659623.677, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 274113520713.21634, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 268620977582.17148, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 268927308309.11826, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 291102803101.93744, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 297554322499.26166, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 275285371342.0834, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 270500219669.26224, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269455263299.4938, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

```

```

packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 293390700602.9906, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 299724423882.7319, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 275126170208.67175, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 270267583340.30392, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269382447458.56412, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 293109101936.22723, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 299452447939.0266, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 274841772404.84204, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269830678356.61826, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269256329279.5095, tolerance: 635955900.9288123

```

```

    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 292577554374.5457, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 298943887810.3167, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 274051946790.85733, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 268573717281.67993, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 268915978578.34824, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 291042350452.3709, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 297484634292.5844, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 274840164312.18774, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 270621725492.51758, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:

```

```

Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269731153808.08066, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 293432778873.37506, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 299665315171.12964, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 274320207459.15918, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 270089455409.16162, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269461868093.74542, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 292811730158.3288, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 299004317909.24976, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 273940590197.94852, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269428587209.82605, tolerance: 655974723.0327224
    positive)

```



```

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 269193156672.4067, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 292033002979.416, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 298242533193.7177, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 273151756027.88977, tolerance: 657118734.8147435
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 267998525106.42767, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 268738473819.79742, tolerance: 635955900.9288123
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 290312550120.06335, tolerance: 642749145.1567798
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 296613343354.8944, tolerance: 689033704.679551
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 221086170647.84283, tolerance: 655974723.0327224
    positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.

```

```
Duality gap: 204650042677.63684, tolerance: 635955900.9288123
positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 232586340970.15808, tolerance: 642749145.1567798
positive)
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:476: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations.
Duality gap: 186705495812.65787, tolerance: 689033704.679551
positive)
```

1.11.1 Results

```
[61]: # let's get the predictions
x_elastictrain= grid_elasticnet.predict(X_train)
x_elastictest = grid_elasticnet.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_elastictrain)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train,
→x_elastictrain))))
print('train r2: {}'.format(r2_score(y_train, x_elastictrain)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_elastictest)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_elastictest))))
print('test r2: {}'.format(r2_score(y_test, x_elastictest)))

print('Training set score: ', grid_elasticnet_train_score)
print('Test score: ', grid_elasticnet_test_score)

#find best parameters
print('Best parameters: ', grid_elasticnet.best_params_)
print('Best cross-validation score:', grid_elasticnet.best_score_)
```

```
train mse: 560460520.2398213
train rmse: 23674.04739878294
train r2: 0.9102377304004107
```

```
test mse: 872749277.5992303
test rmse: 29542.33026691074
test r2: 0.8730013625797748
Training set score: 0.9102377304004107
Test score: 0.8730013625797748
Best parameters: {'alpha': 0.1, 'l1_ratio': 0.2}
Best cross-validation score: 0.8840700291378194
```

1.12 Decision Tree Regression

```
[62]: from sklearn.tree import DecisionTreeRegressor
      from sklearn.model_selection import GridSearchCV
      decisiontree = DecisionTreeRegressor(random_state = 0)

      param_DTree = {"max_depth": range(1,30),
                     "max_leaf_nodes": range(2,20),
                     "min_samples_split":range(2,30)}

      grid_dtree = GridSearchCV(decisiontree,param_DTree,cv=5,scoring='r2')
      grid_dtree.fit(X_train,y_train)

[62]: GridSearchCV(cv=5, error_score=nan,
                  estimator=DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse',
                                                    max_depth=None, max_features=None,
                                                    max_leaf_nodes=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    presort='deprecated',
                                                    random_state=0, splitter='best'),
                  iid='deprecated', n_jobs=None,
                  param_grid={'max_depth': range(1, 30),
                              'max_leaf_nodes': range(2, 20),
                              'min_samples_split': range(2, 30)},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                  scoring='r2', verbose=0)
```

1.12.1 Results

```
[63]: # let's get the predictions
      x_dtreetrain= grid_dtree.predict(X_train)
      x_dtreetest = grid_dtree.predict(X_test)

      print('train mse: {}'.format(mean_squared_error(y_train, x_dtreetrain)))
      print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_dtreetrain))))
      print('train r2: {}'.format(r2_score(y_train, x_dtreetrain)))
      print()
      print('test mse: {}'.format(mean_squared_error(y_test, x_dtreetest)))
      print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_dtreetest))))
      print('test r2: {}'.format(r2_score(y_test, x_dtreetest)))
      print()
      print('Best parameters: ', grid_dtree.best_params_)
```

```
print("Accuracy on training set: {:.4f}".format(grid_dtreescore(X_train,
→y_train)))
print("Accuracy on test set: {:.3f}".format(grid_dtreescore(X_test, y_test)))
```

```
train mse: 1105977258.816959
train rmse: 33256.23638984061
train r2: 0.8228688278802175
```

```
test mse: 1548654925.2624905
test rmse: 39352.953196202325
test r2: 0.7746465446714791
```

```
Best parameters: {'max_depth': 6, 'max_leaf_nodes': 19, 'min_samples_split':
29}
```

```
Accuracy on training set: 0.8229
```

```
Accuracy on test set: 0.775
```

1.13 SVM

```
[64]: from sklearn.svm import SVC
svc_kernel = SVC()

#define a list of parameters
param_svc_kernel = [{'C': [1,10,100,1000,10000],
                        'gamma':['auto'],
                        'kernel':['rbf']},
                    {'C': [1,10,100,1000,10000],
                        'gamma':['auto'],
                        'kernel':['poly']},
                    {'C': [1,10,100,1000,10000],
                        'gamma':['auto'],
                        'kernel':['sigmoid',]}]

#apply grid search
grid_svc_kernel = GridSearchCV(svc_kernel, param_svc_kernel, cv=5,
→n_jobs=2,scoring='r2')
grid_svc_kernel.fit(X_train, y_train)
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\model_selection\_split.py:667: UserWarning: The least populated
class in y has only 1 members, which is less than n_splits=5.
% (min_groups, self.n_splits)), UserWarning)
```

```
[64]: GridSearchCV(cv=5, error_score=nan,
                estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                    class_weight=None, coef0=0.0,
```

```

        decision_function_shape='ovr', degree=3,
        gamma='scale', kernel='rbf', max_iter=-1,
        probability=False, random_state=None, shrinking=True,
        tol=0.001, verbose=False),
iid='deprecated', n_jobs=2,
param_grid=[{'C': [1, 10, 100, 1000, 10000], 'gamma': ['auto'],
            'kernel': ['rbf']},
            {'C': [1, 10, 100, 1000, 10000], 'gamma': ['auto'],
            'kernel': ['poly']},
            {'C': [1, 10, 100, 1000, 10000], 'gamma': ['auto'],
            'kernel': ['sigmoid']}],
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring='r2', verbose=0)

```

1.13.1 Results

```

[65]: # let's get the predictions
x_svctrain= grid_svc_kernel.predict(X_train)
x_svctest = grid_svc_kernel.predict(X_test)

print('train mse: {}'.format(mean_squared_error(y_train, x_svctrain)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_svctrain))))
print('train r2: {}'.format(r2_score(y_train, x_svctrain)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_svctest)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_svctest))))
print('test r2: {}'.format(r2_score(y_test, x_svctest)))
print()
print('Best parameters: ', grid_svc_kernel.best_params_)
print('train score: ', grid_svc_kernel.score(X_train, y_train))
print('test score: ', grid_svc_kernel.score(X_test, y_test))

```

```

train mse: 61834.09436834094
train rmse: 248.66462226931466
train r2: 0.9999900967714073

```

```

test mse: 1296973682.7054794
test rmse: 36013.520831841466
test r2: 0.8112700924524573

```

```

Best parameters: {'C': 100, 'gamma': 'auto', 'kernel': 'poly'}
train score: 0.9999900967714073
test score: 0.9999900967714073

```

2 Ensemble

2.1 Bagging

2.1.1 With Descision Tree Regressor

```
[66]: from sklearn.linear_model import LogisticRegression
      from sklearn.tree import DecisionTreeRegressor
      from sklearn.ensemble import BaggingRegressor
```

```
[67]: bag_dtree1 = BaggingRegressor(base_estimator=DecisionTreeRegressor(),
      ↪bootstrap=True, random_state=0, oob_score=True)

      bag_dtree1_param = {
          'base_estimator__max_depth': range(1,10),
          'max_samples': [0.8,1],
          'n_estimators': range(1,80,5)}
      bag_dtree1_grid = GridSearchCV(bag_dtree1, bag_dtree1_param,cv=5,
      ↪return_train_score=True)
      bag_dtree1_grid.fit(X_train,y_train)
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
    warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
    warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
    warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
    warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
    warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
```

OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```


reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
```

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```


reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
```

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```


OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.  
    warn("Some inputs do not have OOB scores. "  
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.  
    warn("Some inputs do not have OOB scores. "  
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.  
    warn("Some inputs do not have OOB scores. "  
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.  
    warn("Some inputs do not have OOB scores. "  
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.  
    warn("Some inputs do not have OOB scores. "  
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.  
    warn("Some inputs do not have OOB scores. "  
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.  
    warn("Some inputs do not have OOB scores. "  
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.  
    warn("Some inputs do not have OOB scores. "
```

reliable oob estimates.

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
```

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
```

OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```


reliable oob estimates.

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-  
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have  
OOB scores. This probably means too few estimators were used to compute any  
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
```



```
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
```

```
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
[67]: GridSearchCV(cv=5, error_score=nan,
  estimator=BaggingRegressor(base_estimator=DecisionTreeRegressor(ccp_alpha=0.0,
  criterion='mse',
  max_depth=None,
  max_features=None,
  max_leaf_nodes=None,
  min_impurity_decrease=0.0,
  min_impurity_split=None,
  min_samples_leaf=1,
  min_samples_split=2,
  min_weight_fraction_leaf=0.0,
  presort='deprecated',
  random_state=None,
  splitter='be...

  bootstrap_features=False,
  max_features=1.0, max_samples=1.0,
  n_estimators=10, n_jobs=None,
  oob_score=True, random_state=0,
  verbose=0, warm_start=False),

  iid='deprecated', n_jobs=None,
  param_grid={'base_estimator__max_depth': range(1, 10),
              'max_samples': [0.8, 1],
              'n_estimators': range(1, 80, 5)},
  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
  scoring=None, verbose=0)
```

Results

```
[68]: # let's get the predictions
x_bagtree_train= bag_dtrees1_grid.predict(X_train)
x_bagtree_test = bag_dtrees1_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_bagtree_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train,
↪x_bagtree_train))))
print('train r2: {}'.format(r2_score(y_train, x_bagtree_train)))
```

```

print()
print('test mse: {}'.format(mean_squared_error(y_test, x_bagtree_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_bagtree_test))))
print('test r2: {}'.format(r2_score(y_test, x_bagtree_test)))
print()
print(f'Best Mean Cross Validation Score is {bag_dtree1_grid.best_score_}')
print(f'Best parameters {bag_dtree1_grid.best_params_}')
print(f'Train score is {bag_dtree1_grid.score(X_train,y_train)}')
print(f'Test score is {bag_dtree1_grid.score(X_test,y_test)}')

```

```

train mse: 238564809.90829495
train rmse: 15445.543367207738
train r2: 0.9617919228729973

```

```

test mse: 982627004.5486312
test rmse: 31346.881895152364
test r2: 0.8570124388836232

```

```

Best Mean Cross Validation Score is 0.8701118447613622
Best parameters {'base_estimator__max_depth': 9, 'max_samples': 0.8,
'n_estimators': 61}
Train score is 0.9617919228729972
Test score is 0.8570124388836233

```

2.1.2 With HW2 Best Model - Lasso

```

[69]: from sklearn.linear_model import LinearRegression
      from sklearn.linear_model import Lasso

```

```

[70]: bag_lasso = BaggingRegressor(base_estimator=Lasso(), bootstrap=True,
    ↪random_state=0, oob_score=True)

bag_lasso_param = {'base_estimator__alpha': [1, 1e4, 1e5]}

bag_lasso_grid = GridSearchCV(bag_lasso, bag_lasso_param,cv=5,
    ↪return_train_score=True)
bag_lasso_grid.fit(X_train,y_train)

```

```

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
  warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.

```

```
warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
    warn("Some inputs do not have OOB scores. "
```

OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. ")
```

```
[70]: GridSearchCV(cv=5, error_score=nan,
                  estimator=BaggingRegressor(base_estimator=Lasso(alpha=1.0,
                                                                    copy_X=True,
                                                                    fit_intercept=True,
                                                                    max_iter=1000,
                                                                    normalize=False,
                                                                    positive=False,
                                                                    precompute=False,
                                                                    random_state=None,
                                                                    selection='cyclic',
                                                                    tol=0.0001,
                                                                    warm_start=False),
                                                                    bootstrap=True,
                                                                    bootstrap_features=False,
                                                                    max_features=1.0, max_samples=1.0,
                                                                    n_estimators=10, n_jobs=None,
                                                                    oob_score=True, random_state=0,
                                                                    verbose=0, warm_start=False),
                  iid='deprecated', n_jobs=None,
                  param_grid={'base_estimator__alpha': [1, 10000.0, 100000.0]},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring=None, verbose=0)
```

Results

```
[71]: # let's get the predictions
x_baglasso_train= bag_lasso_grid.predict(X_train)
x_baglasso_test = bag_lasso_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_baglasso_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train,
→x_baglasso_train))))
print('train r2: {}'.format(r2_score(y_train, x_baglasso_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_baglasso_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_baglasso_test))))
print('test r2: {}'.format(r2_score(y_test, x_baglasso_test)))
print()
print(f'Best Mean Cross Validation Score is {bag_lasso_grid.best_score_}')
print(f'Best parameters {bag_lasso_grid.best_params_}')
print(f'Train score is {bag_lasso_grid.score(X_train,y_train)}')
print(f'Test score is {bag_lasso_grid.score(X_test,y_test)}')
```

train mse: 565445909.6571118

train rmse: 23779.106578194056

train r2: 0.9094392801032474

test mse: 853696103.4052668

test rmse: 29218.07836606074

test r2: 0.8757738967121659

Best Mean Cross Validation Score is 0.8849435910242107

Best parameters {'base_estimator__alpha': 1}

Train score is 0.9094392801032474

Test score is 0.8757738967121659

2.2 Pasting

2.2.1 With Decision Tree Regressor

```
[72]: paste_dtrees1 = BaggingRegressor(base_estimator=DecisionTreeRegressor(),
    ↪bootstrap=False, random_state=0, oob_score=False)

paste_dtrees1_param = {
    'base_estimator__max_depth': range(1,10),
    'max_samples': [0.8,1],
    'n_estimators': range(1,80,5)}
paste_dtrees1_grid = GridSearchCV(paste_dtrees1, paste_dtrees1_param,cv=5,
    ↪return_train_score=True)
paste_dtrees1_grid.fit(X_train,y_train)

[72]: GridSearchCV(cv=5, error_score=nan,
    estimator=BaggingRegressor(base_estimator=DecisionTreeRegressor(ccp_alpha=0.0,
    criterion='mse',
    max_depth=None,
    max_features=None,
    max_leaf_nodes=None,
    min_impurity_decrease=0.0,
    min_impurity_split=None,
    min_samples_leaf=1,
    min_samples_split=2,
    min_weight_fraction_leaf=0.0,
    presort='deprecated',
    random_state=None,
    splitter='be...

                                bootstrap_features=False,
                                max_features=1.0, max_samples=1.0,
                                n_estimators=10, n_jobs=None,
                                oob_score=False, random_state=0,
                                verbose=0, warm_start=False),
    iid='deprecated', n_jobs=None,
    param_grid={'base_estimator__max_depth': range(1, 10),
                'max_samples': [0.8, 1],
                'n_estimators': range(1, 80, 5)},
    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
    scoring=None, verbose=0)
```

Results

```
[73]: # let's get the predictions
x_pastetree_train= paste_dtrees1_grid.predict(X_train)
x_pastetree_test = paste_dtrees1_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_pastetree_train)))
```

```

print('train mse: {}'.format(sqrt(mean_squared_error(y_train,
↳x_pastetree_train))))
print('train r2: {}'.format(r2_score(y_train, x_pastetree_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_pastetree_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test,
↳x_pastetree_test))))
print('test r2: {}'.format(r2_score(y_test, x_pastetree_test)))
print()
print(f'Best Mean Cross Validation Score is {paste_dtree1_grid.best_score_}')
print(f'Best parameters {paste_dtree1_grid.best_params_}')
print(f'Train score is {paste_dtree1_grid.score(X_train,y_train)}')
print(f'Test score is {paste_dtree1_grid.score(X_test,y_test)}')

```

```

train mse: 151317451.74681643
train rmse: 12301.115874050469
train r2: 0.9757652904918109

```

```

test mse: 954334002.3152863
test rmse: 30892.296811912292
test r2: 0.8611295121650201

```

```

Best Mean Cross Validation Score is 0.8572589293630738
Best parameters {'base_estimator__max_depth': 9, 'max_samples': 0.8,
'n_estimators': 36}
Train score is 0.9757652904918109
Test score is 0.8611295121650201

```

2.2.2 With HW 2 Best Model - Lasso

```

[74]: paste_lasso = BaggingRegressor(base_estimator=Lasso(), bootstrap=False,
↳random_state=0, oob_score=False)

paste_dtree2_param = {'base_estimator__alpha': [1, 1e4, 1e5]}

paste_lasso_grid = GridSearchCV(paste_lasso, paste_dtree2_param,cv=5,
↳return_train_score=True)
paste_lasso_grid.fit(X_train,y_train)

```

```

[74]: GridSearchCV(cv=5, error_score=nan,
            estimator=BaggingRegressor(base_estimator=Lasso(alpha=1.0,
                                                    copy_X=True,
                                                    fit_intercept=True,
                                                    max_iter=1000,
                                                    normalize=False,
                                                    positive=False,
                                                    precompute=False,

```

```

        random_state=None,
        selection='cyclic',
        tol=0.0001,
        warm_start=False),

        bootstrap=False,
        bootstrap_features=False,
        max_features=1.0, max_samples=1.0,
        n_estimators=10, n_jobs=None,
        oob_score=False, random_state=0,
        verbose=0, warm_start=False),

        iid='deprecated', n_jobs=None,
        param_grid={'base_estimator__alpha': [1, 10000.0, 100000.0]},
        pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
        scoring=None, verbose=0)

```

Results

```

[75]: # let's get the predictions
x_pastelasso_train= paste_lasso_grid.predict(X_train)
x_pastelasso_test = paste_lasso_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_pastelasso_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train,
↳x_pastelasso_train))))
print('train r2: {}'.format(r2_score(y_train, x_pastelasso_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_pastelasso_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test,
↳x_pastelasso_test))))
print('test r2: {}'.format(r2_score(y_test, x_pastelasso_test)))
print()
print(f'Best Mean Cross Validation Score is {paste_lasso_grid.best_score_}')
print(f'Best parameters {paste_lasso_grid.best_params_}')
print(f'Train score is {paste_lasso_grid.score(X_train,y_train)}')
print(f'Test score is {paste_lasso_grid.score(X_test,y_test)}')

```

```

train mse: 552539222.3512007
train rmse: 23506.15286156373
train r2: 0.9115063901025299

```

```

test mse: 857849072.1153542
test rmse: 29289.0606219345
test r2: 0.87516957496597

```

```

Best Mean Cross Validation Score is 0.881849542095595
Best parameters {'base_estimator__alpha': 1}

```


Train score is 0.9115063901025299
Test score is 0.87516957496597

2.3 Random Forest

```
[76]: from sklearn.ensemble import RandomForestRegressor
```

```
[77]: rf =RandomForestRegressor(random_state=42)
      rf_param = {
          'n_estimators': range(1,300,10),#[200, 500],
          'max_features': ['auto', 'sqrt', 'log2'],
          'max_depth' : range(1,20,2)#[2,4,5,6,7,8],
      }

      rf_grid = GridSearchCV(rf, rf_param,cv=5, return_train_score=True)
      rf_grid.fit(X_train,y_train)
```

```
[77]: GridSearchCV(cv=5, error_score=nan,
                  estimator=RandomForestRegressor(bootstrap=True, ccp_alpha=0.0,
                                                    criterion='mse', max_depth=None,
                                                    max_features='auto',
                                                    max_leaf_nodes=None,
                                                    max_samples=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    n_estimators=100, n_jobs=None,
                                                    oob_score=False, random_state=42,
                                                    verbose=0, warm_start=False),
                  iid='deprecated', n_jobs=None,
                  param_grid={'max_depth': range(1, 20, 2),
                              'max_features': ['auto', 'sqrt', 'log2'],
                              'n_estimators': range(1, 300, 10)},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring=None, verbose=0)
```

2.3.1 Results

```
[78]: # let's get the predictions
      x_rfgrid_train= rf_grid.predict(X_train)
      x_rfgrid_test = rf_grid.predict(X_test)

      # check model performance:

      print('train mse: {}'.format(mean_squared_error(y_train, x_rfgrid_train)))
```

```

print('train mse: {}'.format(sqrt(mean_squared_error(y_train,
↳x_rfgrid_train))))
print('train r2: {}'.format(r2_score(y_train, x_rfgrid_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_rfgrid_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_rfgrid_test))))
print('test r2: {}'.format(r2_score(y_test, x_rfgrid_test)))
print()
print(f'Best Mean Cross Validation Score is {rf_grid.best_score_}')
print(f'Best parameters {rf_grid.best_params_}')
print(f'Train score is {rf_grid.score(X_train,y_train)}')
print(f'Test score is {rf_grid.score(X_test,y_test)}')

```

```

train mse: 122538951.61327057
train rmse: 11069.731325252233
train r2: 0.9803743992414402

```

```

test mse: 932265620.8565763
test rmse: 30533.025085251156
test r2: 0.8643408060007887

```

```

Best Mean Cross Validation Score is 0.8767927179403199
Best parameters {'max_depth': 13, 'max_features': 'sqrt', 'n_estimators': 291}
Train score is 0.9803743992414402
Test score is 0.8643408060007887

```

2.4 Adaboost Regressor

```
[79]: from sklearn.ensemble import AdaBoostRegressor
```

```

[80]: ada_
↳AdaBoostRegressor(base_estimator=DecisionTreeRegressor(),random_state=42,learning_rate=1.
↳0)
ada_param = {
    'base_estimator__max_depth':range(1,20,2),
    'n_estimators': range(1,300,10),
    'loss':['linear', 'square', 'exponential']
}

ada_grid = GridSearchCV(ada, ada_param,cv=5, return_train_score=True)
ada_grid.fit(X_train,y_train)

```

```

[80]: GridSearchCV(cv=5, error_score=nan,
    estimator=AdaBoostRegressor(base_estimator=DecisionTreeRegressor(ccp_alpha=0.0,
    criterion='mse',
    max_depth=None,
    max_features=None,

```

```

max_leaf_nodes=None,
min_impurity_decrease=0.0,
min_impurity_split=None,
min_samples_leaf=1,
min_samples_split=2,
min_weight_fraction_leaf=0.0,
presort='deprecated',
random_state=None,
splitter='best'),

                                learning_rate=1.0, loss='linear',
                                n_estimators=50, random_state=42),
                                iid='deprecated', n_jobs=None,
                                param_grid={'base_estimator__max_depth': range(1, 20, 2),
                                              'loss': ['linear', 'square', 'exponential'],
                                              'n_estimators': range(1, 300, 10)},
                                pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                                scoring=None, verbose=0)

```

2.4.1 Results

```

[81]: # let's get the predictions
x_adagrid_train= ada_grid.predict(X_train)
x_adagrid_test = ada_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_adagrid_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train,
↳x_adagrid_train))))
print('train r2: {}'.format(r2_score(y_train, x_adagrid_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_adagrid_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_adagrid_test))))
print('test r2: {}'.format(r2_score(y_test, x_adagrid_test)))
print()
print(f'Best Mean Cross Validation Score is {ada_grid.best_score_}')
print(f'Best parameters {ada_grid.best_params_}')
print(f'Train score is {ada_grid.score(X_train,y_train)}')
print(f'Test score is {ada_grid.score(X_test,y_test)}')

```

```

train mse: 59039229.255843244
train rmse: 7683.698930583059
train r2: 0.9905443915815029

```

```

test mse: 888472231.5703418
test rmse: 29807.251325312467
test r2: 0.8707134274513209

```

```
Best Mean Cross Validation Score is 0.8819499238825339
Best parameters {'base_estimator__max_depth': 17, 'loss': 'square',
'n_estimators': 291}
Train score is 0.9905443915815029
Test score is 0.8707134274513209
```

2.5 Gradient Boost Regressor

```
[82]: from sklearn.ensemble import GradientBoostingRegressor
```

```
[83]: gboostparam = {'max_depth':range(1,12),
                    'n_estimators':range(45,60,5),
                    'loss':['ls', 'lad', 'huber', 'quantile'],
                    'max_features':['auto','sqrt','log2']}

gboost = GridSearchCV(estimator = GradientBoostingRegressor(learning_rate=0.1,
                                                            min_samples_split=2,
                                                            min_samples_leaf=1,
                                                            subsample=1,
                                                            #
                                                            max_features='sqrt',
                                                            random_state=10),
                    param_grid = gboostparam,
                    n_jobs=-1,cv=5)

gboost.fit(X_train,y_train)
```

```
[83]: GridSearchCV(cv=5, error_score=nan,
                  estimator=GradientBoostingRegressor(alpha=0.9, ccp_alpha=0.0,
                                                         criterion='friedman_mse',
                                                         init=None, learning_rate=0.1,
                                                         loss='ls', max_depth=3,
                                                         max_features=None,
                                                         max_leaf_nodes=None,
                                                         min_impurity_decrease=0.0,
                                                         min_impurity_split=None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=2,
                                                         min_weight_fraction_leaf=0.0,
                                                         n_estimators=100,
                                                         n_iter_n...
                                                         presort='deprecated',
                                                         random_state=10, subsample=1,
                                                         tol=0.0001,
                                                         validation_fraction=0.1,
                                                         verbose=0, warm_start=False),
```

```

iid='deprecated', n_jobs=-1,
param_grid={'loss': ['ls', 'lad', 'huber', 'quantile'],
            'max_depth': range(1, 12),
            'max_features': ['auto', 'sqrt', 'log2'],
            'n_estimators': range(45, 60, 5)},
pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
scoring=None, verbose=0)

```

2.5.1 Results

```

[84]: # let's get the predictions
x_gboost_train= gboost.predict(X_train)
x_gboost_test = gboost.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_gboost_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train,
→x_gboost_train))))
print('train r2: {}'.format(r2_score(y_train, x_gboost_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_gboost_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_gboost_test))))
print('test r2: {}'.format(r2_score(y_test, x_gboost_test)))
print()
print(f'Best Mean Cross Validation Score is {gboost.best_score_}')
print(f'Best parameters {gboost.best_params_}')
print(f'Train score is {gboost.score(X_train,y_train)}')
print(f'Test score is {gboost.score(X_test,y_test)}')

```

```

train mse: 210402353.29194453
train rmse: 14505.252610414771
train r2: 0.9663023672880685

```

```

test mse: 1106890497.8616946
test rmse: 33269.96389931457
test r2: 0.8389301617200744

```

```

Best Mean Cross Validation Score is 0.8918768459494615
Best parameters {'loss': 'ls', 'max_depth': 5, 'max_features': 'log2',
'n_estimators': 55}
Train score is 0.9663023672880685
Test score is 0.8389301617200744

```

2.6 Extra Trees

```
[85]: from sklearn.ensemble import ExtraTreesRegressor
```

```
[86]: extree= ExtraTreesRegressor(random_state=42)
etr_param = {
    'n_estimators': range(2,250,5),
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth' : range(2,20,5),
}
etr_grid = GridSearchCV(extree,
    ↪etr_param,cv=5,n_jobs=-1,return_train_score=True, )
etr_grid.fit(X_train,y_train)
```

```
[86]: GridSearchCV(cv=5, error_score=nan,
                  estimator=ExtraTreesRegressor(bootstrap=False, ccp_alpha=0.0,
                                                  criterion='mse', max_depth=None,
                                                  max_features='auto',
                                                  max_leaf_nodes=None,
                                                  max_samples=None,
                                                  min_impurity_decrease=0.0,
                                                  min_impurity_split=None,
                                                  min_samples_leaf=1,
                                                  min_samples_split=2,
                                                  min_weight_fraction_leaf=0.0,
                                                  n_estimators=100, n_jobs=None,
                                                  oob_score=False, random_state=42,
                                                  verbose=0, warm_start=False),
                  iid='deprecated', n_jobs=-1,
                  param_grid={'max_depth': range(2, 20, 5),
                              'max_features': ['auto', 'sqrt', 'log2'],
                              'n_estimators': range(2, 250, 5)},
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring=None, verbose=0)
```

2.6.1 Results

```
[87]: # let's get the predictions
x_etr_train= etr_grid.predict(X_train)
x_etr_test = etr_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_etr_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_etr_train))))
print('train r2: {}'.format(r2_score(y_train, x_etr_train)))
print()
```

```

print('test mse: {}'.format(mean_squared_error(y_test, x_etr_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_etr_test))))
print('test r2: {}'.format(r2_score(y_test, x_etr_test)))
print()
print(f'Best Mean Cross Validation Score is {etr_grid.best_score_}')
print(f'Best parameters {etr_grid.best_params_}')
print(f'Train score is {etr_grid.score(X_train,y_train)}')
print(f'Test score is {etr_grid.score(X_test,y_test)}')

```

```

train mse: 5741824.059904766
train rmse: 2396.210353851424
train r2: 0.9990804005979974

```

```

test mse: 929787656.6985651
test rmse: 30492.419659623032
test r2: 0.8647013884495182

```

```

Best Mean Cross Validation Score is 0.8804371492768851
Best parameters {'max_depth': 17, 'max_features': 'sqrt', 'n_estimators': 162}
Train score is 0.9990804005979974
Test score is 0.8647013884495182

```

2.7 XGBoost Regressor

```
[88]: # pip install xgboost
```

```
[89]: from xgboost import XGBRegressor
```

```

[90]: xgbr= XGBRegressor(random_state=42)
xgbr_param = {
    'max_depth' : range(2,20,2),
    'n_estimators' : range(2,200,10),
    'learning_rate' : [0.1],
    'min_child_weight' : range(1,8,1),
    'subsample': [0.6,0.7,0.8,0.9,1]
}
xgbr_grid = GridSearchCV(xgbr, xgbr_param,cv=5,n_jobs=-1,
    ↪return_train_score=True)
xgbr_grid.fit(X_train,y_train)

```

```

[90]: GridSearchCV(cv=5, error_score=nan,
    estimator=XGBRegressor(base_score=None, booster=None,
        colsample_bylevel=None,
        colsample_bynode=None,
        colsample_bytree=None, gamma=None,
        gpu_id=None, importance_type='gain',
        interaction_constraints=None,

```

```

        learning_rate=None, max_delta_step=None,
        max_depth=None, min_child_weight=None,
        missing=nan, monotone_constraints=None,
        n_estima...
        reg_lambda=None, scale_pos_weight=None,
        subsample=None, tree_method=None,
        validate_parameters=False, verbosity=None),
iid='deprecated', n_jobs=-1,
param_grid={'learning_rate': [0.1], 'max_depth': range(2, 20, 2),
            'min_child_weight': range(1, 8),
            'n_estimators': range(2, 200, 10),
            'subsample': [0.6, 0.7, 0.8, 0.9, 1]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=None, verbose=0)

```

```

[91]: # let's get the predictions
x_xgbr_train= xgbr_grid.predict(X_train)
x_xgbr_test = xgbr_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_xgbr_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_xgbr_train))))
print('train r2: {}'.format(r2_score(y_train, x_xgbr_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_xgbr_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_xgbr_test))))
print('test r2: {}'.format(r2_score(y_test, x_xgbr_test)))
print()
print(f'Best Mean Cross Validation Score is {xgbr_grid.best_score_}')
print(f'Best parameters {xgbr_grid.best_params_}')
print(f'Train score is {xgbr_grid.score(X_train,y_train)}')
print(f'Test score is {xgbr_grid.score(X_test,y_test)}')

```

```

train mse: 98739443.27826326
train rmse: 9936.772276663245
train r2: 0.9841860823241139

```

```

test mse: 847173787.4212183
test rmse: 29106.249971805337
test r2: 0.8767229954557099

```

```

Best Mean Cross Validation Score is 0.8984774624760332
Best parameters {'learning_rate': 0.1, 'max_depth': 4, 'min_child_weight': 1,
'n_estimators': 192, 'subsample': 0.8}
Train score is 0.9841860823241139
Test score is 0.8767229954557099

```


2.8 Summary

```
[92]: regressors={#'LinearReg':lr,
    'Ridge':grid_ridge,
    'Lasso':grid_lasso,
    'LinearSGD':grid_linearsgd,
    'Polynomial':grid_poly,
    'elasticnet':grid_elasticnet,
    'GridLRRidge':grid_search_linear,
    'GridLRLasso':grid_search_linear2,
    'GridRF':grid_rf,
    # 'KNN':knnreg,
    'GridElastic':grid_elasticnet,
    'GridDecisionTree':grid_dtree,
    'SVC':grid_svc_kernel,
    'BaggingDTree':bag_dtree1_grid,
    'BaggingLasso':bag_lasso_grid,
    'PastingDTree':paste_dtree1_grid,
    'PastingLasso':paste_lasso_grid,
    'RandomForest':rf_grid,
    'AdaBoost':rf_grid,
    'GradientBoostR':gboost,
    'ExtraTrees':etr_grid,
    'XGBoost':xgbr_grid,
}
```

```
[93]: regressors.keys()
```

```
[93]: dict_keys(['Ridge', 'Lasso', 'LinearSGD', 'Polynomial', 'elasticnet',
    'GridLRRidge', 'GridLRLasso', 'GridRF', 'GridElastic', 'GridDecisionTree',
    'SVC', 'BaggingDTree', 'BaggingLasso', 'PastingDTree', 'PastingLasso',
    'RandomForest', 'AdaBoost', 'GradientBoostR', 'ExtraTrees', 'XGBoost'])
```

```
[94]: results_mean_std = []
for key, value in regressors.items():
    mean = value.cv_results_['mean_test_score'][value.best_index_]
    std=value.cv_results_['std_test_score'][value.best_index_]

    results_mean_std.append({
        "model": key,
        "mean": mean,
        "std": std
    })
```

```
[95]: # Create a Pandas DataFrame with the mean+std results
accuracy_df = pd.DataFrame(results_mean_std, columns=['model', 'mean', 'std'])
```

```
[96]: # Show the accuracy dataframe

accuracy_df.sort_values(by=['mean'], inplace=True, ascending=False)
accuracy_df
```

```
[96]:
```

	model	mean	std
19	XGBoost	8.984775e-01	8.966995e-03
17	GradientBoostR	8.918768e-01	1.466126e-02
12	BaggingLasso	8.849436e-01	1.222205e-02
0	Ridge	8.843424e-01	3.178374e-02
1	Lasso	8.843050e-01	3.179290e-02
4	elasticnet	8.840700e-01	1.039538e-02
8	GridElastic	8.840700e-01	1.039538e-02
5	GridLRRidge	8.818970e-01	1.141903e-02
6	GridLRLasso	8.818495e-01	1.142073e-02
14	PastingLasso	8.818495e-01	1.142073e-02
18	ExtraTrees	8.804371e-01	1.394497e-02
15	RandomForest	8.767927e-01	1.611594e-02
16	AdaBoost	8.767927e-01	1.611594e-02
2	LinearSGD	8.765910e-01	1.366770e-02
11	BaggingDTree	8.701118e-01	1.681030e-02
13	PastingDTree	8.572589e-01	1.829952e-02
9	GridDecisionTree	7.573303e-01	2.847110e-02
10	SVC	7.104416e-01	2.673301e-02
7	GridRF	-1.253123e+21	2.506247e+21
3	Polynomial	-9.790049e+22	1.950468e+23

```
[97]: # Create a prediction of all models on the test set
predictions_all = {}
for key, value in regressors.items():
    # Get best estimator
    best_model = value.best_estimator_

    # Predict test labels
    predictions = best_model.predict(X_test)

    # Save predictions to a list
    predictions_all[key] = predictions
```

```
[98]: # Creat a DataFrame for the predictions
pred = pd.DataFrame(predictions_all)
```

```
[99]: pred.head()
```

```
[99]:
```

	Ridge	Lasso	LinearSGD	Polynomial	elasticnet	\
0	259427.918908	259358.714345	269968.912408	-2.401626e+15	263655.245673	
1	130977.047829	130977.454966	136605.961715	-8.888887e+14	130804.224834	

2	118352.334635	118328.412013	120844.488906	2.467426e+14	120690.792966
3	219592.797874	219614.219922	219618.652969	1.198736e+15	217482.466827
4	91161.872462	91146.795565	93556.522122	2.906841e+14	92806.958086

	GridLRRidge	GridLRLasso	GridRF	GridElastic	\
0	259427.918908	259358.714345	258576.640044	263655.245673	
1	130977.047829	130977.454966	130385.243135	130804.224834	
2	118352.334635	118328.412013	117989.111328	120690.792966	
3	219592.797874	219614.219922	220890.140789	217482.466827	
4	91161.872462	91146.795565	91502.241892	92806.958086	

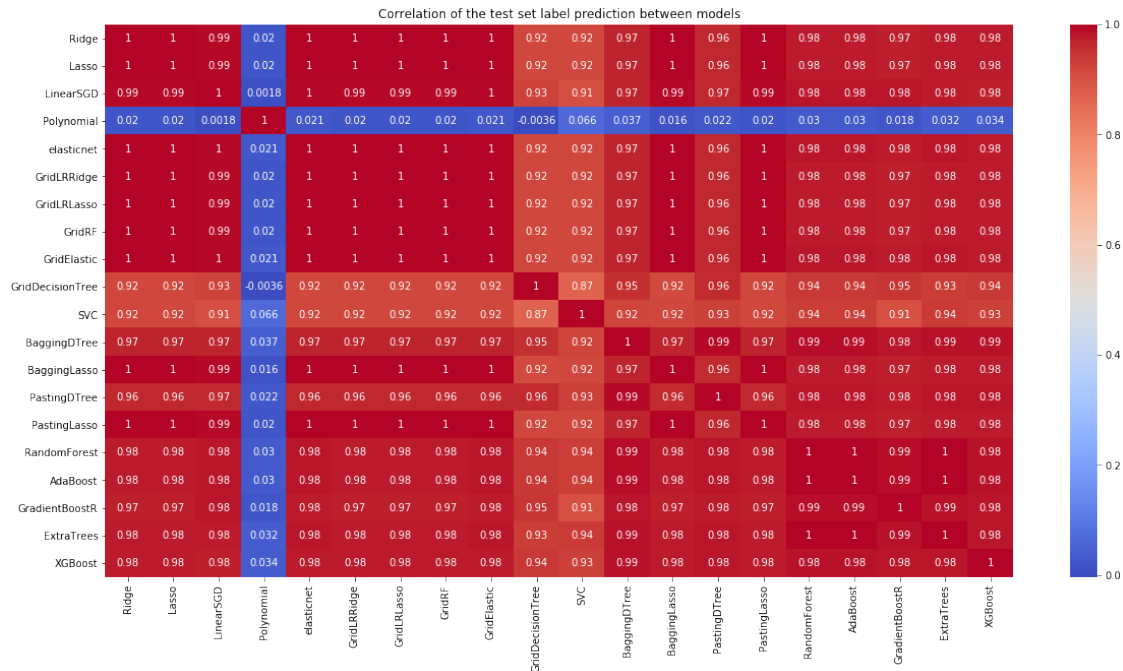
	GridDecisionTree	SVC	BaggingDTree	BaggingLasso	PastingDTree	\
0	286141.333333	237000	231628.891408	253304.203849	260380.403015	
1	148118.800000	105000	138410.074763	135019.758648	142482.141309	
2	118911.000000	135000	121417.290889	118712.599413	118725.602426	
3	229747.629630	178000	234395.674515	218284.457610	239981.004528	
4	118911.000000	113000	97929.805069	88489.935148	98999.445234	

	PastingLasso	RandomForest	AdaBoost	GradientBoostR	ExtraTrees	\
0	259358.714345	252395.782594	252395.782594	279172.467178	256034.145142	
1	130977.454966	140416.801495	140416.801495	144826.895238	136406.530809	
2	118328.412013	125970.464256	125970.464256	126477.951094	132510.120732	
3	219614.219922	209285.267185	209285.267185	210454.548365	201347.542099	
4	91146.795565	100051.392538	100051.392538	100665.764154	96968.065484	

	XGBoost
0	273605.375000
1	134782.234375
2	118255.343750
3	221607.421875
4	93816.242188

```
[100]: # Plot a heatmap of all correlations for easier visualization
fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(20,10))
g = sns.heatmap(pred.corr(), annot=True, cmap='coolwarm', ax=ax)
g.set_title('Correlation of the test set label prediction between models')
```

```
[100]: Text(0.5, 1, 'Correlation of the test set label prediction between models')
```



```
[101]: def get_redundant_pairs(df):
        '''Get diagonal and lower triangular pairs of correlation matrix'''
        pairs_to_drop = set()
        cols = df.columns
        for i in range(0, df.shape[1]):
            for j in range(0, i+1):
                pairs_to_drop.add((cols[i], cols[j]))
        return pairs_to_drop

def get_top_abs_correlations(df, n=5):
    au_corr = df.corr().abs().unstack()
    labels_to_drop = get_redundant_pairs(df)
    au_corr = au_corr.drop(labels=labels_to_drop).sort_values(ascending=True)
    return au_corr[0:n]

[102]: print("Top Least Correlations")
        print(get_top_abs_correlations(pred, 5))
```

```
Top Least Correlations
LinearSGD    Polynomial    0.001773
Polynomial  GridDecisionTree 0.003570
            BaggingLasso    0.015869
            GradientBoostR  0.018109
            GridRF          0.019594
dtype: float64
```

2.9 Stacking

2.10 Voting top 5

```
[103]: from sklearn.ensemble import VotingRegressor
```

```
[104]: vclf1 = VotingRegressor( estimators=[
                                                ('XGBoost', xgbr_grid.best_estimator_),
                                                ('GradientBoost', gboost.best_estimator_),
                                                ('BaggingLasso', bag_lasso_grid.best_estimator_),
                                                ('PastingDTree', bag_lasso_grid.best_estimator_),
                                                ('Ridge', grid_ridge.best_estimator_)
                                            ] )

vclf1_param = { }
vclf1_grid = GridSearchCV(vclf1, vclf1_param, cv=5, return_train_score=True)
vclf1_grid.fit(X_train, y_train)
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
```

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-

packages\sklearn\ensemble_bagging.py:1056: UserWarning: Some inputs do not have OOB scores. This probably means too few estimators were used to compute any reliable oob estimates.

```
warn("Some inputs do not have OOB scores. "
```

```
[104]: GridSearchCV(cv=5, error_score=nan,
                  estimator=VotingRegressor(estimators=[('XGBoost',
                                                         XGBRegressor(base_score=0.5,
                                                                       booster=None,
                                                                       gamma=0,
                                                                       gpu_id=-1,
                                                                       importance_type='gain',
                                                                       interaction_constraints=None,
                                                                       learning_rate=0.1,
                                                                       max_delta_step=0,
                                                                       max_depth=4,
                                                                       min_child_weight=1,
                                                                       missing=nan,
                                                                       monotone_constraints=
```

```

n_jobs=None,
oob_score=True,
random_state=0,
warm_start=False)),
verbose=0,
('Ridge',
 Ridge(alpha=1, copy_X=True,
        fit_intercept=True,
        max_iter=None,
        normalize=False,
        random_state=None,
        solver='auto',
        tol=0.001))),
n_jobs=None, weights=None),
iid='deprecated', n_jobs=None, param_grid={},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=None, verbose=0)

```

```

[105]: # let's get the predictions
x_vclf1_train= vclf1_grid.predict(X_train)
x_vclf1_test = vclf1_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_vclf1_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_vclf1_train))))
print('train r2: {}'.format(r2_score(y_train, x_vclf1_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_vclf1_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_vclf1_test))))
print('test r2: {}'.format(r2_score(y_test, x_vclf1_test)))
print()
print(f'Best Mean Cross Validation Score is {vclf1_grid.best_score_}')
print(f'Best parameters {vclf1_grid.best_params_}')
print(f'Train score is {vclf1_grid.score(X_train,y_train)}')
print(f'Test score is {vclf1_grid.score(X_test,y_test)}')

```

```

train mse: 317633603.336003
train rmse: 17822.2782869083
train r2: 0.9491284183151113

```

```

test mse: 822879018.3587837
test rmse: 28685.867920611774
test r2: 0.8802582634262037

```

```

Best Mean Cross Validation Score is 0.8999858356957512
Best parameters {}

```

Train score is 0.9491284183151113
Test score is 0.8802582634262037

2.11 Voting Least Correlated

```
[106]: vclf3 = VotingRegressor(estimators=
        [ ('LinearSGD', grid_linearsgd.best_estimator_),
          ('Polynomial', grid_poly.best_estimator_),
          ('GridDecisionTree', grid_dtree.best_estimator_),
          ('BaggingLasso', bag_lasso_grid.best_estimator_),
          ('GradientBoostR', gboost.best_estimator_)
        ] )

vclf3_param = { }
vclf3_grid = GridSearchCV(vclf3, vclf3_param, cv=5, return_train_score=True)
vclf3_grid.fit(X_train, y_train)
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```

```
C:\Users\nabhs\Anaconda3\envs\buan6341_2020\lib\site-
packages\sklearn\ensemble\_bagging.py:1056: UserWarning: Some inputs do not have
OOB scores. This probably means too few estimators were used to compute any
reliable oob estimates.
```

```
warn("Some inputs do not have OOB scores. "
```



```
[106]: GridSearchCV(cv=5, error_score=nan,
                  estimator=VotingRegressor(estimators=[('LinearSGD',
                                                         Pipeline(memory=None,
                                                                    steps=[('scaler',
MinMaxScaler(copy=True,
              feature_range=(0,
                              1))),
                                                         ('sgd_reg',
SGDRegressor(alpha=0.01,
              average=False,
              early_stopping=False,
              epsilon=0.1,
              eta0=0.05,
              fit_intercept=True,
              l1_ratio=0.15,
              learning_rate='invscaling',
              loss='squared_loss',
              max_iter=1...
min_samples_split=2,
min_weight_fraction_leaf=0.0,
n_estimators=55,
n_iter_no_change=None,
presort='deprecated',
random_state=10,
subsample=1,
tol=0.0001,
validation_fraction=0.1,
verbose=0,
warm_start=False))],
                  n_jobs=None, weights=None),
        iid='deprecated', n_jobs=None, param_grid={},
        pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
        scoring=None, verbose=0)
```

```
[107]: # let's get the predictions
x_vclf3_train= vclf3_grid.predict(X_train)
x_vclf3_test = vclf3_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_vclf3_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train, x_vclf3_train))))
print('train r2: {}'.format(r2_score(y_train, x_vclf3_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_vclf3_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_vclf3_test))))
print('test r2: {}'.format(r2_score(y_test, x_vclf3_test)))
```

```

print()
print(f'Best Mean Cross Validation Score is {vclf3_grid.best_score_}')
print(f'Best parameters {vclf3_grid.best_params_}')
print(f'Train score is {vclf3_grid.score(X_train,y_train)}')
print(f'Test score is {vclf3_grid.score(X_test,y_test)}')

```

```

train mse: 282747004.4125732
train rmse: 16815.08264661739
train r2: 0.954715788316905

```

```

test mse: 8.592991133805444e+28
test rmse: 293138041437911.0
test r2: -1.2504142866314648e+19

```

```

Best Mean Cross Validation Score is -2.7187395878116203e+20
Best parameters {}
Train score is 0.954715788316905
Test score is -1.2504142866314648e+19

```

2.12 Stacking Top 5

<https://towardsdatascience.com/stacking-made-easy-with-sklearn-e27a0793c92b>

```
[108]: from sklearn.ensemble import StackingRegressor
```

```
[109]: sregf1 = StackingRegressor(estimators=
    [ ('GradientBoost', gboost.best_estimator_),
      ('BaggingLasso', bag_lasso_grid.best_estimator_),
      ('PastingDTree', bag_lasso_grid.best_estimator_),
      ('Ridge', grid_ridge.best_estimator_)
    ] , final_estimator=XGBRegressor())

sregf1_param = {
    'final_estimator__learning_rate' : [0.1],
    'final_estimator__max_depth': [4],
    'final_estimator__min_child_weight': [1],
    'final_estimator__n_estimators': [192],
    'final_estimator__subsample': [0.8],
    'n_jobs': [-1]
}

sregf1_grid = GridSearchCV(sregf1, sregf1_param, cv=5, return_train_score=True )
sregf1_grid.fit(X_train,y_train)

```

```
[109]: GridSearchCV(cv=5, error_score=nan,
    estimator=StackingRegressor(cv=None,
    estimators=[('GradientBoost',
    GradientBoostingRegressor(alpha=0.9,
    ccp_alpha=0.0,
    criterion='friedman_mse',

```

```

init=None,
learning_rate=0.1,
loss='ls',
max_depth=5,
max_features='log2',
max_leaf_nodes=None,
min_impurity_decrease=0.0,
min_impurity_split=None,
min_samples_leaf=1,
min_samples_split...

n_jobs=None, passthrough=False,
verbose=0),

iid='deprecated', n_jobs=None,
param_grid={'final_estimator__learning_rate': [0.1],
            'final_estimator__max_depth': [4],
            'final_estimator__min_child_weight': [1],
            'final_estimator__n_estimators': [192],
            'final_estimator__subsample': [0.8], 'n_jobs': [-1]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=None, verbose=0)

```

```

[110]: # let's get the predictions
x_sregf1_train= sregf1_grid.predict(X_train)
x_sregf1_test = sregf1_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_sregf1_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train,
↳x_sregf1_train))))
print('train r2: {}'.format(r2_score(y_train, x_sregf1_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_sregf1_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_sregf1_test))))
print('test r2: {}'.format(r2_score(y_test, x_sregf1_test)))
print()
print(f'Best Mean Cross Validation Score is {sregf1_grid.best_score_}')
print(f'Best parameters {sregf1_grid.best_params_}')
print(f'Train score is {sregf1_grid.score(X_train,y_train)}')
print(f'Test score is {sregf1_grid.score(X_test,y_test)}')

```

```

train mse: 408279864.3139851
train rmse: 20205.93636320735
train r2: 0.9346106890152514

```

```

test mse: 825035315.361602
test rmse: 28723.427987648025

```

```
test r2: 0.8799444885675358
```

Best Mean Cross Validation Score is 0.87774060712915

```
Best parameters {'final_estimator__learning_rate': 0.1,
'final_estimator__max_depth': 4, 'final_estimator__min_child_weight': 1,
'final_estimator__n_estimators': 192, 'final_estimator__subsample': 0.8,
'n_jobs': -1}
```

Train score is 0.9346106890152514

Test score is 0.8799444885675358

2.13 Stacking Least Correlated

```
[111]: sregf2 = StackingRegressor(estimators=
                                     [ ('LinearSGD', grid_linearsgd.best_estimator_),
                                       ('Polynomial', grid_poly.best_estimator_),
                                       ('GridDecisionTree', grid_dtree.best_estimator_),
                                       ('BaggingLasso', bag_lasso_grid.best_estimator_),
                                       ('GradientBoostR', gboost.best_estimator_)
                                     ],
                                     final_estimator=XGBRegressor())
sregf2_param = {
    'final_estimator__learning_rate' : [0.1],
    'final_estimator__max_depth': [4],
    'final_estimator__min_child_weight': [1],
    'final_estimator__n_estimators': [192],
    'final_estimator__subsample': [0.8],
    'n_jobs': [-1]
}
sregf2_grid = GridSearchCV(sregf2, sregf2_param, cv=5, return_train_score=True)
sregf2_grid.fit(X_train, y_train)
```

```
[111]: GridSearchCV(cv=5, error_score=nan,
                    estimator=StackingRegressor(cv=None,
                                                  estimators=[ ('LinearSGD',
                                                                Pipeline(memory=None,
                                                                      steps=[ ('scaler',
                                                                MinMaxScaler(copy=True,
                                                                      feature_range=(0,
                                                                      1)))],
                                                                ('sgd_reg',
                                                                SGDRegressor(alpha=0.01,
                                                                      average=False,
                                                                      early_stopping=False,
                                                                      epsilon=0.1,
                                                                      eta0=0.05,
                                                                      fit_intercept=True,
                                                                      l1_ratio=0.15,
```

```

learning_rate='invscaling',
loss='squared_loss',...

n_jobs=None, passthrough=False,
verbose=0),

iid='deprecated', n_jobs=None,
param_grid={'final_estimator__learning_rate': [0.1],
            'final_estimator__max_depth': [4],
            'final_estimator__min_child_weight': [1],
            'final_estimator__n_estimators': [192],
            'final_estimator__subsample': [0.8], 'n_jobs': [-1]},
pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
scoring=None, verbose=0)

```

```

[112]: # let's get the predictions
x_sregf2_train= sregf2_grid.predict(X_train)
x_sregf2_test = sregf2_grid.predict(X_test)

# check model performance:

print('train mse: {}'.format(mean_squared_error(y_train, x_sregf2_train)))
print('train rmse: {}'.format(sqrt(mean_squared_error(y_train,
→x_sregf2_train))))
print('train r2: {}'.format(r2_score(y_train, x_sregf2_train)))
print()
print('test mse: {}'.format(mean_squared_error(y_test, x_sregf2_test)))
print('test rmse: {}'.format(sqrt(mean_squared_error(y_test, x_sregf2_test))))
print('test r2: {}'.format(r2_score(y_test, x_sregf2_test)))
print()
print(f'Best Mean Cross Validation Score is {sregf2_grid.best_score_}')
print(f'Best parameters {sregf2_grid.best_params_}')
print(f'Train score is {sregf2_grid.score(X_train,y_train)}')
print(f'Test score is {sregf2_grid.score(X_test,y_test)}')

```

```

train mse: 355270884.9509614
train rmse: 18848.630850832677
train r2: 0.9431005043099072

```

```

test mse: 685715224.4673892
test rmse: 26186.1647529261
test r2: 0.9002177356076231

```

```

Best Mean Cross Validation Score is 0.8665230017810973
Best parameters {'final_estimator__learning_rate': 0.1,
'final_estimator__max_depth': 4, 'final_estimator__min_child_weight': 1,
'final_estimator__n_estimators': 192, 'final_estimator__subsample': 0.8,
'n_jobs': -1}
Train score is 0.9431005043099072

```

Test score is 0.9002177356076231